

Computer Architecture Introduction

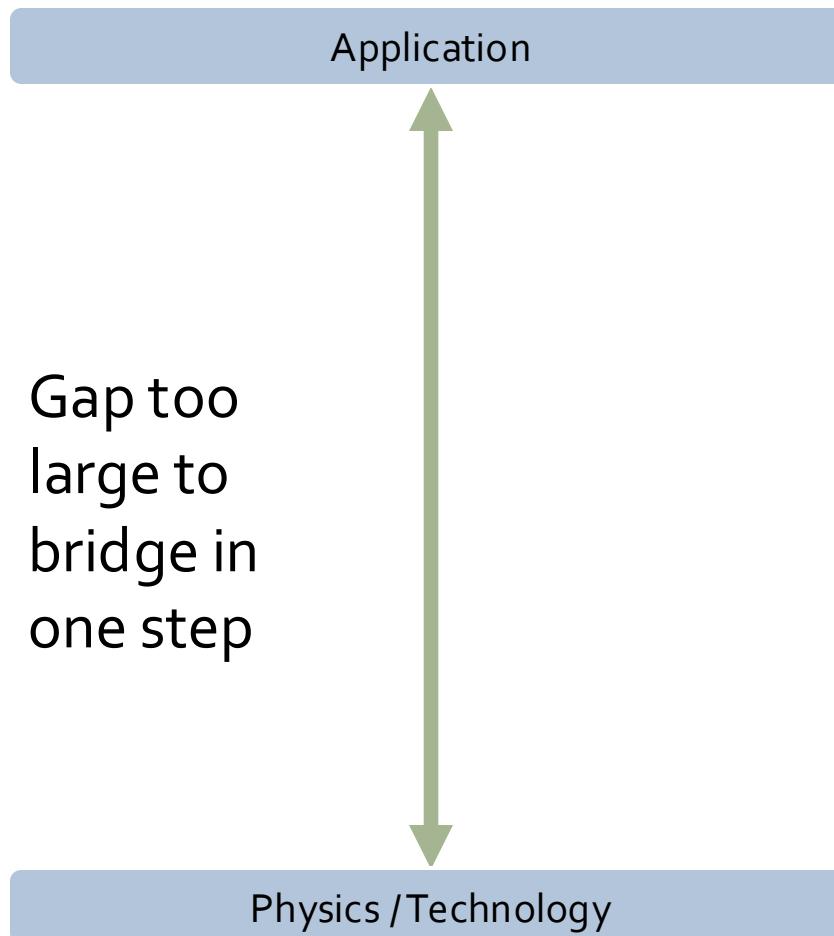
Ting-Jung Chang

NYCU CS

Agenda

- What is Computer Architecture?
- Course Admin
- ISA Review

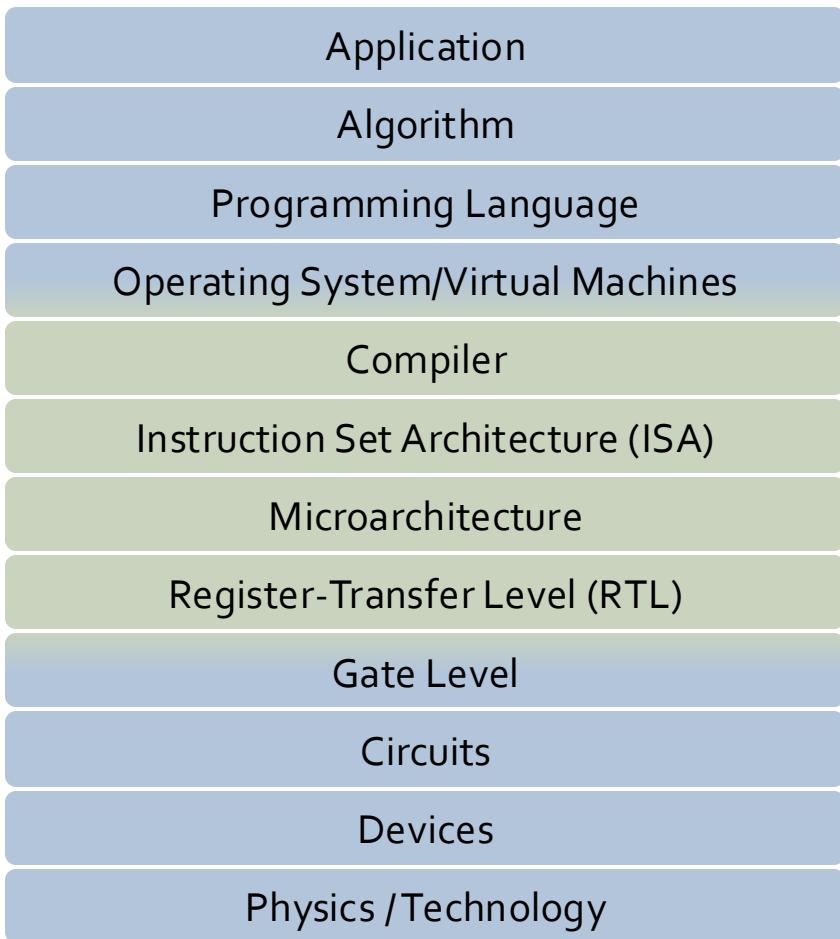
What is Computer Architecture?



In its broadest definition, computer architecture is the **design of the abstraction/implementation layers** that allow us to execute information processing **applications** efficiently using manufacturing **technologies**.

The Computer Systems Stack

Computer Architecture



Sort an array of numbers

$2,6,3,8,4,5 \rightarrow 2,3,4,5,6,8$

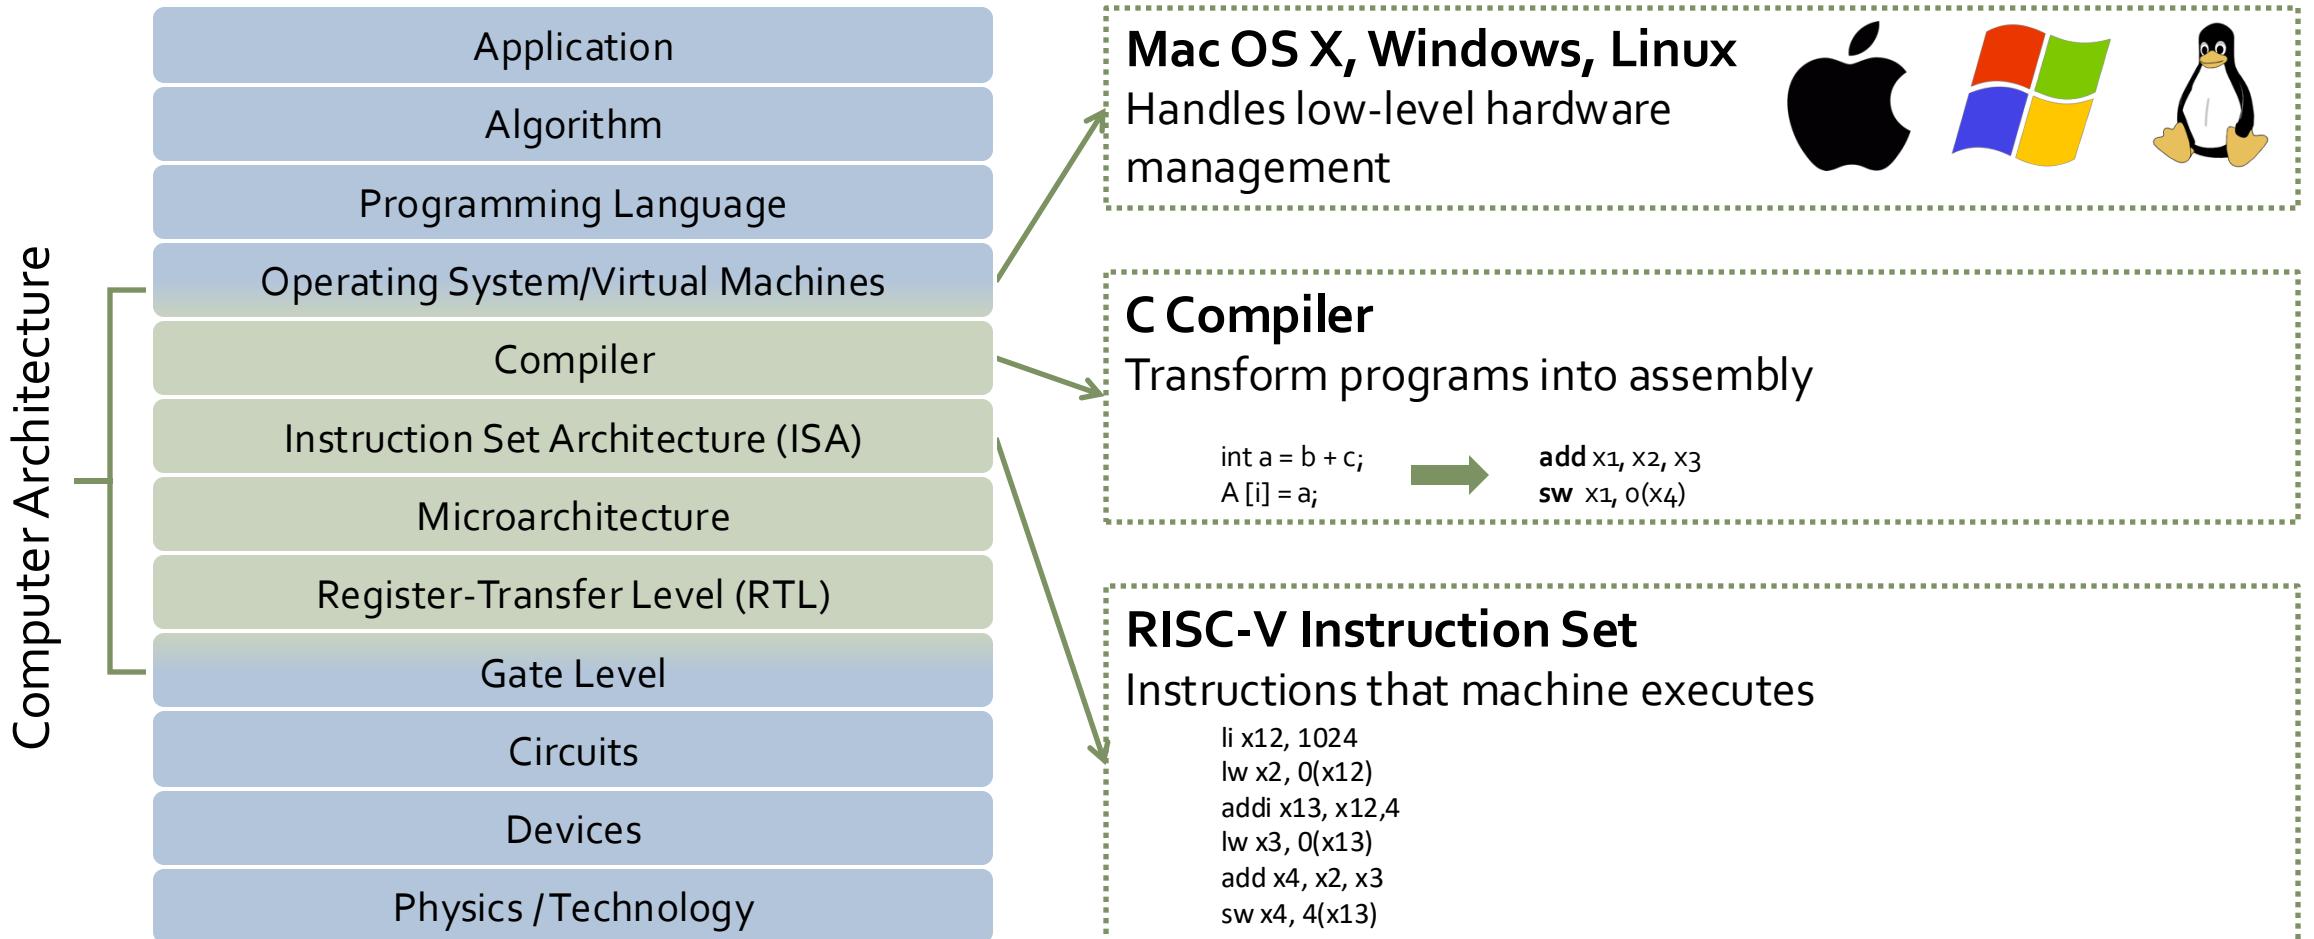
Insertion sort algorithm

1. Find minimum number in input array
2. Move minimum number into output array
3. Repeat steps 1 and 2 until finished

C implementation of insertion sort

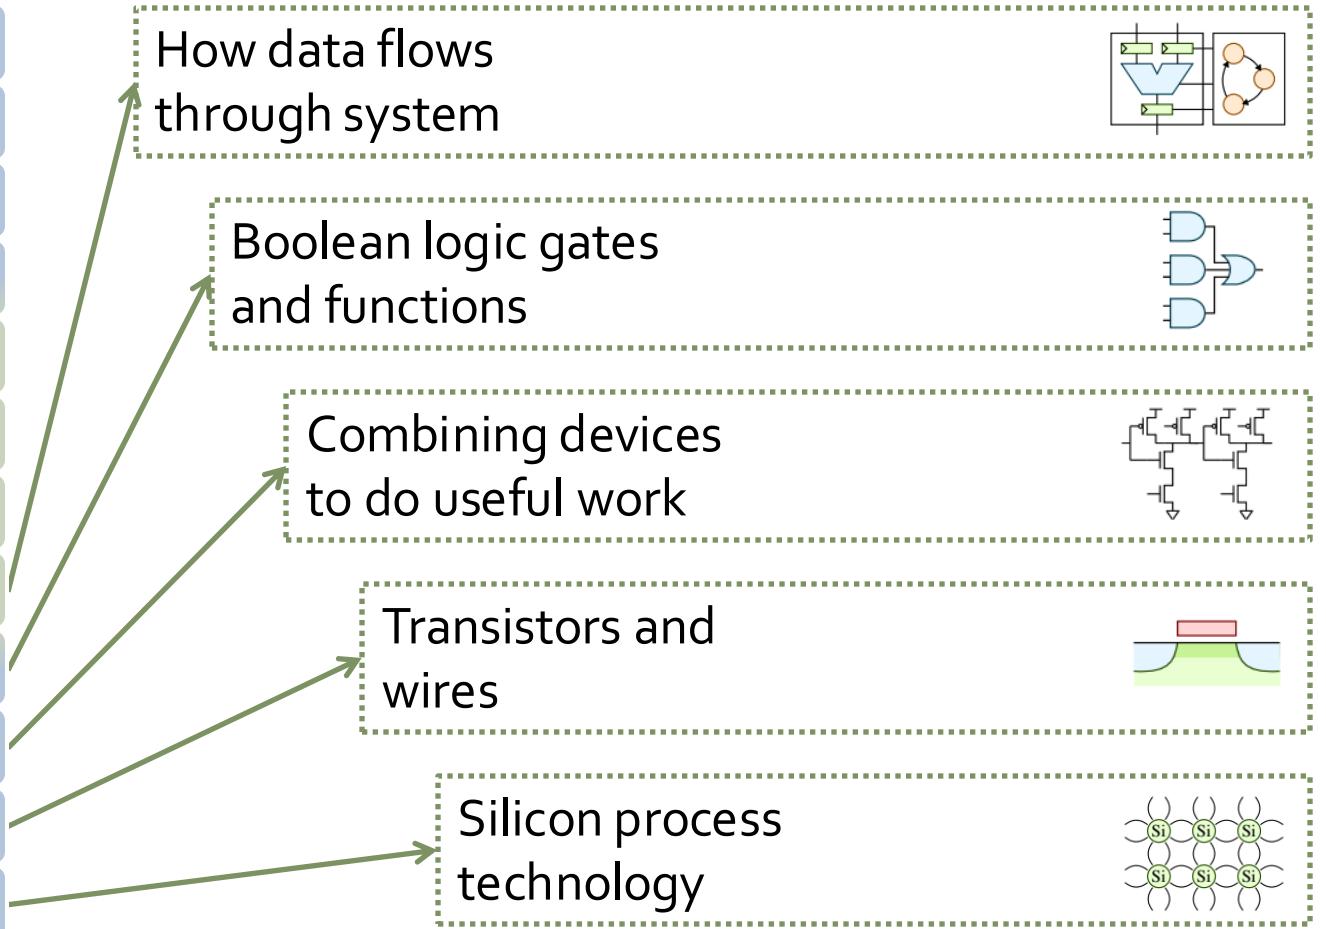
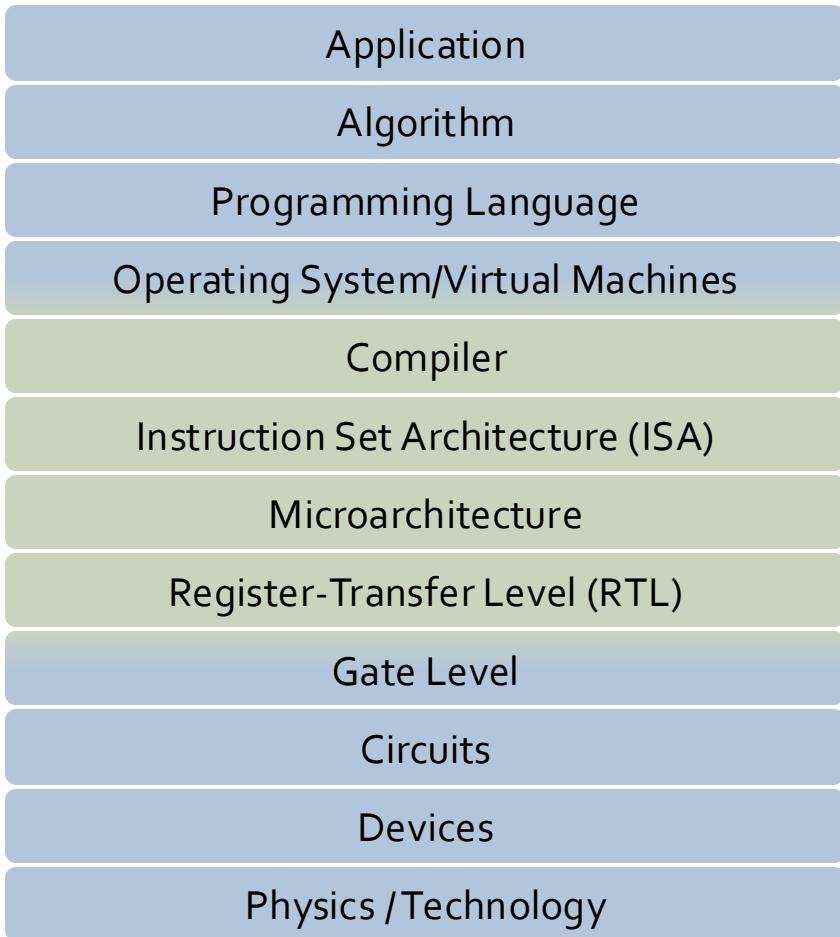
```
void isort( int b[], int a[], int n ) {
    for( int idx, k = 0; k < n; k++ ){
        int min = 100;
        for( int i = 0; i < n; i++ ){
            if( a[i] < min ){
                min = a[i];
                idx = i;
            }
        }
        b[k] = min;
        a[idx] = 100;
    }
}
```

The Computer Systems Stack

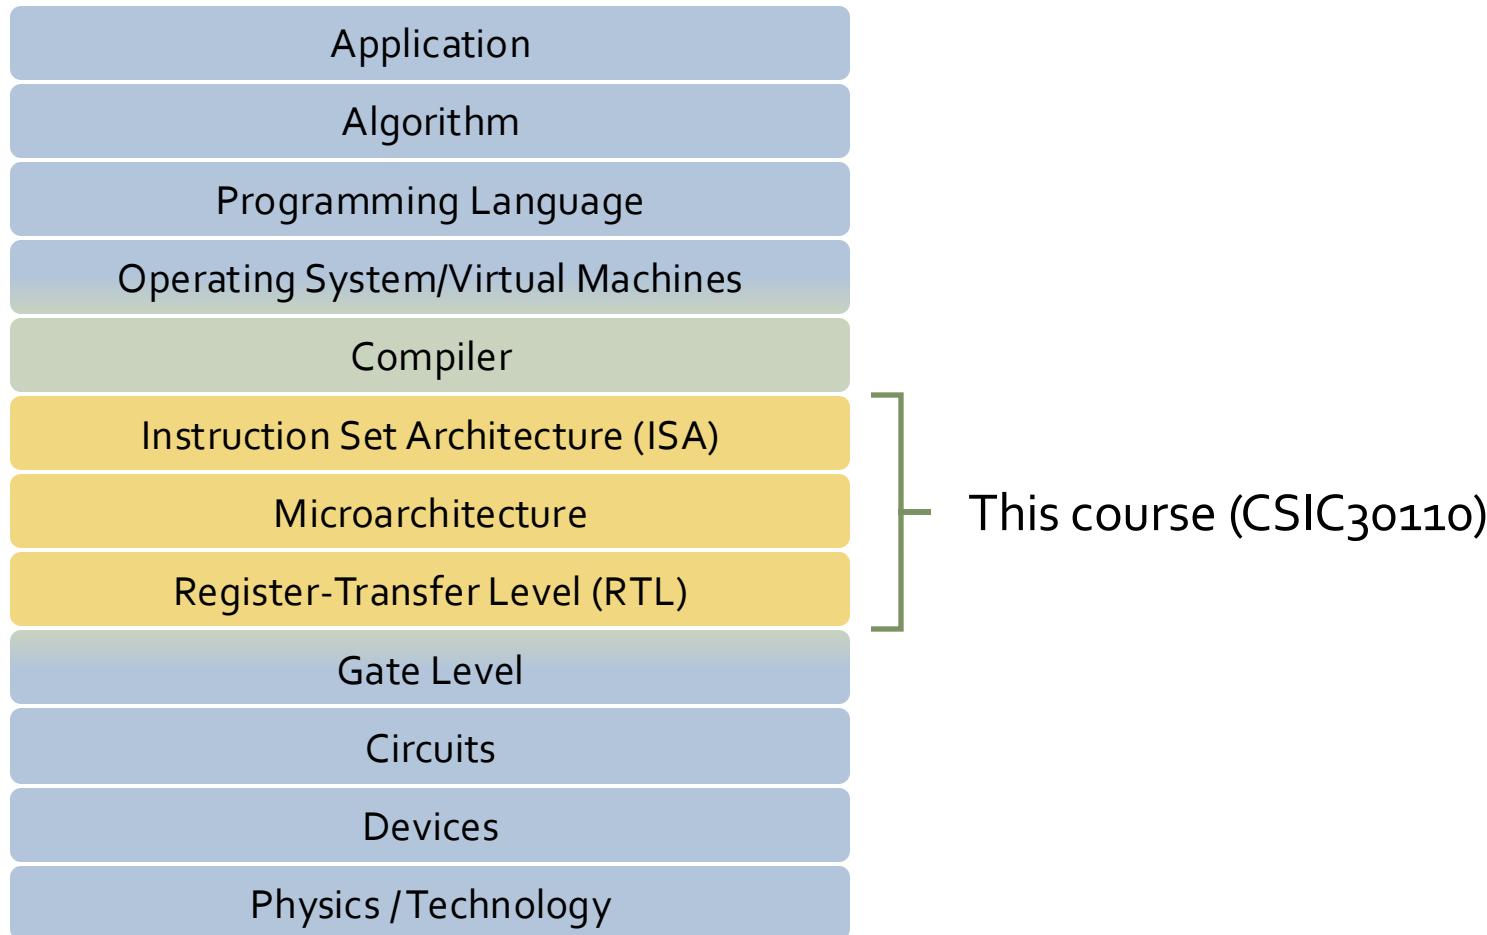


The Computer Systems Stack

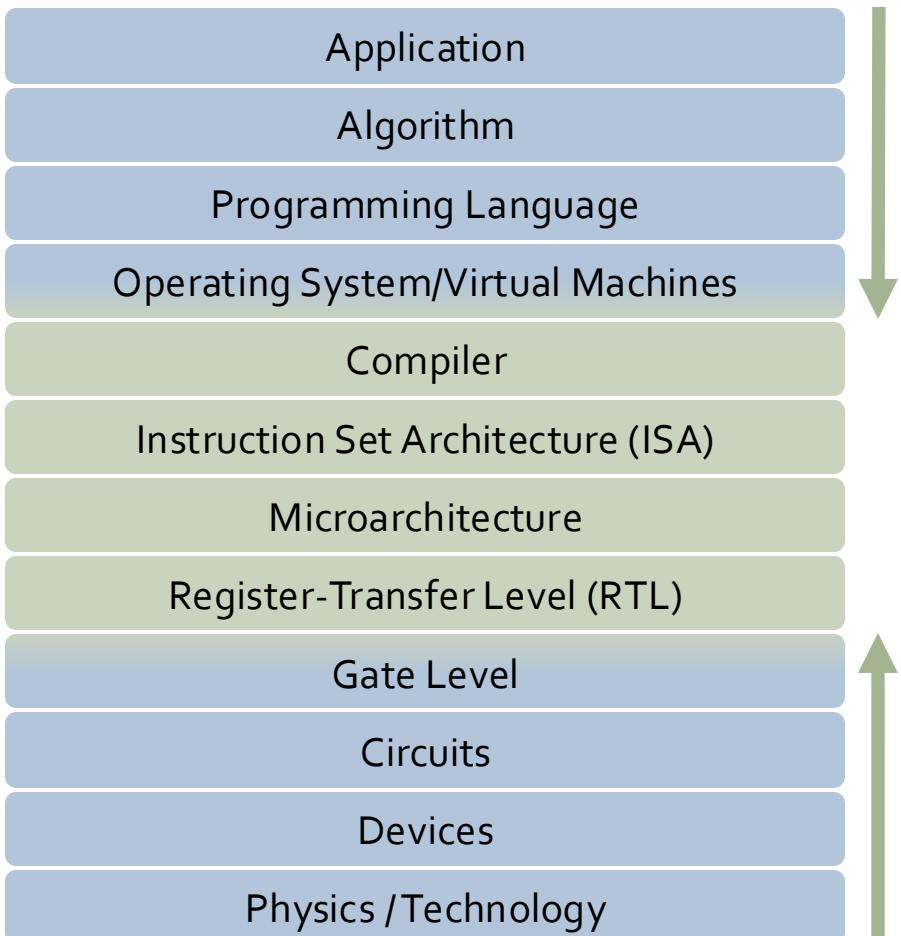
Computer Architecture



The Computer Systems Stack



Architecture is Constantly Changing



Application Requirements:

- Suggest how to improve architecture
- Provide revenue to fund development

Architecture provides feedback to guide application and technology research directions

Technology Constraints:

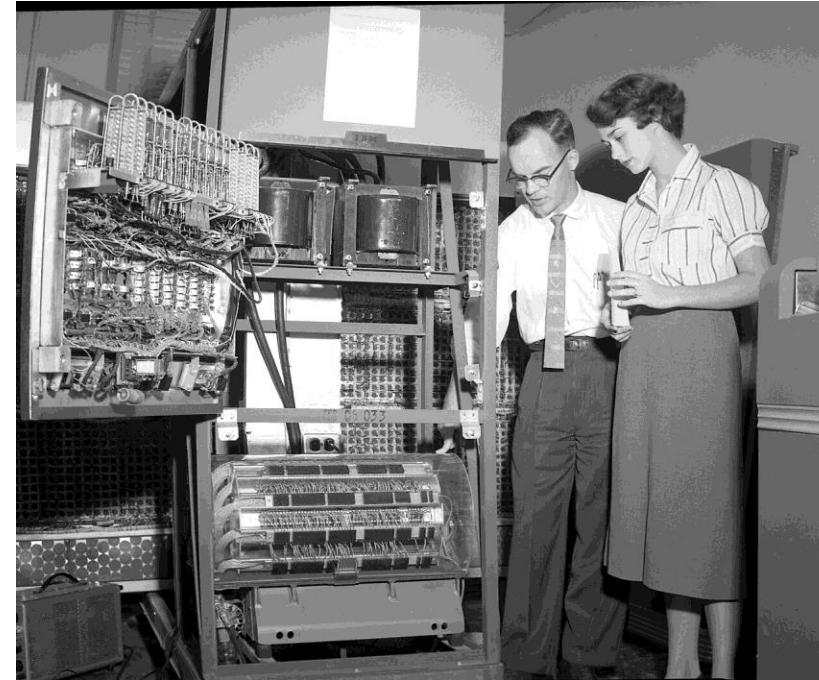
- Restrict what can be done efficiently
- New technologies make new arch possible

Computers Then...



IBM 650, 1962, NCTU

- The first mass-produced computer
- Almost 2,000 produced

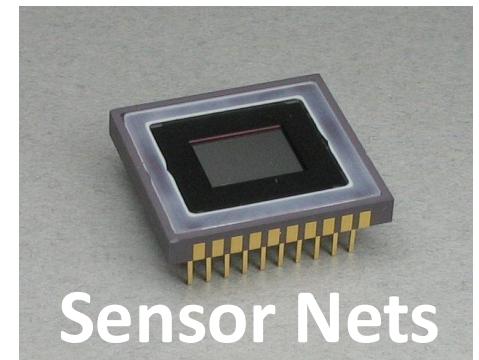


[Cushing Memorial Library and Archives, Texas A&M,
Creative Commons Attribution 2.0 Generic]

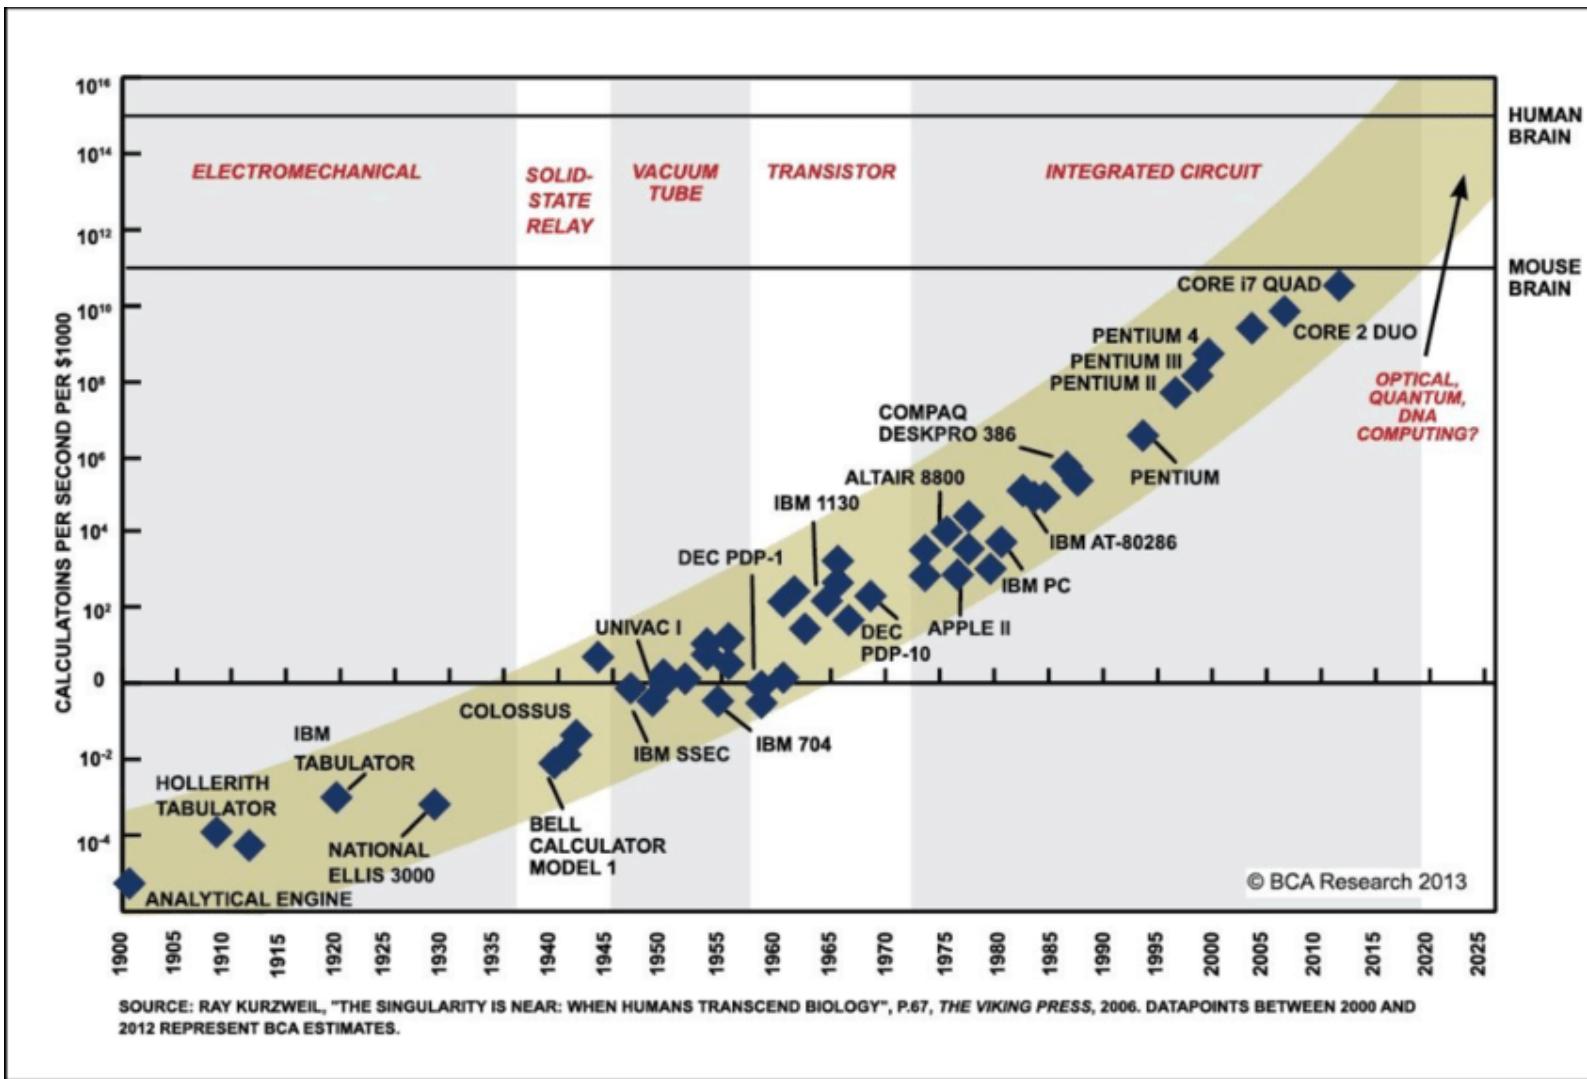
Computers Then...



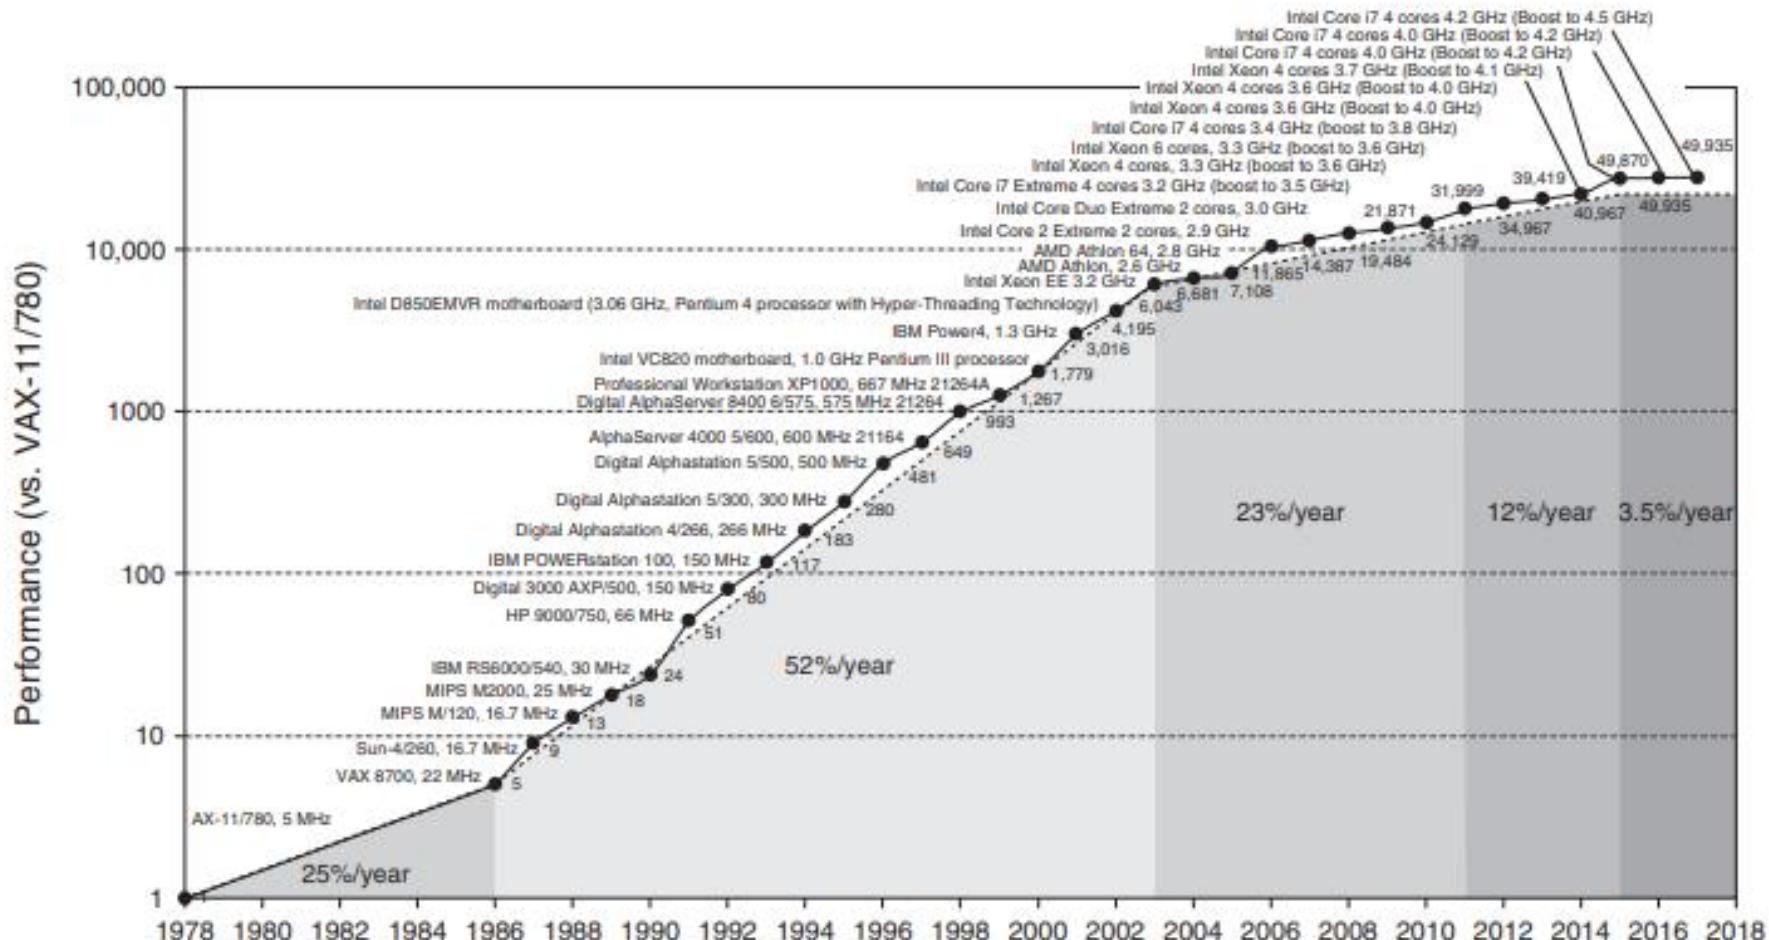
Computers Now...



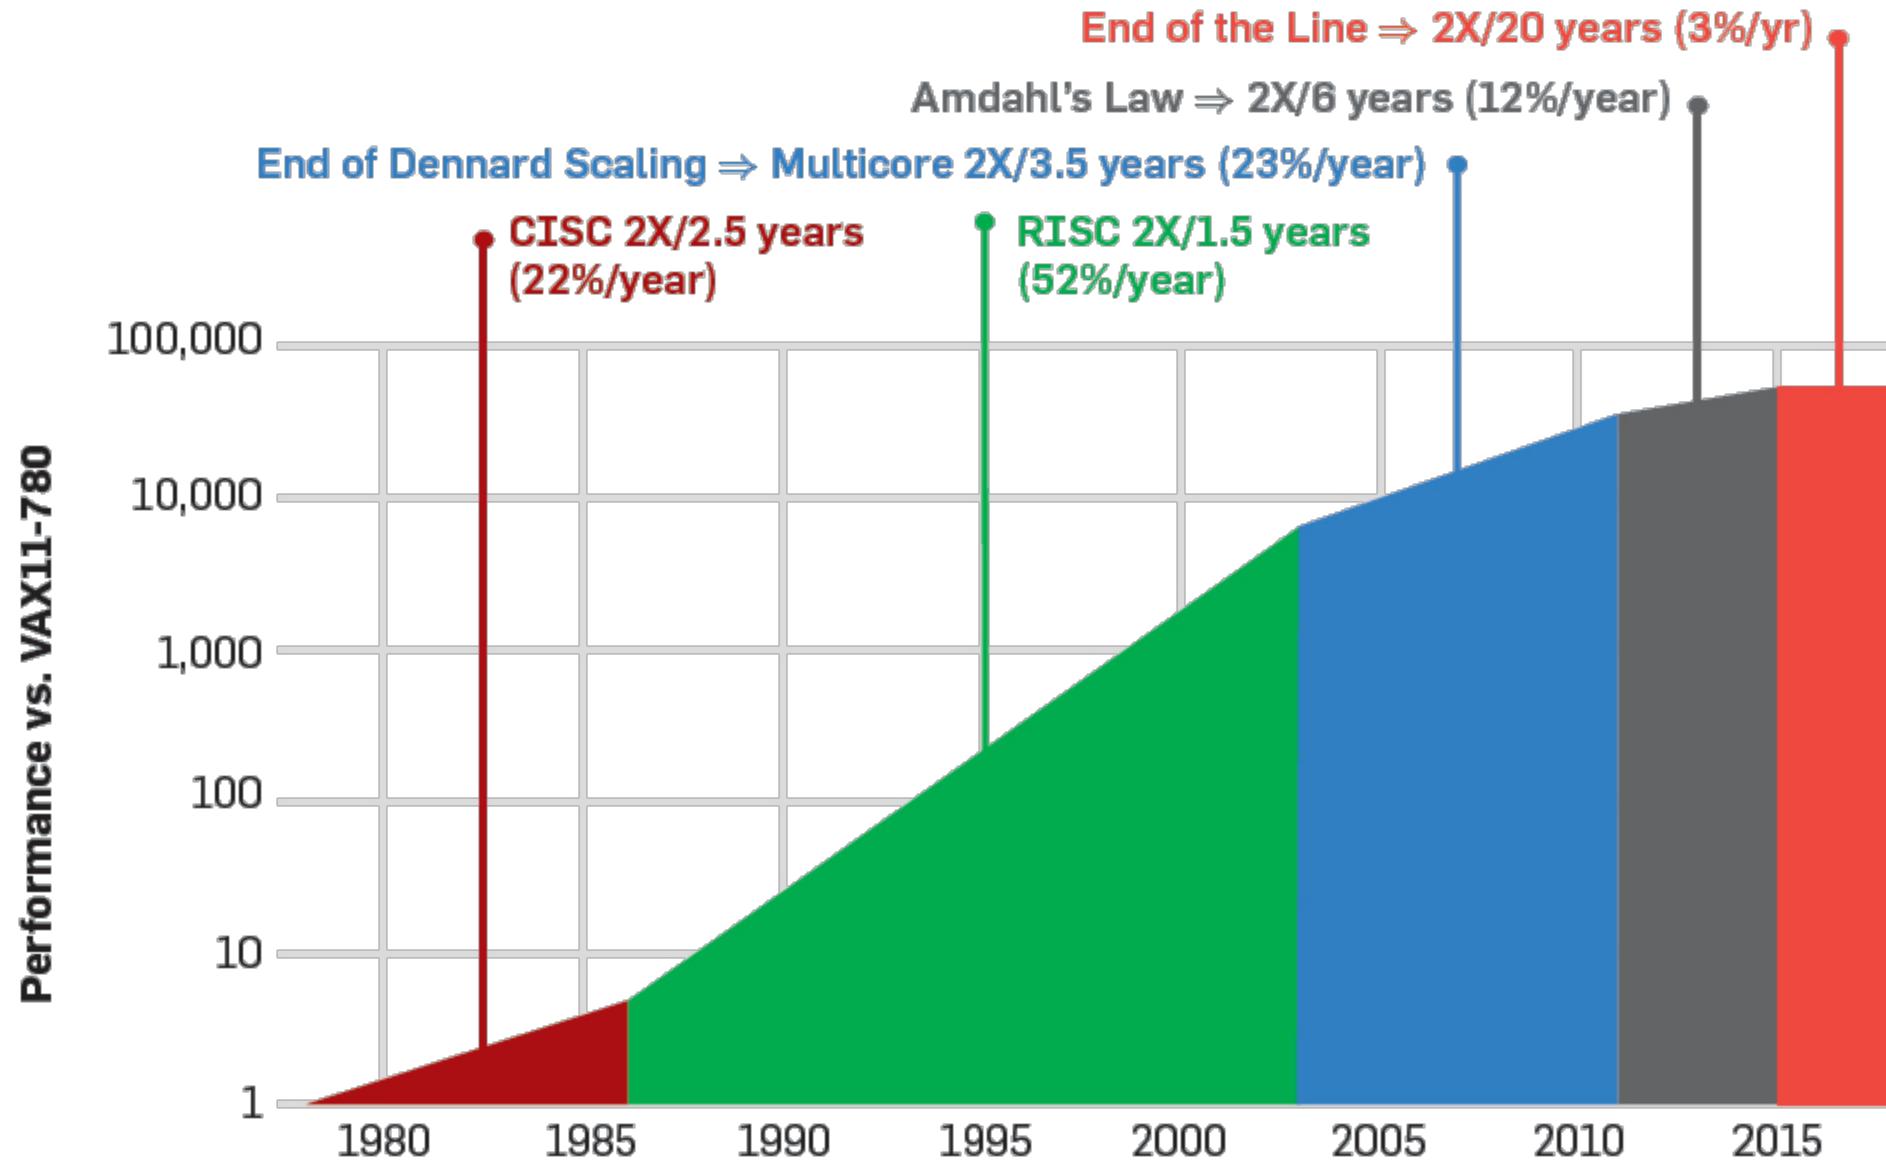
Moore's Law



Sequential Processor Performance

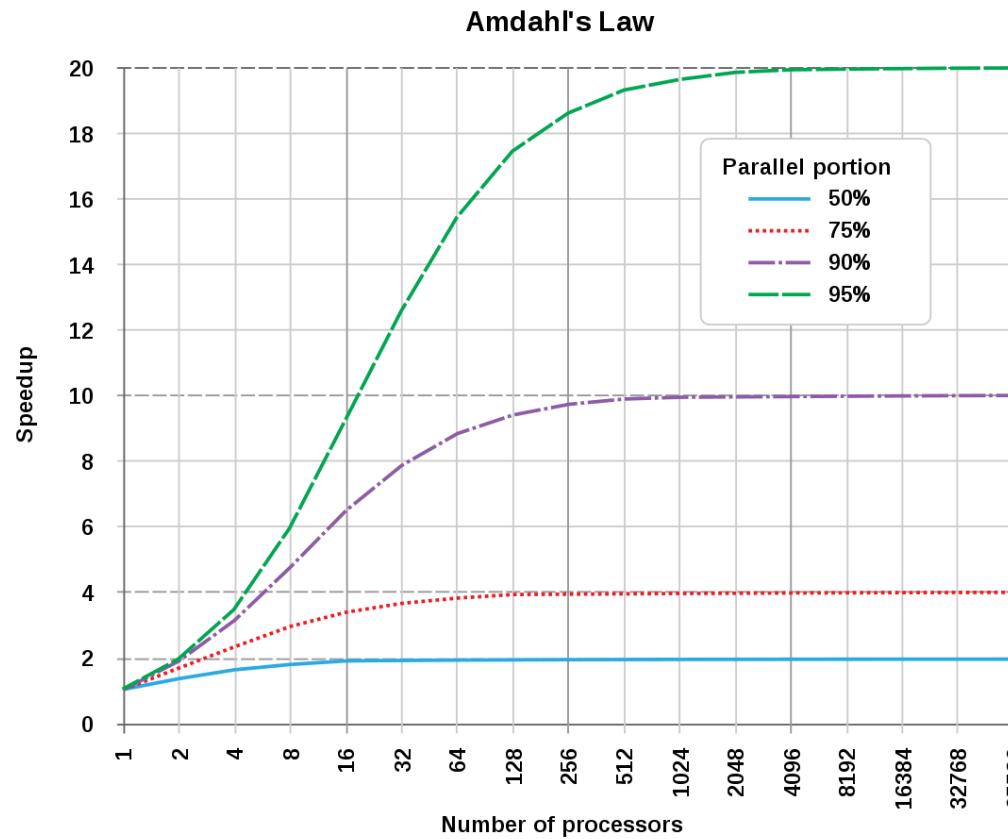


From Hennessy and Patterson Ed. 6 Image Copyright © 2019, Elsevier Inc. All rights reserved.



Amdahl's Law

- Speedup =
$$\frac{\text{Execution time for entire task without using the enhancement}}{\text{Execution time for entire task using the enhancement when possible}}$$



Upheaval in Computer Design

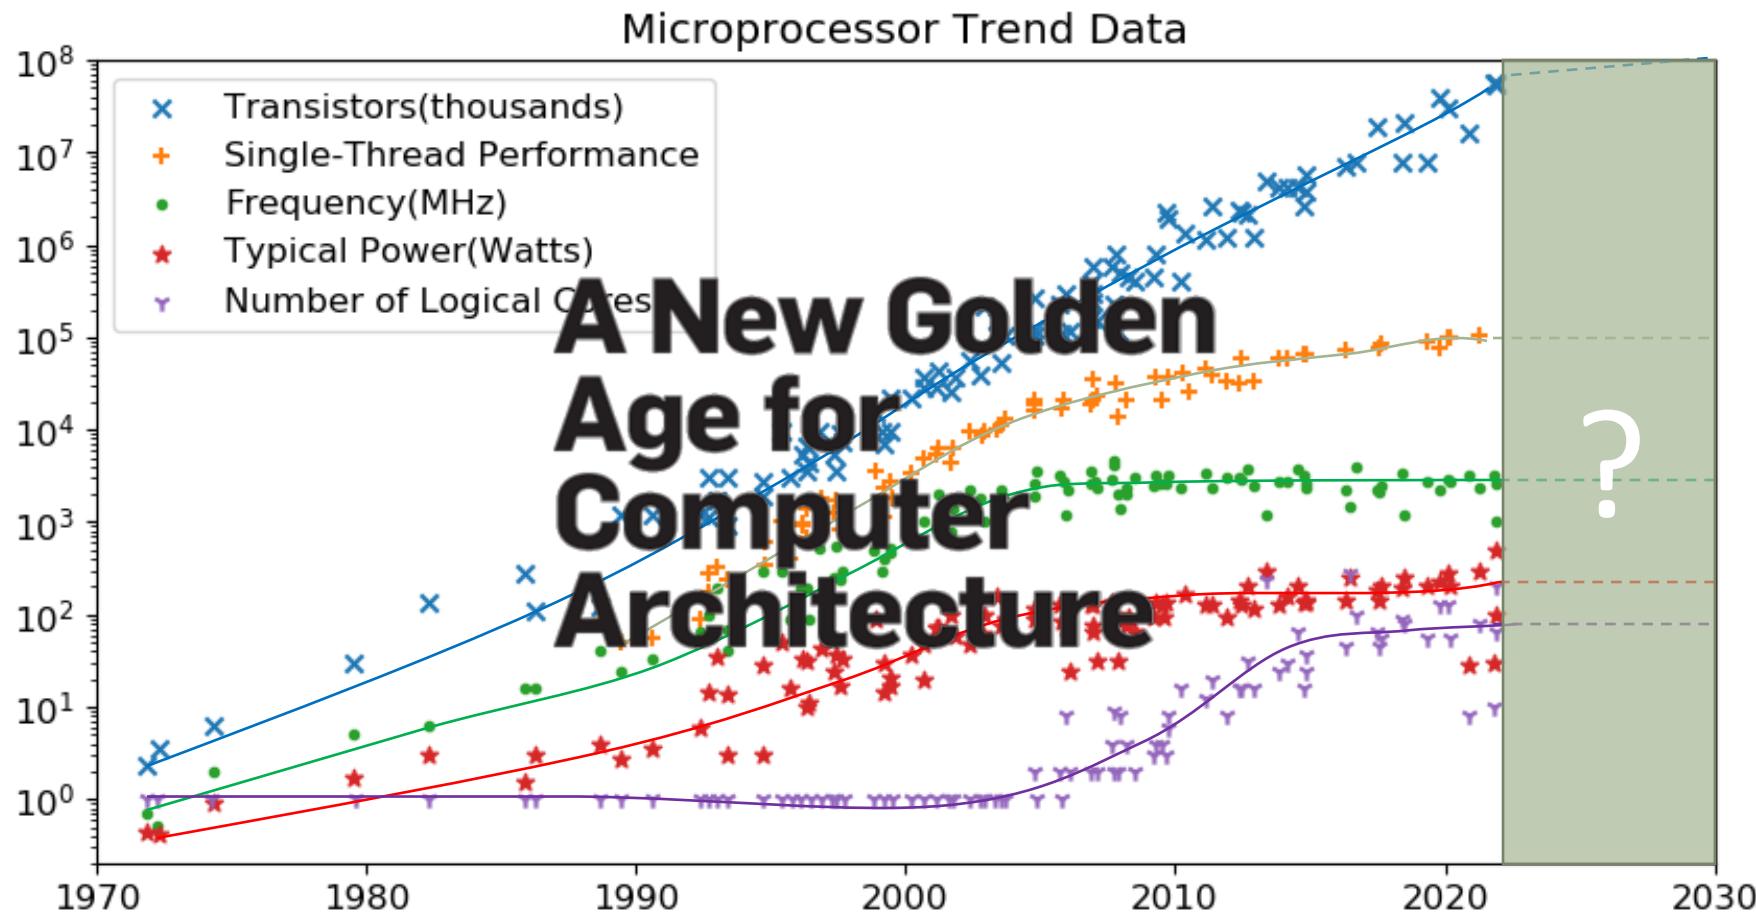
- Most of last 50 years, Moore's Law ruled
 - Scaling improved performance/energy without changing software model
- Last decade, technology scaling slowed/stopped
 - Dennard scaling is over (supply voltage ~fixed)
 - Moore's Law (cost/transistor) over?
 - No competitive replacement for CMOS in anytime soon
 - Energy efficiency is the main limiter
- 2020s shift
 - AI/ML drives compute demand
- No “free lunch” for software developers, must consider
 - Parallelism and heterogeneity are mandatory

Today's Dominant Target Systems

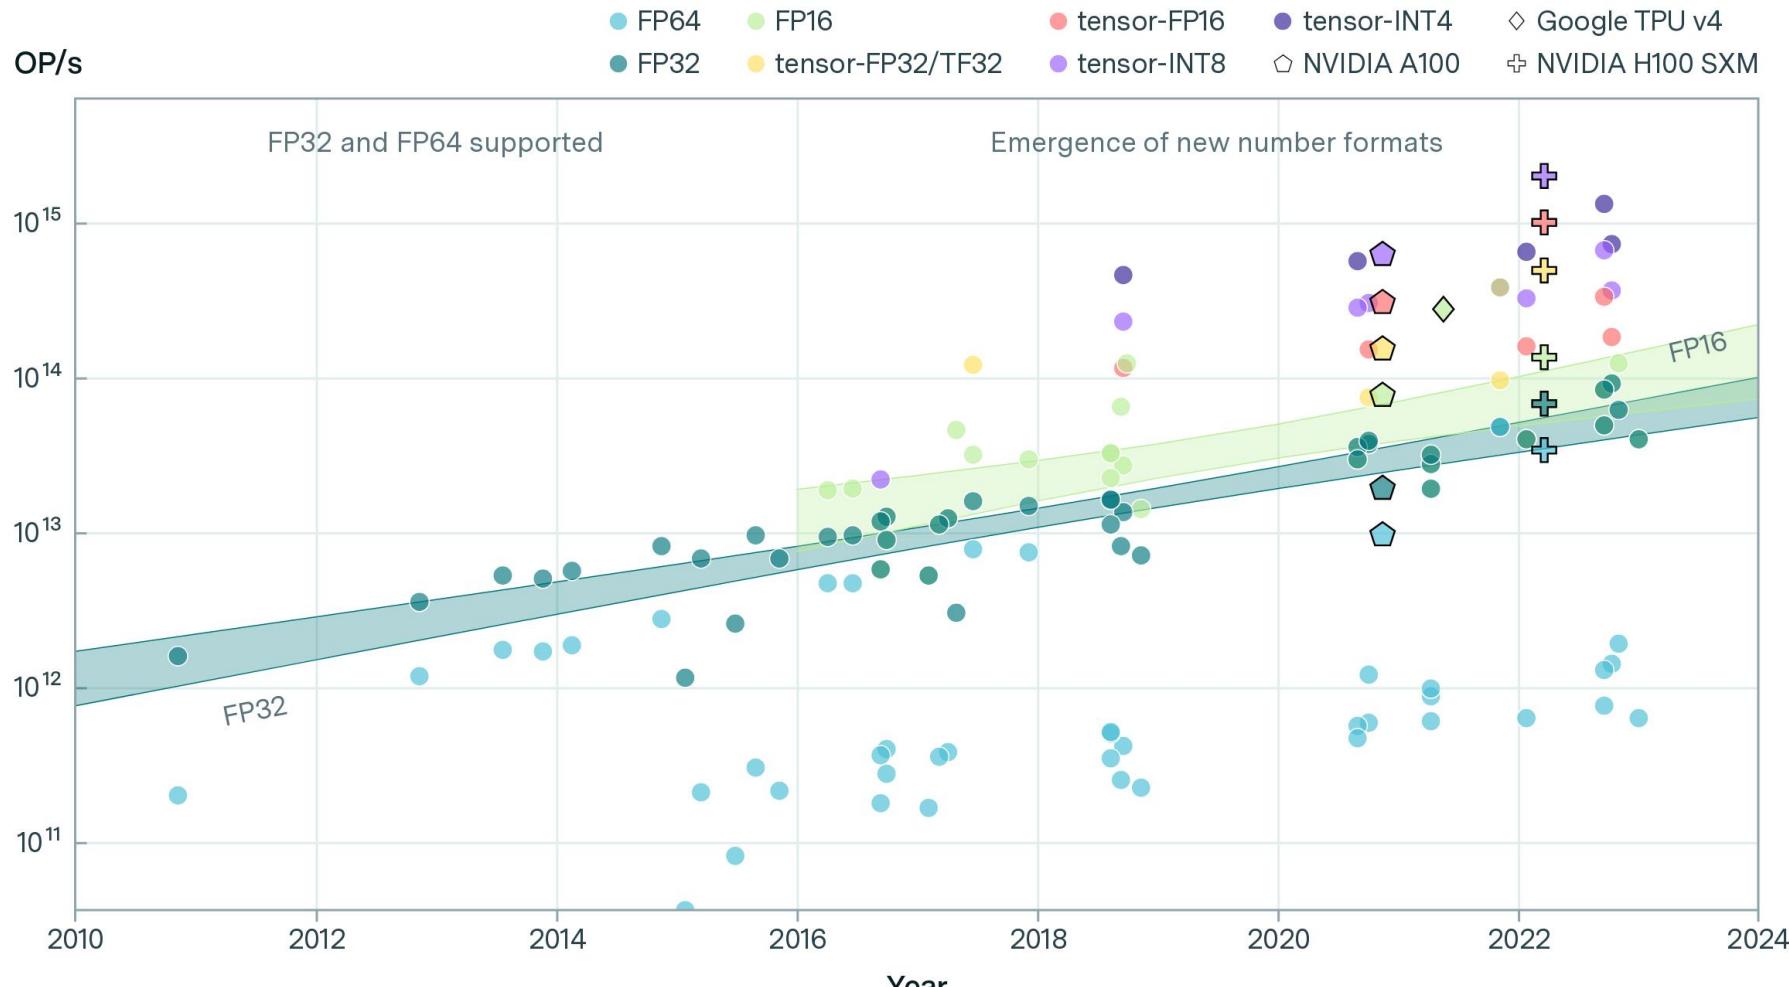
- Mobile (smartphone/tablet)
 - >1 billion sold/year
 - Dominated by ARM ISA in SoCs
 - Ship with AI/Neural engines + accelerators (vision, audio, security, sensors)
- Warehouse-Scale Computers (WSCs)
 - 100k+ cores per warehouse, cloud datacenters
 - Dominated by x86 ISA (server CPUs) + custom accelerators
 - Energy & carbon footprint now a key bottleneck
- Embedded and Edge
 - Consumer electronics, automotive, IoT
 - Strong RISC-V growth in microcontrollers & AI edge chips
 - Edge AI (TinyML, on-device generative AI) expanding rapidly

Beyond Moore's Law?

Number of accelerators?
Parallelization and specialization?



Trends in Machine Learning Hardware?



[Image Credit: Epoch AI]

The Verticalization of Silicon

- Shift: from buying chips → to building differentiated chips
 - Apple: ditched x86, built M-series SoCs for Mac/iPad
 - Google, Amazon, Microsoft, Alibaba: datacenter CPUs & AI accelerators
 - Tesla: custom chips for autonomous driving & EVs
 - OpenAI: in-house AI accelerator to reduce Nvidia dependence?
- End-system value/profit justifies cost of chip design
 - can be >>\$100M engineering cost to develop a new advanced chip!

Big Tech's Homegrown Chips

Company	Chip	Launched
Amazon	Graviton	2018
Google	Axion	2024
Microsoft	Cobalt	2023
Amazon	Trainium	2022
Amazon	Inferentia	2019
Google	TPU	2015/2017
Microsoft	MAIA	2023
Meta	MTIA	2023

Agenda

- What is Computer Architecture?
- Course Admin
- ISA Review

Course Administration

- Instructor: Prof. Ting-Jung Chang (tingchang@cs.nycu.edu.tw)
 - Office Hours: T 15:30-16:30 @ EC707 by appointment
- Lectures: Tuesday 13:20 – 15:10 @ EDB27
- Text: Computer Architecture: A Quantitative Approach **6th Edition**
- Prerequisite: Computer Organization, or equivalent
- Course Webpage: e3

TAs

- 趙家逸 (chaiyi.cs14@nycu.edu.tw)
- 謝宥逸 (yuyi92025.cs14@nycu.edu.tw)
- 吳柏頡 (pchw.cs14@nycu.edu.tw)
- 邵品翔 (amkingo916.cs14@nycu.edu.tw)
- Location: EC118
- Please email in advance
- Check e3 for final updates (office hours and other info)

Course Structure

- Midterm (20%) (~10/21)
- Final (30%) (~12/16)
- Labs (50%)
 - 5 design labs (Verilog)
- Ungraded Problem Sets (0%)
 - Intended to help you learn the material
 - Feel free to discuss with other students and instructors
 - Useful for exam preparation

Lab Academic Integrity

- Do your own work
- Be careful on AI usage
 - Explore ideas or debug with AI
 - But you must check, adapt, and make it your own
 - Work that lacks originality — AI or otherwise — will be penalized
- No public sharing of labs (GitHub, Google Drive, etc.)
- Integrity matters: violations can result in penalties, including failing the course

Course Administration

- Your goal today: decide if you're coming back
- Notices (all materials available on e3)
 - Labo released
 - Verilog tips as supplementary material

Computer Organization

- RISC-I (1982) fabbed in 5 μm NMOS, with a die area of 77 mm^2 , ran at 1 MHz. This chip is probably the first VLSI RISC
- Basic pipelined processor
- ~50,000 transistors

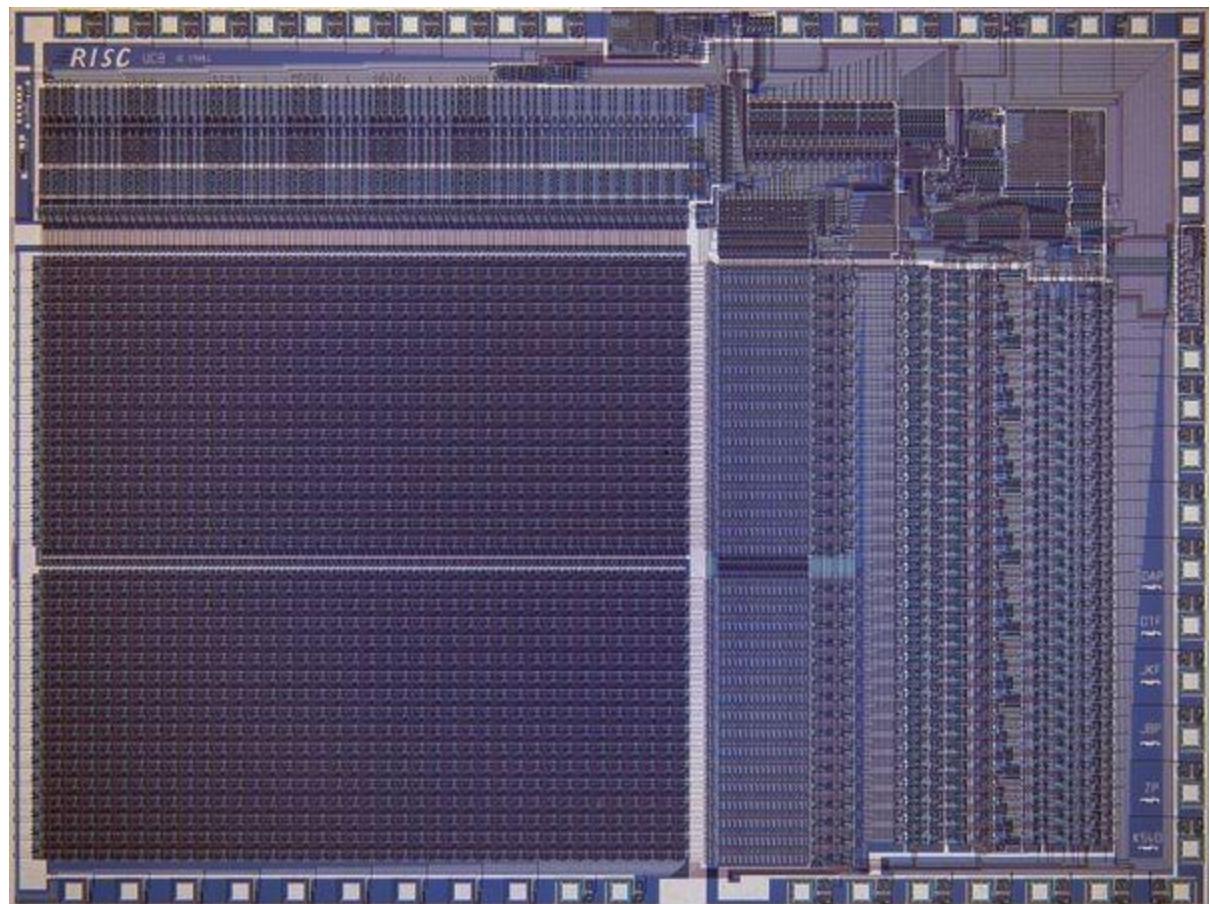


Photo of Berkeley RISC I, © University of California (Berkeley)

Computer Architecture

~731,000,000 transistors

- Instruction Level Parallelism
 - Superscalar
 - Very Long Instruction Word (VLIW)
- Long Pipelines (Pipeline Parallelism)
- Advanced Memory and Caches
- Data Level Parallelism
 - Vector
 - GPU
- Thread Level Parallelism
 - Multithreading
 - Multiprocessor
 - Multicore
 - Manycore
- More...

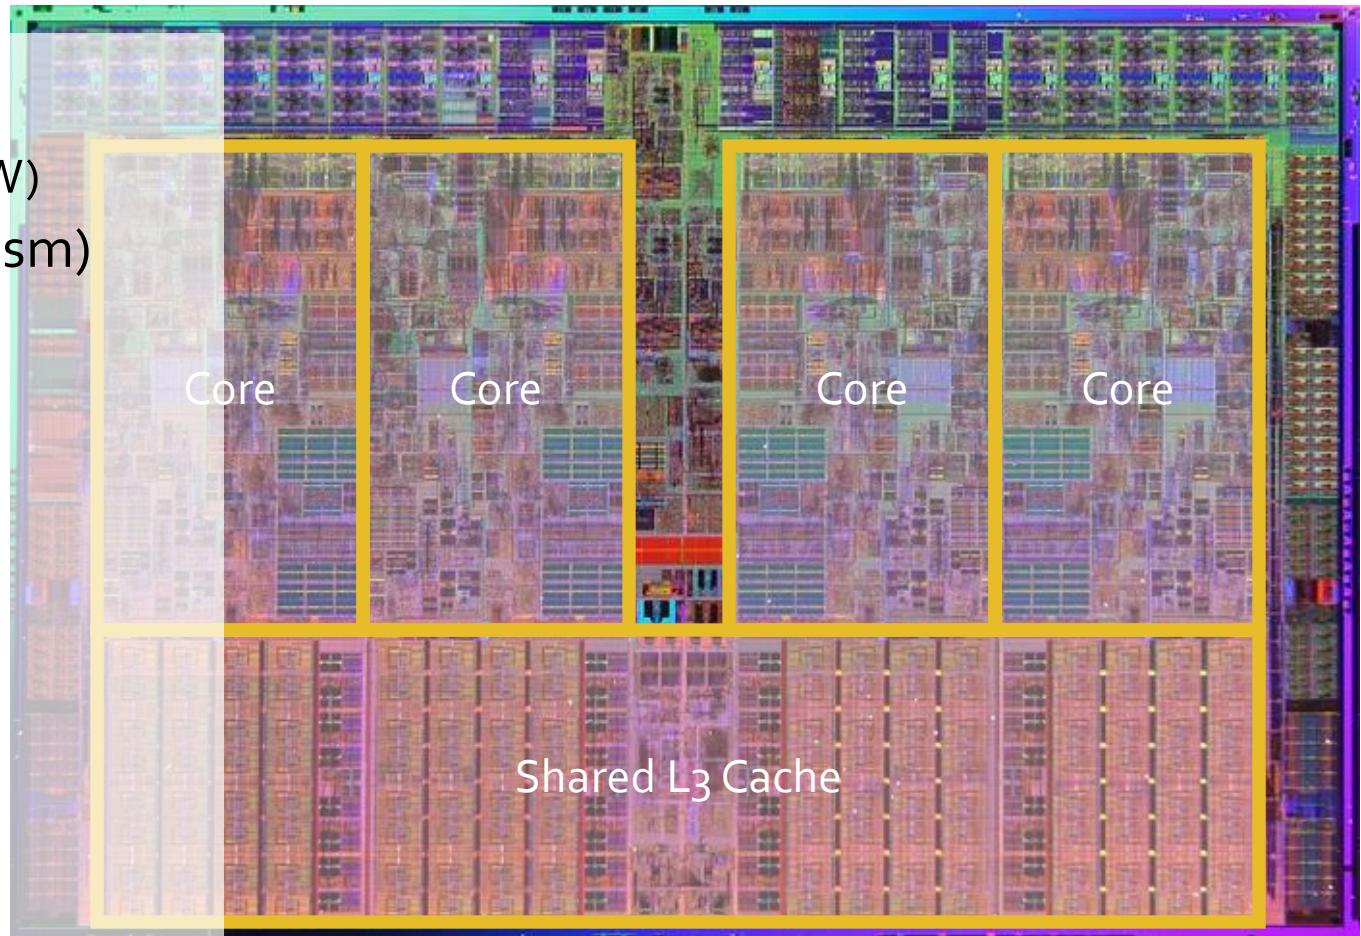


Photo of Intel Nehalem Processor, Original Core i7, © Intel 30

Agenda

- What is Computer Architecture?
- Course Admin
- ISA Review

Architecture vs. Microarchitecture

- Architecture/Instruction Set Architecture:
 - Programmer visible state (memory and register)
 - Operations (Instructions and how they work)
 - Execution Semantics (Interrupts)
 - Input/Output
 - Data Types/Sizes
- Microarchitecture/Organization:
 - Tradeoffs on how to implement ISA for some metric (speed, energy, cost)
 - Examples: pipeline depth, number of pipelines, cache size, silicon area, peak power, execution ordering, bus widths, ALU widths

Software Developments

up to 1955	Libraries of numerical routines <ul style="list-style-type: none">• Floating point operations• Transcendental functions• Matrix manipulation, equation solvers
1955-60	High level languages – Fortran 1956 Operating Systems <ul style="list-style-type: none">• Assemblers, loaders, linkers, compilers• Accounting programs to keep track of usage and charges

Machines required **experienced** operators

- Most users could not be expected to understand these programs, much less write them
- Machines had to be sold with a lot of resident software

Compatibility Problem at IBM

- By early 1960's, IBM had 4 incompatible lines of computers!
 - 701 → 7094
 - 650 → 7074
 - 702 → 7080
 - 1401 → 7010
- Each system had its own
 - Instruction set
 - I/O system and secondary storage (magnetic tapes, drums and disks)
 - Assemblers, compilers, libraries
 - Market niche business, scientific, real time, ...

IBM 360!

IBM 360: A General-Purpose Register (GPR) Machine

- Processor State
 - 16 general-purpose 32-bit registers
 - May be used as index and base register
 - Register 0 has some special properties
 - 4 floating point 64-bit registers
 - A program status word (PSW)
 - PC, condition codes, control flags
- A 32-bit machine with 24-bit addresses
 - But no instruction contains a 24-bit address!
- Data formats
 - 8-bit bytes, 16-bit half words, 32-bit words, 64-bit double words

The IBM 360 is why bytes are 8-bits long today!

IBM 360: Initial Implementations

	Model 30	Model 70
Storage	8K – 64 KB	256K – 512 KB
Datapath	8-bit	64-bit
Circuit Delay	30 nsec/level	5 nsec/level
Local Store	Main Store	Transistor Registers
Control Store	Read only 1μsec	Conventional circuits

- IBM 360 instruction set architecture (ISA) completely hid the underlying technological differences between various models.
- Milestone: The first true ISA designed as portable hardware-software interface!

With minor modifications it still survives today!

IBM 360: Over 50 years later...

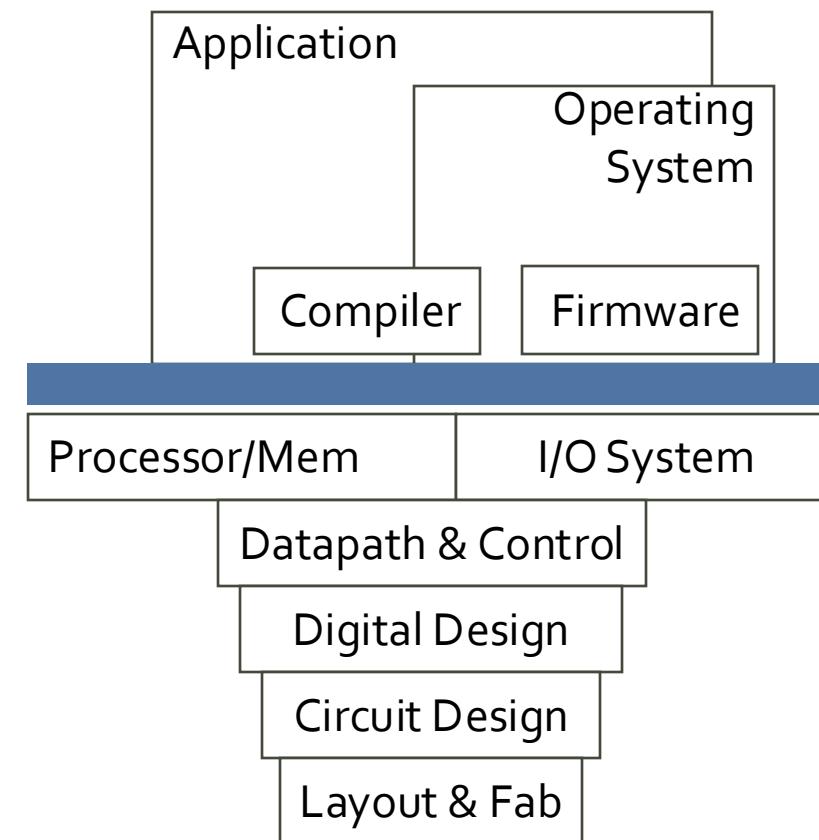
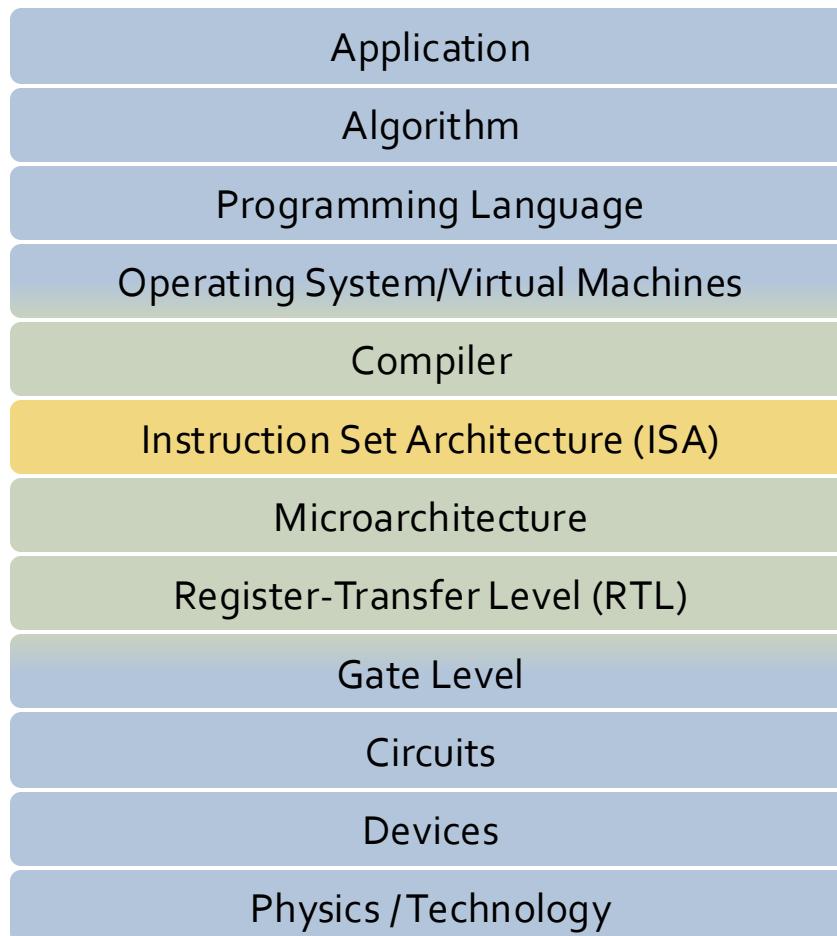
The zSeries z16 Microprocessor



Image Credit: IBM
© International Business
Machines Corporation.

- 5.2GHz in 7nm Samsung technology
- 22.5 Billion transistors in 530 mm²
- 64-bit virtual addressing – original S/360 was 24-bit, and S/370 was 31-bit extension
- 8 cores + L2s per chip
- 128 KB L1, 32 MB L2, 256MB L3, 2048 MB L4
- Integrated on-chip AI accelerator

Instruction Set Architecture (ISA)

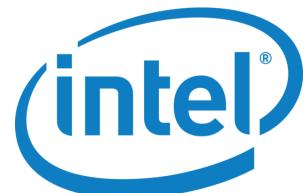


Instruction Set Architecture (ISA)

- The ISA is functional abstraction of the processor (a “mental model”)
 - What operations can be performed
 - How to name storage locations
 - The format (bit pattern) of the instructions
- ISA typically does NOT define
 - Timing of the operations
 - Power used by operations
 - How operations/storage are implemented
- Many implementations possible for a given ISA

Different Instruction Set Architecture

- ARM
 - Family of ISAs developed by ARM
 - Widely used in mobile and low-power devices
 - Now expanding into desktop, datacenters and cloud servers
- x86
 - Family of ISAs developed by Intel (and AMD)
 - Used in general-purpose computing systems (desktops and servers)
- RISC-V
 - Open standard ISAs developed at UC-Berkeley
 - Mostly used in embedded systems
 - Early adoption in AI accelerators and research chips

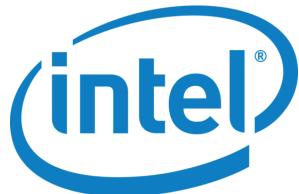


Different Instruction Set Architecture

- ARM
 - Family of ISAs developed by Arm Holdings
 - Widely used in mobile devices
 - Now expanding into desktop and server markets
- x86
 - Family of ISAs developed by Intel
 - Used in general-purpose computing
- RISC-V
 - Open standard ISAs developed by RISC-V International
 - Mostly used in embedded systems
 - Early adoption in AI applications

The collage consists of four news snippets arranged vertically:

- Arm gives up on killing off Qualcomm's vital chip license** (The Register, 2 weeks ago)
Arm has given up on terminating one of its key licenses with Qualcomm, leaving the latter free to continue producing homegrown Arm-compatible chips.
- Qualcomm's legal win could reshape Arm licensing landscape** (DIGITIMES Asia, Dec 25, 2024)
Following Qualcomm's recent victory in its patent dispute with Arm, industry analysts are examining the broader implications for...
- Qualcomm says Arm has withdrawn license breach notice** (Reuters, 2 weeks ago)
Qualcomm Chief Executive Officer Cristiano Amon said on Wednesday that Arm Holdings has withdrawn a threat to terminate Qualcomm's license...
- Qualcomm wins a legal battle over Arm chip licensing** (The Verge, Dec 20, 2024)
Qualcomm wins a legal battle over Arm chip licensing. The jury sided with Qualcomm after Arm argued the chipmaker breached a licensing agreement...
- Arm and Qualcomm's licensing dispute could nudge market dynamics and future competition** (Canalys, Oct 30, 2024)
Recent news about Arm notifying Qualcomm of its intention to terminate its technology licensing agreement has raised considerable concerns...



servers



Same Architecture/Different Microarchitecture

Intel Xeon

- x86 instruction set
- 32+ cores
- 200W+
- Decode 6 instructions/cycle/core
- Large shared caches (tens of MB L3)
- Out-of-order
- ~ 2.0-3.5 GHz

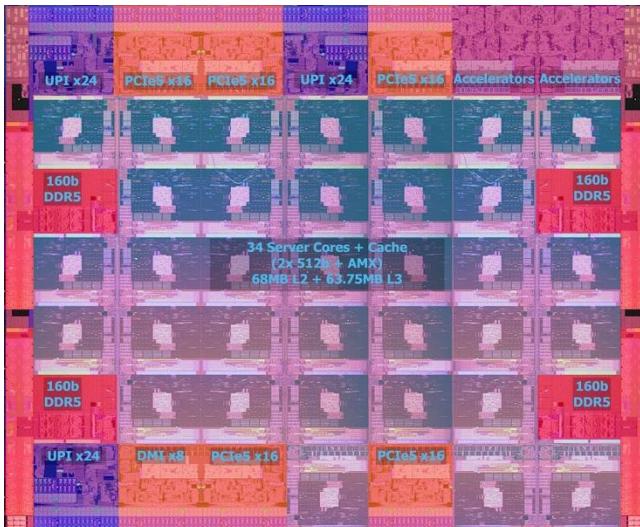


Image Credit: Intel

Intel Atom

- x86 instruction set
- Few cores (1-8)
- 2W
- Decode 2 instructions/cycle/core
- tens of KB L1, <2 MB L2
- In-order
- 1.6GHz

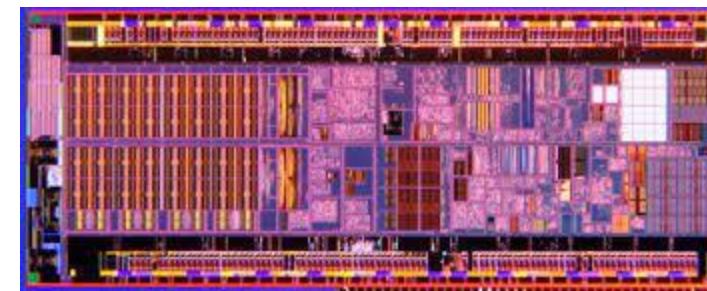


Image Credit: Intel

Same Architecture/Different Microarchitecture

Intel Xeon

- x86 instruction set
- 32+ cores
- 200W+
- Decode 6 instructions/cycle/core
- Large shared caches (tens of MB L3)
- Out-of-order
- ~ 2.0-3.5 GHz

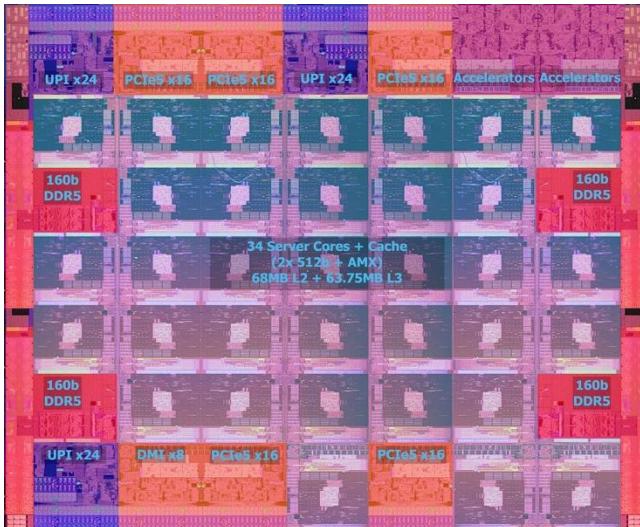


Image Credit: Intel

Apple M2 Ultra

- ARM instruction set
- Up to 24 cores
- ~90–100 W
- Decode 8+ instructions/cycle/core
- Large shared caches across clusters
- Out-of-order
- ~3.5 GHz peak frequency

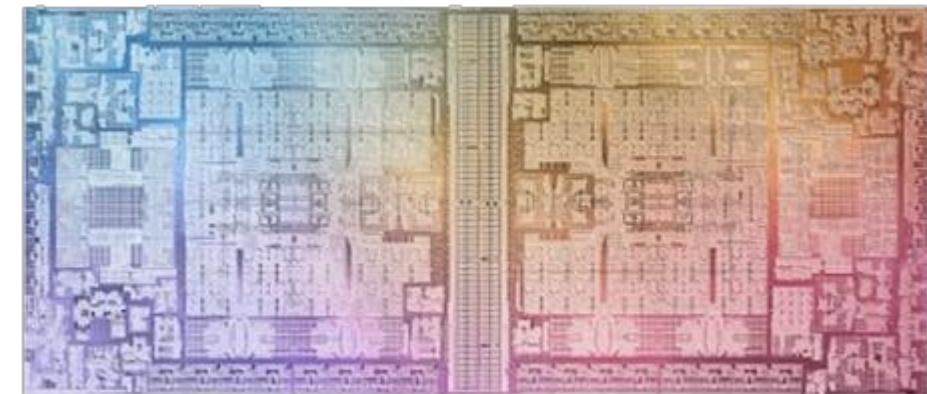


Image Credit: Apple

The Standard Structure of An Instruction

- An instruction typically has an operator (op or opcode), one or two source operands (src), and one destination operand (dest).

R1	=	R2	+	R3
----	---	----	---	----

goto	target
------	--------

R1	=	*addr
----	---	-------

if	(R1 > *addr)
----	--------------

*addr	=	R1
-------	---	----

if	(R1 == 0)	{goto	target}
----	-----------	-------	---------

What would these look like in RISC-V assembly?

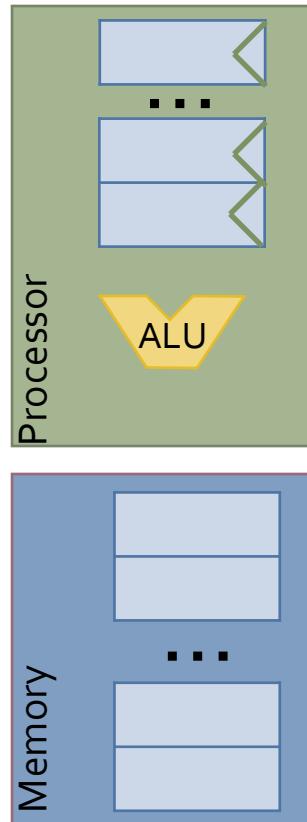
Key ISA Decisions

- Operands
 - How many?
 - Location?
 - Addressing mode?
 - Types?
- Operations
 - What kind?
 - How many?
- Instruction format
 - Length(s) of bit pattern
 - Which bits mean what

ISA Classification

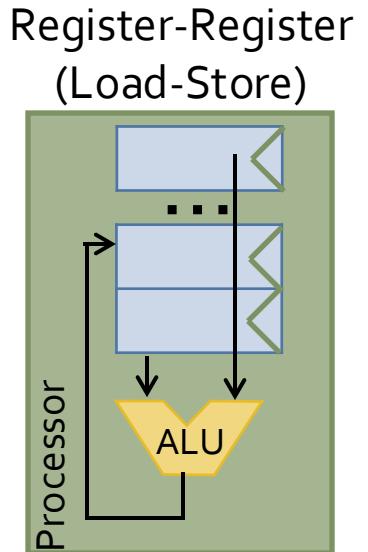
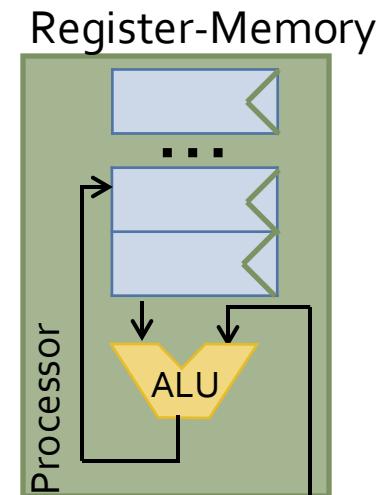
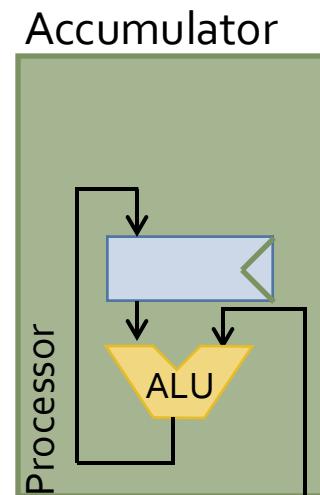
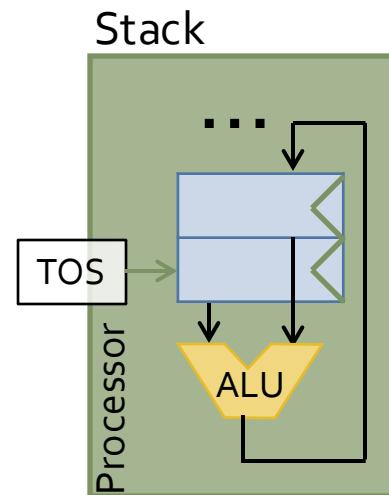
Operands

- Operands may be from
 - Stack
 - Accumulator
 - Register
 - Register and Memory



Where do operands come from
and where do results go?

Machine Models of ISA



Number Explicitly
Named Operands:

0

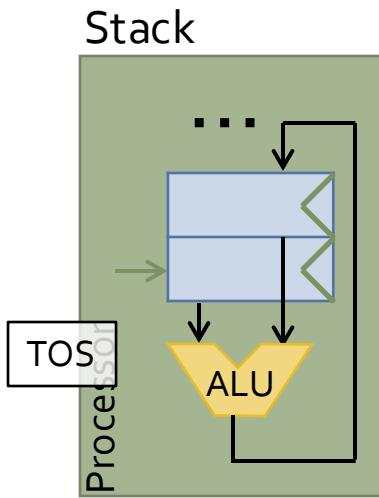
1

2 or 3

2 or 3

Summary: Machine Model

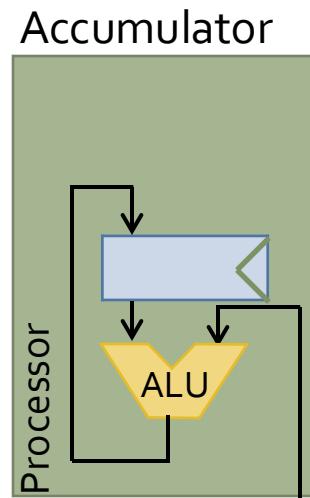
Stack



$C = A + B$

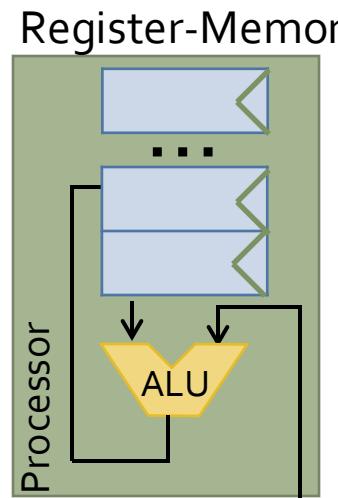
Push A
Push B
Add
Pop C

Accumulator



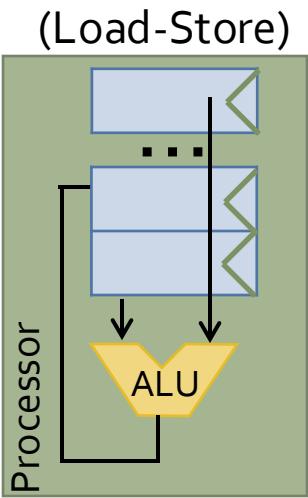
Load A
Add [B]
Store C

Register-Memory
(Load-Store)



Load R₁, [A]
Add R₃, R₁, [B]
Store R₃, [C]

Register-Register
(Load-Store)



Load R₁, [A]
Load R₂, [B]
Add R₃, R₁, R₂
Store R₃, [C]

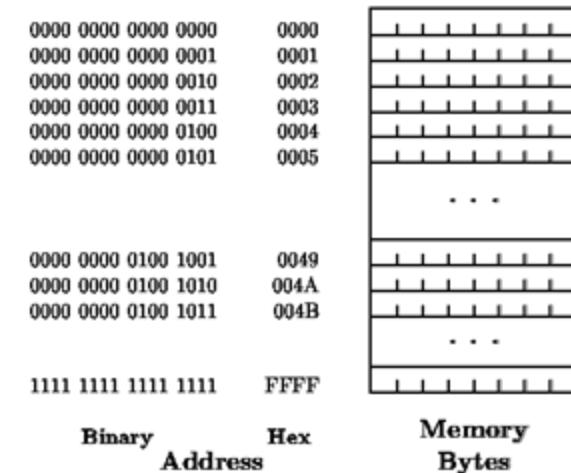
ISA Classification

Operand Addressing Mode

- Addressing mode specifies how operands are located from memory and/or registers.
- Common addressing modes
 - Register
 - Immediate
 - Register indirect
 - Displacement
 - Indexed
 - Direct
 - Memory indirect
 - Auto-increment and auto-decrement
 - Scaled

Memory Addressing

- Objects have byte addresses
 - the number of bytes counted from the beginning of memory
- Object Length:
 - bytes (8 bits), half words (16 bits)
 - words (32 bits), and double words (64 bits).
 - The type is implied in opcode, e.g.,
 - `lb` – load byte
 - `lw` – load word, etc.
- Addressing mode specifies how operands are located from memory and/or registers



Memory Addressing

- Byte Ordering
 - Little Endian: Stores the least significant byte (LSB) at the lowest memory address
 - Ex: ox12345678
 - Big Endian: Stores the most significant byte (MSB) at the lowest memory address
 - Ex: ox12345678
 - Problem occurs when exchanging data among machines with different orderings

Address	0x100	0x101	0x102	0x103
	0x78	0x56	0x34	0x12

Address	0x100	0x101	0x102	0x103
	0x12	0x34	0x56	0x78

Summary: Addressing Modes

Addressing Mode	Instruction	Function
Register	Add R ₄ , R ₃ , R ₂	Regs[R ₄] <- Regs[R ₃] + Regs[R ₂] **
Immediate	Add R ₄ , R ₃ , #5	Regs[R ₄] <- Regs[R ₃] + 5 **
Displacement	Add R ₄ , R ₃ , 100(R ₁)	Regs[R ₄] <- Regs[R ₃] + Mem[100 + Regs[R ₁]]
Register Indirect	Add R ₄ , R ₃ , (R ₁)	Regs[R ₄] <- Regs[R ₃] + Mem[Regs[R ₁]]
Absolute/Direct	Add R ₄ , R ₃ , (0x475)	Regs[R ₄] <- Regs[R ₃] + Mem[0x475]
Memory Indirect	Add R ₄ , R ₃ , @(R ₁)	Regs[R ₄] <- Regs[R ₃] + Mem[Mem[Regs[R ₁]]]
PC relative	Add R ₄ , R ₃ , 100(PC)	Regs[R ₄] <- Regs[R ₃] + Mem[100 + PC]
Scaled	Add R ₄ , R ₃ , 100(R ₁)[R ₅]	Regs[R ₄] <- Regs[R ₃] + Mem[100 + Regs[R ₁] + Regs[R ₅] * 4]

** May not actually access memory!

Type And Size of Operands

- Types
 - Binary Integer
 - Binary Coded Decimal (BCD)
 - Floating Point
 - IEEE 754
 - Cray Floating Point
 - Intel Extended Precision (80-bit)
 - Packed Vector Data
 - Addresses
- Width
 - Binary Integer (8-bit, 16-bit, 32-bit, 64-bit)
 - Floating Point (32-bit, 40-bit, 64-bit, 80-bit)
 - Addresses (16-bit, 24-bit, 32-bit, 48-bit, 64-bit)

Key ISA Decisions

- Operands
 - How many?
 - Location?
 - Addressing mode?
 - Types?
- Operations
 - What kind?
 - How many?
- Instruction format
 - Length(s) of bit pattern
 - Which bits mean what

Classes of Instructions

- Arithmetic & Logical
 - Integer arithmetic: add, sub, mul, div, shift
 - Logical operation: and, or, xor, not
- Data Transfer
 - copy, load, store (w/ memory addressing)
- Control Flow
 - branch, jump, call, return, trap
- Floating Point
- System: OS and memory management
- Graphics: pixel and vertex, compression/decompression operations
- String
 - move, compare, search

Key ISA Decisions

- Operands
 - How many?
 - Location?
 - Addressing mode?
 - Types?
- Operations
 - What kind?
 - How many?
- Instruction format
 - Length(s) of bit pattern
 - Which bits mean what

Encoding an Instruction Set

- Opcode: specifying the operation
- # of operands:
 - addressing mode
 - address specifier tells what addressing mode is used
 - Load-store computer
 - Only one memory operand
 - Only one or two addressing modes
- The architecture must balance several competing forces when encoding the instruction set:
 - # of registers && addressing modes
 - Size of registers && addressing mode fields
 - Average instruction size && average program size.
 - Easy to handle in pipeline implementation

Example: x86 Instruction Encoding

Instruction Prefixes	Opcode	ModR/M	Scale, Index, Base	Displacement	Immediate
Up to four Prefixes (1 byte each)	1,2, or 3 bytes	1 byte (if needed)	1 byte (if needed)	0,1,2, or 4 bytes	0,1,2, or 4 bytes

Possible instructions 1 to 15 bytes long

MIPS64 Instruction Encoding

Basic instruction formats

R	opcode	rs	rt	rd	shamt	funct
	31	26 25	21 20	16 15	11 10	6 5 0
I	opcode	rs	rt		immediate	
	31	26 25	21 20	16 15		
J	opcode			address		
	31	26 25				

Floating-point instruction formats

FR	opcode	fmt	ft	fs	fd	funct
	31	26 25	21 20	16 15	11 10	6 5 0
FI	opcode	fmt	ft		immediate	
	31	26 25	21 20	16 15		

RISC-V Instruction Encoding

- Restrictions
 - 4 bytes per instruction
 - Different instructions have different parameters (registers, immediates, ...)
 - Various fields should be encoded to consistent locations
 - Simpler decoding circuitry
- RISC-V uses 6 “types” of instruction encoding

Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

ISA Classification

Instruction Length

- **Fixed Width:** Every Instruction has same width
 - + Easy to decode (**RISC** Architectures: MIPS, PowerPC, SPARC, ARM, RISC-V...)
 - Wasted bits in instructions (Why is this bad?)
 - Harder-to-extend ISA (how to add new instructions?)
 - Ex: MIPS, RISC-V, ARM, ...
- **Variable Length:** Instructions can vary in width
 - + Takes less space in memory and caches (**CISC** Architectures: x86, VAX...)
 - More logic to decode a single instruction
 - Harder to decode multiple instructions concurrently
 - Ex: x86, instructions 1-byte up to 18-bytes

ISA Classification

Instruction Length

- **Hybrid**
 - Support 16-bit and 32-bit instructions in RISC. Narrow instructions support fewer operations, smaller address and immediate fields, fewer registers, and two-address format rather than the classic three-address format
 - Claim a code size reduction of up to 40%
 - Ex: ARM Thumb, MIPS16e, and RISC-V
- **Compressed:**
 - PowerPC and some VLIWs (Store instructions compressed, decompress into Instruction Cache)
- **(Very) Long Instruction Word (VLIW):**
 - Multiple instructions in a fixed width bundle
 - Ex: Multiflow, HP/ST Lx, TI C6000

Instruction Length Tradeoffs

- Instructions are eventually encoded with 0s and 1s.
 - Each instruction is encoded into several bytes of binary numbers.
- Tradeoffs
 - Code size (memory space, bandwidth, latency) vs. hardware complexity
 - ISA extensibility and expressiveness
 - Performance? Smaller code vs. imperfect decode

Real World Instruction Sets

RISC? CISC?

Arch	Type	# Oper	# Mem	Data Size	# Regs	Addr Size	Use
Alpha	Reg-Reg	3	0	64-bit	32	64-bit	Workstation
ARM	Reg-Reg	3	0	32/64-bit	16	32/64-bit	Cell Phones, Embedded
MIPS	Reg-Reg	3	0	32/64-bit	32	32/64-bit	Workstation, Embedded
SPARC	Reg-Reg	3	0	32/64-bit	24-32	32/64-bit	Workstation
TI C6000	Reg-Reg	3	0	32-bit	32	32-bit	DSP
IBM 360	Reg-Mem	2	1	8/16/32/64-bit	16	24/31/64-bit	Mainframe
x86	Reg-Mem	2	1	8/16/32/64-bit	4/8/16	16/32/64-bit	Personal Computers, HPC
RISC-V	Reg-Reg	3	0	32/64/128-bit	32	32/64-bit	Embedded

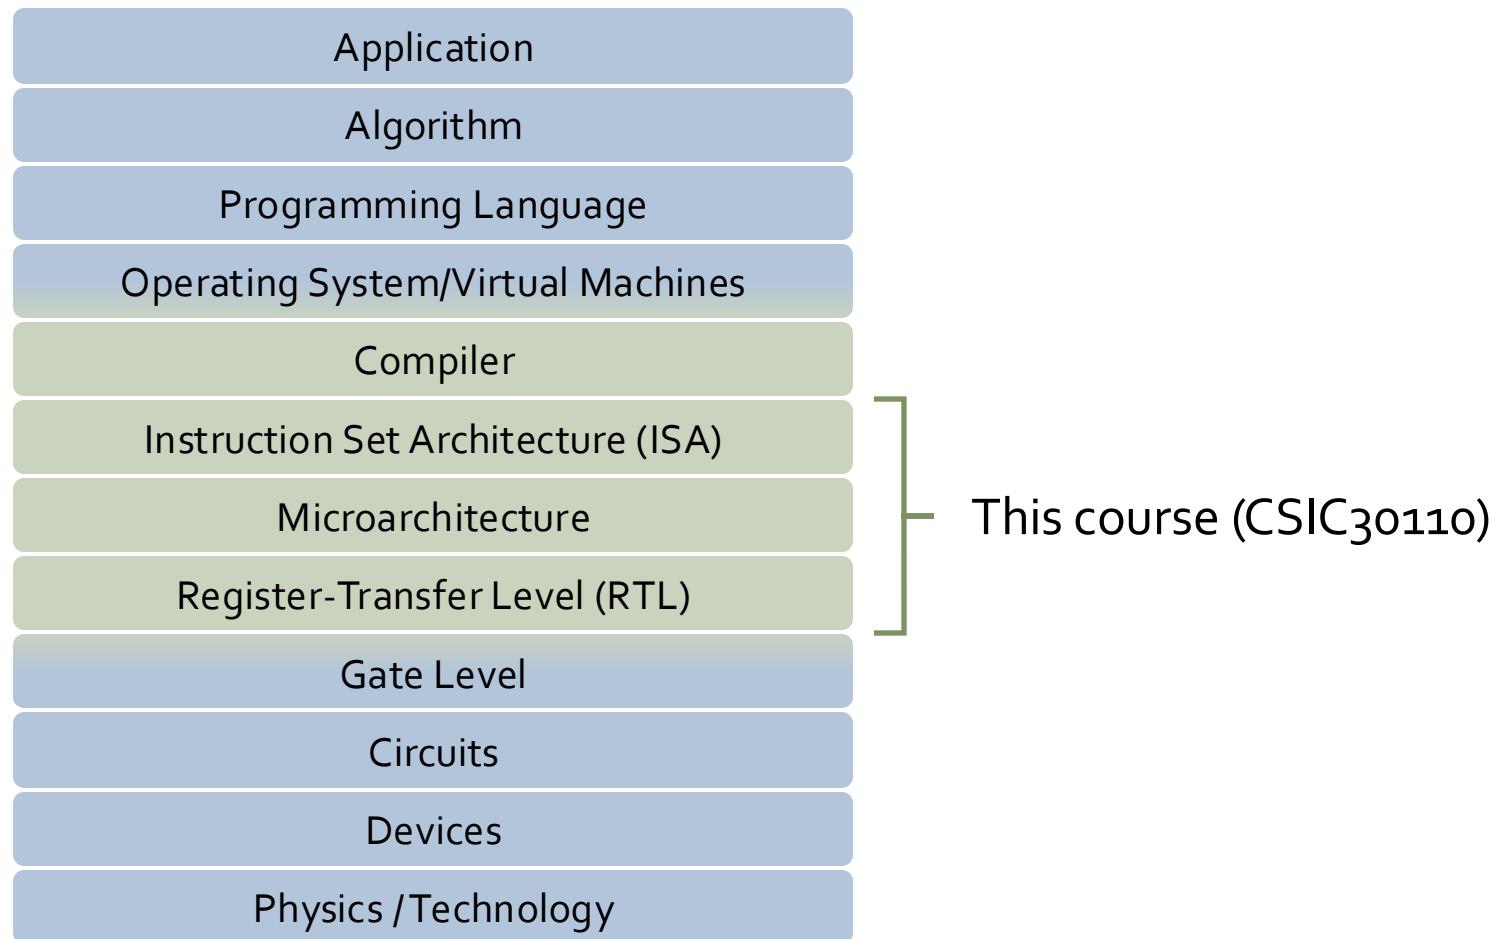
Why the Diversity in ISAs?

- Technology influenced ISA
 - Storage is expensive, tight encoding space
 - Reduced Instruction Set Computer
 - Remove instructions until whole computer fits on die
 - Multicore/Manycore
- Application Influenced ISA
 - Instructions for applications
 - DSP instructions
 - Compiler technology has improved

Recap

- What is Computer Architecture?
- Course Admin
- ISA Review

Recap



Acknowledgements

- These slides contain material developed and copyright by:
 - Arvind (MIT)
 - Krste Asanovic (MIT/UCB)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)
 - Christopher Batten (Cornell)
 - David Wentzlaff (Princeton)
- MIT material derived from course 6.823
- UCB material derived from course CS252
- Cornell material derived from course ECE 4750
- Princeton material derived from course ECE 475