

# **Computer Architecture**

## **Superpipelining and Branch Prediction**

Ting-Jung Chang


NYCU CS

# Recap: Exploiting ILP (Instruction-Level Parallelism)

Common name	Issue structure	Hazard detection	Scheduling	Distinguishing characteristic	Examples
Superscalar (static)	Dynamic	Hardware	Static	In-order execution	Mostly in the embedded space: MIPS and ARM, including the Cortex-A53
Superscalar (dynamic)	Dynamic	Hardware	Dynamic	Some out-of-order execution, but no speculation	None at the present
Superscalar (speculative)	Dynamic	Hardware	Dynamic with speculation	Out-of-order execution with speculation	Intel Core i3, i5, i7; AMD Phenom; IBM Power 7
VLIW/LIW	Static	Primarily software	Static	All hazards determined and indicated by compiler (often implicitly)	Most examples are in signal processing, such as the TI C6x
EPIC	Primarily static	Primarily software	Mostly static	All hazards determined and indicated explicitly by the compiler	Itanium

H&P 6, Fig. 3.19

# Recap: VLIW: Very Long Instruction Word



- Multiple operations packed into one instruction
- Each operation slot is for a fixed function
- Constant operation latencies are specified
- Architecture requires guarantee of:
  - Parallelism within an instruction => no cross-operation RAW check
  - No data use before data ready => no data interlocks

# Recap: Limits of Static Scheduling

- Statically unpredictable branches
- Variable memory latency (unpredictable cache misses)
- Code size explosion
- Compiler complexity
- Despite several attempts, VLIW has failed in general-purpose computing arena.
  - More complex VLIW architectures are close to in-order superscalar in complexity, no real advantage on large complex apps.
- Successful in embedded DSP market
  - Simpler VLIWs with more constrained environment, friendlier code.
  - Google's Tensor Processing Unit is a VLIW machine.

# Course Administration

- PS2 solution released
- PS3 is out
- Midterm 10/21
  - Lecture 1-6/7
  - Ps1 – Ps3
  - 90 minutes closed-book exam (13:20~)

# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Agenda


- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Raising Clock Frequency

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Reducing cycle-time reduces program time
  - Increase pipeline depth
- Reducing CPI reduces program time
  - Wider issue/execute processors
- Reducing number of instructions reduces program time
  - Compiler optimization and complex instructions

# Increasing Frequency Increases CPI



[Sprangle and Carmean ISCA 2002]

# Workload Influences Optimal Pipeline Depth





Fig. 9 shows the distribution of pipeline depth optima for different workload classes. Blue are the Spec (C) workloads, red are the traditional workloads, and green are the modern (C++ and Java) workloads.


# Clock Frequency Improvement



# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Pentium 4



# Pentium III vs. Pentium 4 Pipeline


Basic Pentium III Processor Misprediction Pipeline									
1	2	3	4	5	6	7	8	9	10
Fetch	Fetch	Decode	Decode	Decode	Rename	ROB Rd	Rdy/Sch	Dispatch	Exec
Basic Pentium 4 Processor Misprediction Pipeline									
1	2	3	4	5	6	7	8	9	10
TC Nxt IP	TC Fetch	Drive Alloc	Rename	Que	Sch	Sch	Sch	Disp	Disp
11	12	13	14	15	16	17	18	19	20
RF	RF	Ex	Flgs	Br Ck	Drive	Many more stages...			

- Drive Stage needed just to drive signals between pipelines

# Pentium 4 Pipeline Stages

- It is a 5-issue superscalar processor
- Increasing the number of pipeline stages increases the clock frequency
  - It took the industry 28 years to hit 1 GHz and only 18 months to reach 2 GHz.
  - The price paid for deeper pipelines is that it is very difficult to avoid stalls (That is why when Pentium 4 was introduced its performance was worse than Pentium 3.)

# Double Pumped ALU




- Executes at 2x core frequency!
- Bypass low bits early within pipeline
- Cut carry of 32-bit add
- Calculate Flags in 3<sup>rd</sup> cycle


# Small L1 Cache for Low Latency

- 8KB L1 Cache
- Much smaller than contemporary L1 caches
- Done to allow low latency of 2-cycles.
- Write-through to L2
- Low latency cache especially important to x86 ISA because so few registers requires going to stack often

# Trace Cache



# Pentium 4 Architecture



# Other Deeply Pipelined Processor

Really Deep:

- POWER 7 / POWER 7+ / POWER 8

Deep:

- Intel Core
- Intel Core 2
- Intel Core i3, i5, i7 (Nehalem etc.)
- Intel Core i3, i5, i7 2<sup>nd</sup> Gen. (Sandy/Ivy Bridge)
- POWER 9


# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Problems with Superpipelining

- Hard to build
- Hard to cut critical paths
- More bypassing needed
- Flip-flop overheads become larger part of cycle
- Large pipelining requires more “pipeline” parallelism in program
  - Requires better branch prediction
- Power!!!

# Flip Flops/Latches Dominating Critical Path




- Flip Flop/Latch delay becomes larger portion of cycle
- Setup Time dominates with short cycles
- Can also have Hold Time problems


# Partitioning/Clustering of Structures

- Faster clock frequency makes global communication challenging
- Faster clock frequency makes centralizes structures harder to meet timing (Ex: Issue Queue, ROB)
- Solution, Partition/Cluster resources and add extra latency to cross between clusters. (Ex: 21264)


# Bypassing in Superscalar Pipelines



# Bypassing in Superscalar Pipelines




# Alpha 21264 Is Clustered




# Branch Prediction

- Long Pipes demand **spectacular** branch predictors



# Power Limits Performance




# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- **Classifying ILP Machines**
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Classifying ILP Machines

## Baseline scalar RISC


- Issue parallelism = IP = 1
- Operation latency = OP = 1
- Peak IPC = 1



# Classifying ILP Machines

## Underpipelined Machines


- cycle > operation latency
- issues < 1 instr. per cycle



# Classifying ILP Machines

## Superscalar Machines


- Issue parallelism = IP = n inst / cycle
- Operation latency = OP = 1 cycle
- Peak IPC = n inst / cycle (n x speedup?)



# Classifying ILP Machines

## VLIW Machines (Very Long Instruction Word)


- Issue parallelism = IP =  $n$  inst / cycle
- Operation latency = OP = 1 cycle
- Peak IPC =  $n$  inst / cycle = 1 VLIW / cycle



# Classifying ILP Machines

## VLIW Machines (Very Long Instruction Word)


- Issue parallelism = IP =  $n$  inst / cycle
- Operation latency = OP = 1 cycle
- Peak IPC =  $n$  inst / cycle = 1 VLIW / cycle



# Classifying ILP Machines

## Superpipelined


- Issue parallelism = IP = 1 inst / minor cycle
- Operation latency = OP = m minor cycles
- Peak IPC = m inst / major cycle ( $m \times$  speedup?)



# Classifying ILP Machines

## Superpipelined-Superscalar

- Issue parallelism = IP = n inst / minor cycle
- Operation latency = OP = m minor cycles
- Peak IPC = n x m instr / major cycle




# Superscalar vs. Superpipelined

- Roughly equivalent performance
  - If  $n = m$  then both have about the same IPC
  - Parallelism exposed in space vs. time


# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- **Branch Cost Motivation**
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Longer Frontends Means More Control Flow Penalty




# Longer Pipeline Frontends Amplify Branch Cost



Pentium 3: 10 cycle branch penalty

Pentium 4: 20 cycle branch penalty

# Dual Issue and Branch Cost



# Superscalars Multiply Branch Cost

BEQZ	F	D	I	A0	A1	W
OpA	F	D	I	B0	-	-
OpB		F	D	I	-	-
OpC		F	D	I	-	-
OpD			F	D	-	-
OpE			F	D	-	-
OpF			F	-	-	-
OpG			F	-	-	-
OpH				F	D	I
OpI				F	D	I
				A0	A1	W
				B0	B1	W



Dual-issue Processor  
has twice the  
mispredict penalty

How much work is lost if pipeline doesn't follow correct instruction flow?

~pipeline width x branch penalty

# Average Run-Length between Branches

- Average dynamic instruction mix of SPEC CPU 2017 :

	SPECint	SPECfp
Branches	19%	11%
Loads	24%	26%
Stores	10%	7%
Other	47%	56%

[Limaye et al, ISPASS'18]

- SPECint<sub>17</sub>: perlbench, gcc, mcf, omnetpp, xalancbmk, x264, deepsjeng, leela, exchange2, xz
- SPECfp<sub>17</sub>: bwaves, cactus, lbm, wrf, pop2, imagick, nab, fotonik3d, roms

What is the average run length between branches?

# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Branch Prediction

- Essential in modern processors to mitigate branch delay latencies

Two types of Prediction

1. Predict Branch Outcome
2. Predict Branch/Jump Address

# RISC-V Branches and Jumps

- Each instruction fetch depends on one or two pieces of information from the preceding instruction:
  1. Is the preceding instruction a taken branch?
  2. If so, what is the target address?

Instruction

Taken known?


Target known?

JAL

JALR

BRANCH

# Where is the Branch Information Known?



# Reducing Control Flow Penalty

- Software solutions
  - Eliminate branches
    - loop unrolling Increases run length between branches
  - Reduce resolution time – instruction scheduling
    - Compute the branch condition as early as possible (of limited value)

# Reducing Control Flow Penalty

- Software solutions
  - Eliminate branches
    - loop unrolling Increases run length between branches
  - Reduce resolution time – instruction scheduling
    - Compute the branch condition as early as possible (of limited value)
- Hardware solutions
  - Bypass – usually results are used immediately
  - Change architecture – find something else to do Delay slots
    - replace pipeline bubbles with useful work (requires software cooperation)
  - Speculate (accurately) – branch prediction
    - Speculative execution of instructions beyond the branch

# Branch Delay Slots

(expose control hazard to software)

- Change the ISA semantics so that the instruction that follows a jump or branch is always executed
  - gives compiler the flexibility to put in a useful instruction where normally a pipeline bubble would have resulted.

I <sub>1</sub>	096	ADD	
I <sub>2</sub>	100	BEQZ x1 +200	
I <sub>3</sub>	104	ADD	Delay slot instructions
I <sub>4</sub>	108	ADD	executed regardless of branch
I <sub>5</sub>	304	ADD	outcome

# Branch Prediction


- Motivation:
  - Branch penalties limit performance of deeply pipelined processors
  - Modern branch predictors have high accuracy (>95%) and can reduce branch penalties significantly
- Required hardware support:
  - Prediction structures:
    - Branch history tables, branch target buffers, etc.
  - Mispredict recovery mechanisms:
    - Keep result computation separate from commit
    - Kill instructions following branch in pipeline
    - Restore state to state following branch

# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Static Branch Prediction

Overall probability a branch is taken is ~60-70% but:



# Static Software Branch Prediction

- Extend ISA to enable compiler to tell microarchitecture if branch is likely to be taken or not (Can be up to 80% accurate)

	BR.T	F	D	X	M	W	
OpA		F	-	-	-	-	
Targ		F	D	X	M	W	
BR.NT		F	D	X	M	W	
OpB			F	D	X	M	W
OpC			F	D	X	M	W

What if hint is wrong?

	BR.T	F	D	X	M	W	
OpA		F	-	-	-	-	
Targ		F	-	-	-	-	
OpA			F	D	X	M	W

# Static Hardware Branch Prediction

1. Always Predict Not-Taken
  - What we have been assuming
  - Simple to implement
  - Know fall-through PC in Fetch
  - Poor Accuracy, especially on backward branches
2. Always Predict Taken
  - Difficult to implement because don't know target until Decode
  - Poor accuracy on if-then-else
3. Backward Branch Taken, Forward Branch Not Taken
  - Better Accuracy
  - Difficult to implement because don't know target until Decode

# Agenda

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Dynamic Branch Prediction

Learning based on past behavior


- Temporal correlation
  - The way a branch resolves may be a good predictor of the way it will resolve at the next execution
- Spatial correlation
  - Several branches may resolve in a highly correlated manner (a preferred path of execution)

# Dynamic Hardware Branch Prediction:


## Exploiting Temporal Correlation

- Exploit structure in program: The way a branch resolves may be a good indicator of the way it will resolve the next time it executes (Temporal Correlation)

1-bit Saturating Counter



# 1-bit Saturating Counter




Iteration	Prediction	Actual	Mispredict?
1	NT	T	Y
2	T	T	
3	T	T	
4	T	NT	Y
...			
1	NT	T	Y
2	T	T	
3	T	T	
4	T	NT	Y

For Backward branch in loop

- Assume 4 Iterations
- Assume is executed multiple times

Always 2 Mispredicts


# 2-bit Saturating Counter




Iteration	Prediction	Actual	Mispredict?	State
1	NT	T	Y	Strong NT
2	NT	T	Y	Weak NT
3	T	T		Weak T
4	T	NT	Y	Strong T
...				
1	T	T		Weak T
2	T	T		Strong T
3	T	T		Strong T
4	T	NT	Y	Strong T

Only 1  
Mispredict  
at end of loop

# Other 2-bit FSM Branch Predictors



# Branch History Table (BHT)



4K-entry BHT, 2 bits/entry, ~80-90% correct predictions


# Exploiting Spatial Correlation

[Yeh and Patt, 1992]


```
if (x[i] < 7) then  
    y += 1;  
if (x[i] < 5) then  
    c -= 4;
```

If first condition false, second condition also false

Branch History Register, **BHR**, records the direction of the last N branches executed by the processor (Shift Register)



# Pattern History Table (PHT)



# Two-Level Branch Predictor




Figure 1: The structure of the Two-Level Adaptive Training scheme.






Figure 2: The state transition diagrams of the finite-state machines used for updating the pattern history in the pattern table entry.

# Two-Level Branch Predictor




# Generalized Two-Level Branch Predictor



# Tournament Branch Predictor

[Alpha 21264, 1996]




- Choice predictor learns whether best to use local or global branch history in predicting next branch
- Global history is speculatively updated but restored on mispredict
- Claim 90-100% success on range of applications

# TAGE Predictor


[Seznec & Michaud, 2006]

- Multiple tagged tables, use different global history lengths
- Set of history lengths forms a geometric series

{0, 2, 4, 8, 16, 32, 64, 128, 256, ..., 2048}



# Tagged Geometric History Length (TAGE)




Longest matching table provides the prediction, subject to branch confidence

# Agenda


- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Predicting Target Address



Even with best possible prediction of branch outcome, still have to wait for branch target address to be determined

# Branch Target Buffer (BTB)



Put BTB in Fetch Stage in parallel with PC+4 Speculation logic


# **BTB is only for Control Instructions**

- BTB contains useful information for branch and jump instructions only
  - Do not update it for other instructions
- For all other instructions the next PC is PC+4 !

**How to achieve this effect without decoding the instruction?**

# Consulting BTB Before Decoding

Assume a 128-entry BTB



- The match for  $\text{PC}=1028$  fails and  $1028+4$  is fetched
  - eliminates false predictions after ALU instructions
- BTB contains entries only for control transfer instructions
  - more room to store branch targets

# Uses of Jump Register (JALR)


- Switch statements (jump to address of matching case)  
BTB works well if same case used repeatedly
- Dynamic function call (jump to run-time function address)  
BTB works well if same function usually called, (e.g., in C++ programming, when objects have same type in virtual function call)
- Subroutine returns (jump to return address)  
BTB works well if usually return to the same place  
⇒ Often one function called from many distinct call sites!

How well does BTB work for each of these cases?

# Subroutine Return Stack

- Small structure to accelerate JR for subroutine returns, typically much more accurate than BTBs.

```
fa() { fb(); }
fb() { fc(); }
fc() { fd(); }
```



# Recap

- Motivation for Long Pipelines
- Case Study: Intel Pentium 4
- Challenges for Long Pipelines
- Branch Cost Motivation
- Branch Prediction
  - Outcome
    - Static
    - Dynamic
  - Target Address

# Acknowledgements

- These slides contain material developed and copyright by:
  - Arvind (MIT)
  - Krste Asanovic (MIT/UCB)
  - Joel Emer (Intel/MIT)
  - James Hoe (CMU)
  - John Kubiatowicz (UCB)
  - David Patterson (UCB)
  - Christopher Batten (Cornell)
  - David Wentzlaff (Princeton)
- MIT material derived from course 6.823
- UCB material derived from course CS252
- Cornell material derived from course ECE 4750
- Princeton material derived from course ECE 475