

# **Computer Architecture**

## **Advanced caches**

Ting-Jung Chang  
NYCU CS

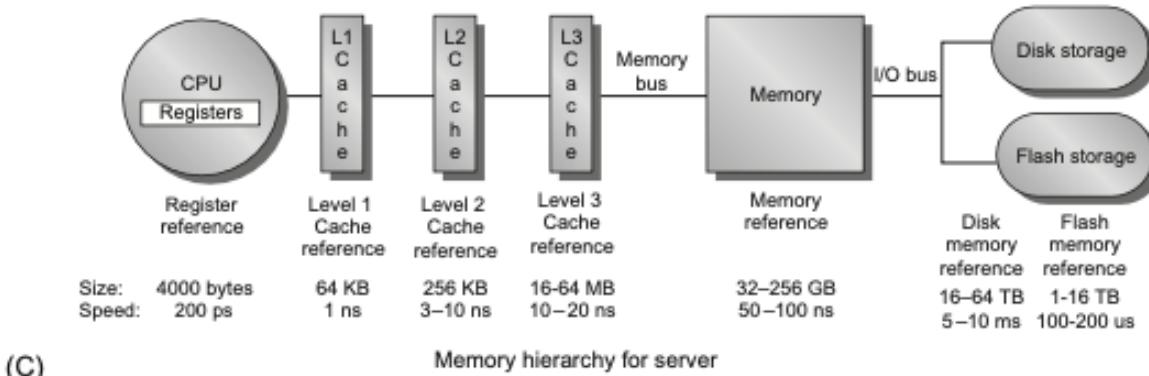
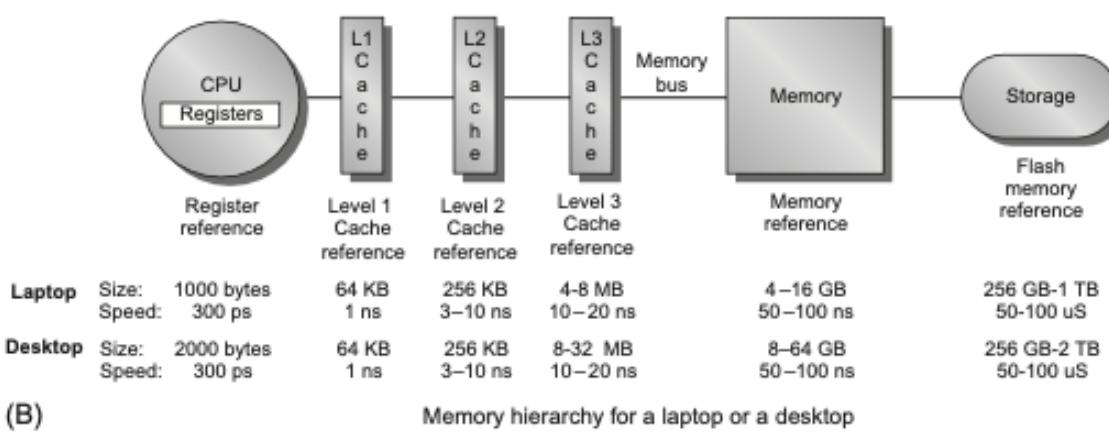
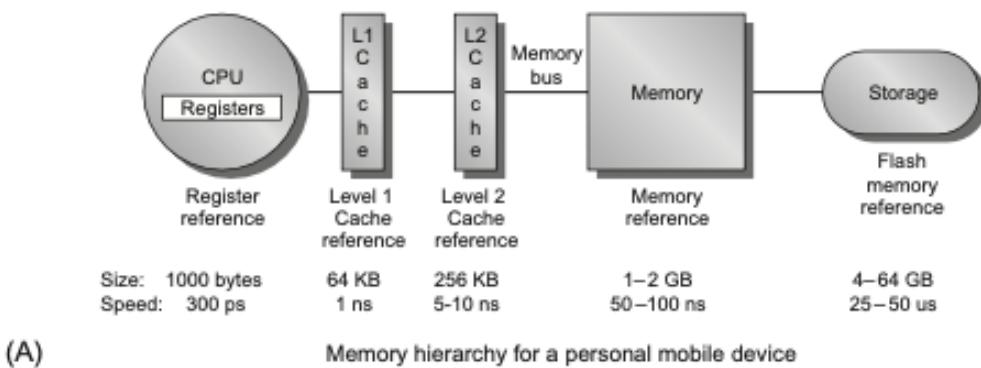
# Agenda

- Review
  - Three C's
  - Basic Cache Optimizations
- Advanced Cache Optimizations
  - Pipelined Cache Write
  - Write Buffer
  - Multilevel Caches
  - Victim Caches
  - Prefetching
    - Hardware
    - Software
  - Multiporting and Banking
  - Software Optimizations
  - Non-Blocking Cache
  - Critical Word First/Early Restart

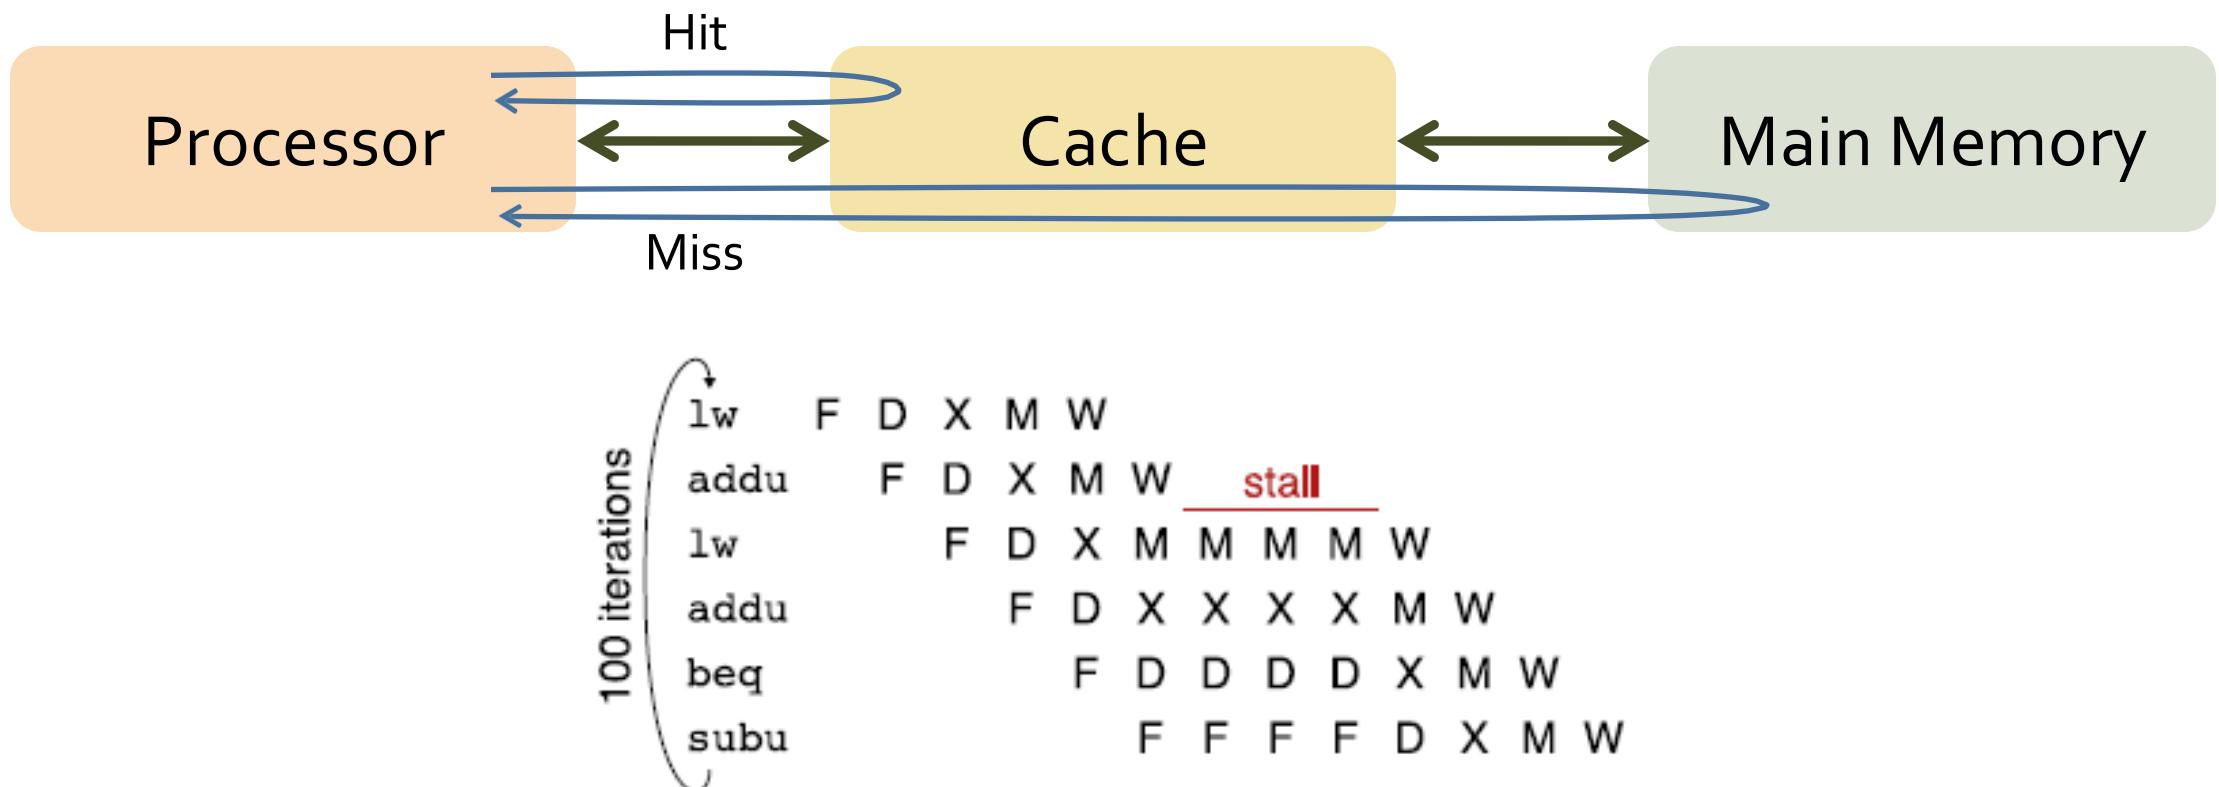
# Summary: Memory Technologies

	Capacity	Latency	Cost/GB
Register	1000s of bits	20 ps	\$\$\$\$
SRAM	~10 KB – 10 MB	1 – 10 ns	~\$1000
DRAM	~ 10 GB	80 ns	~\$10
Flash	~ 100 GB	100 us	~\$1
Hard disk	~ 1TB	10 ms	~\$0.1

- Different technologies have vastly different tradeoffs
- Size is a fundamental limit, even setting cost aside:
  - Small + low latency, high bandwidth, low energy, or
  - Large + high-latency, low bandwidth, high energy
- Can we get the best of both worlds? (large, fast, cheap)

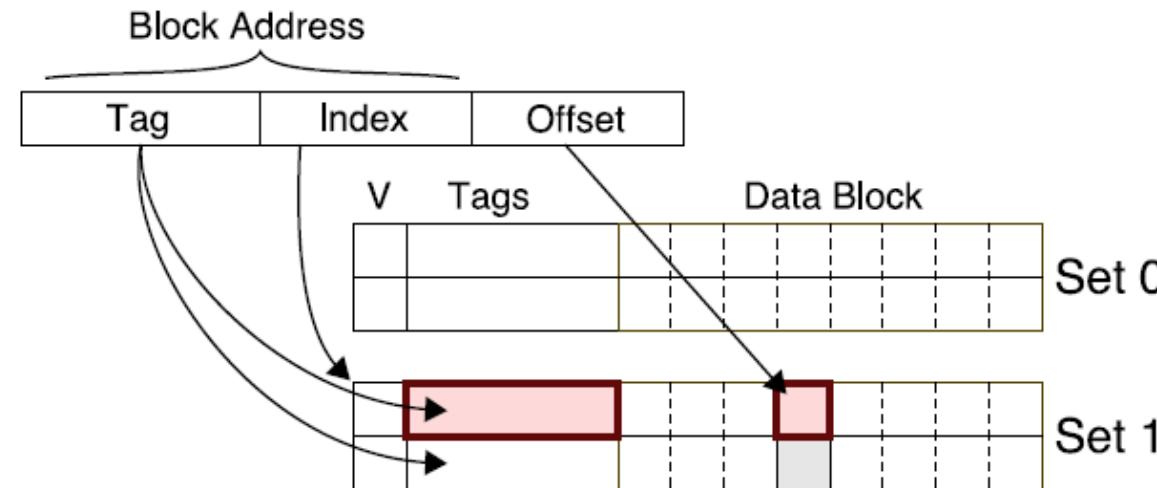


# Average Memory Access Time



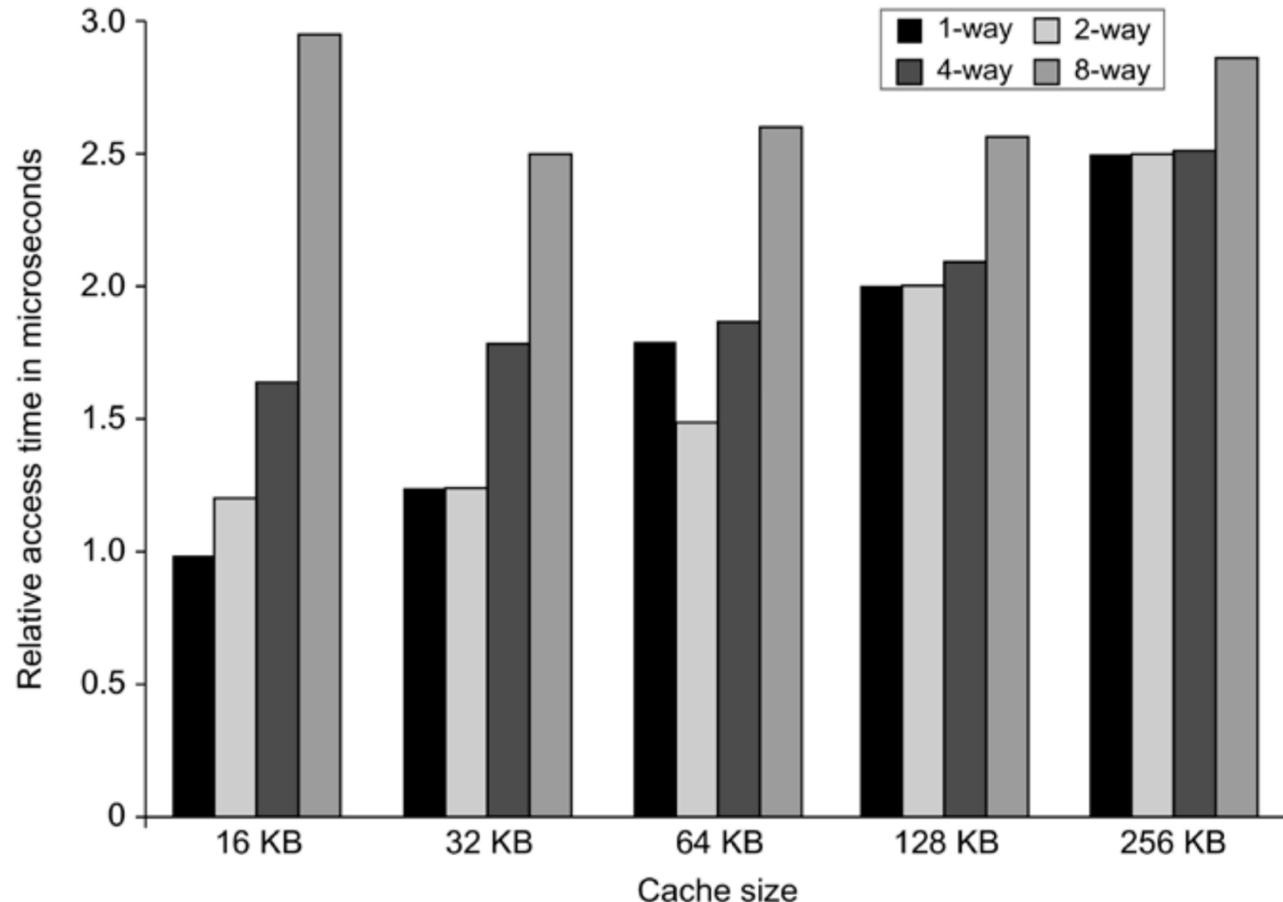
- Average Memory Access Time = Hit Time + ( Miss Rate \* Miss Penalty )

# Categorizing Misses: The Three C's

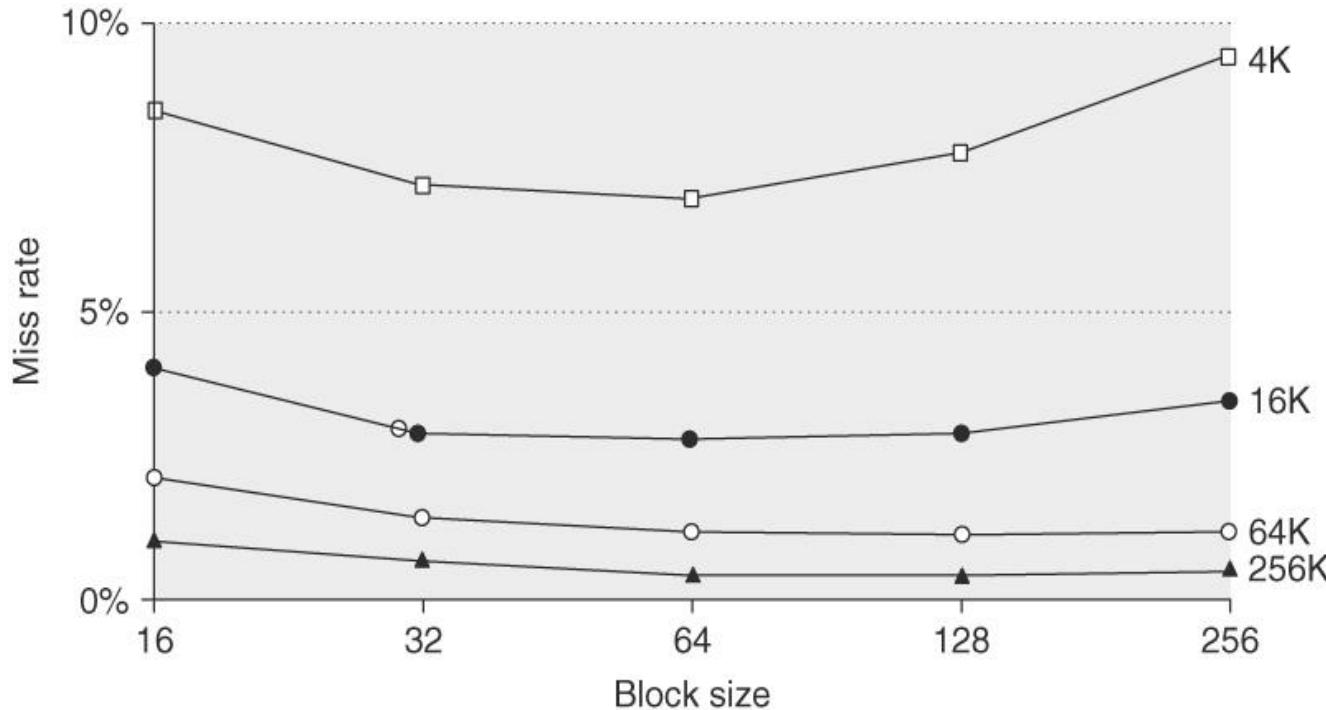


- **Compulsory** – first-reference to a block, occur even with infinite cache
- **Capacity** – cache is too small to hold all data needed by program, occur even under perfect replacement policy (loop over 5 cache lines)
- **Conflict** – misses that occur because of collisions due to less than full associativity (loop over 3 cache lines)

# Reduce Hit Time: Small & Simple Caches

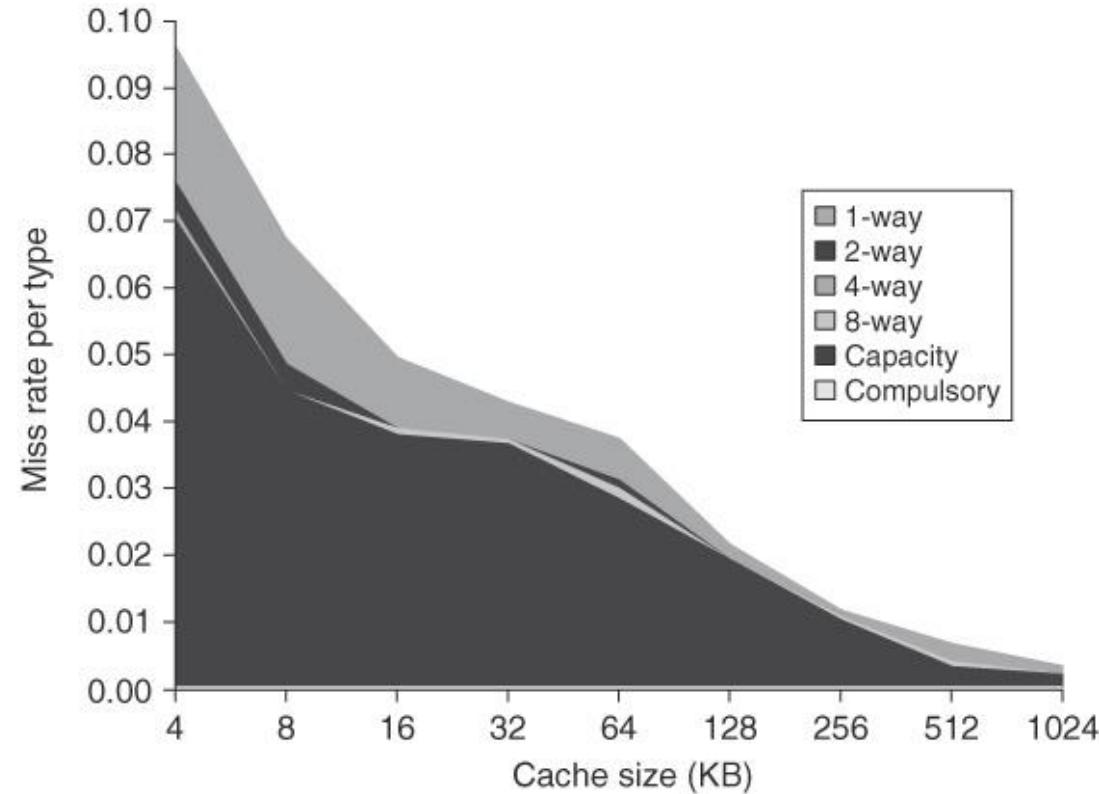


# Reduce Miss Rate: Large Block Size



- Less tag overhead
- Exploit fast burst transfers from DRAM
- Exploit fast burst transfers over wide on-chip busses
- Can waste bandwidth if data is not used
- Fewer blocks -> more conflicts

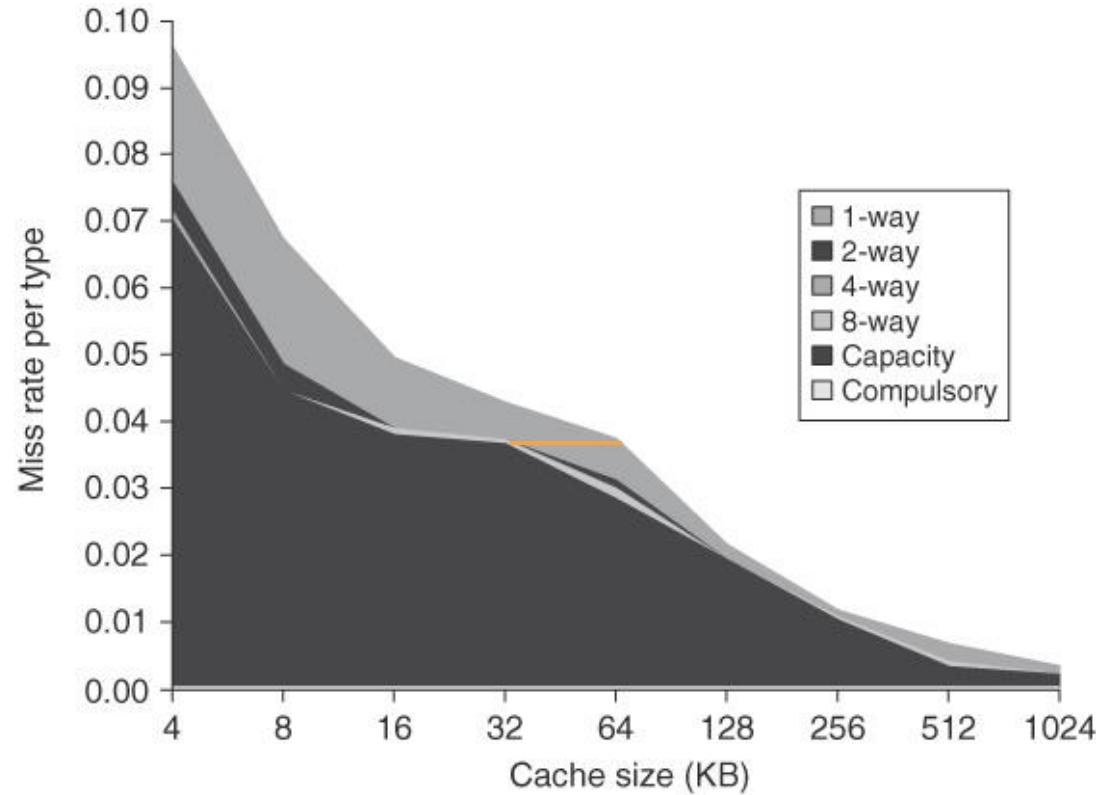
# Reduce Miss Rate: Large Cache Size



Empirical Rule of Thumb:

If cache size is doubled, miss rate usually drops by about  $\sqrt{2}$

# Reduce Miss Rate: High Associativity



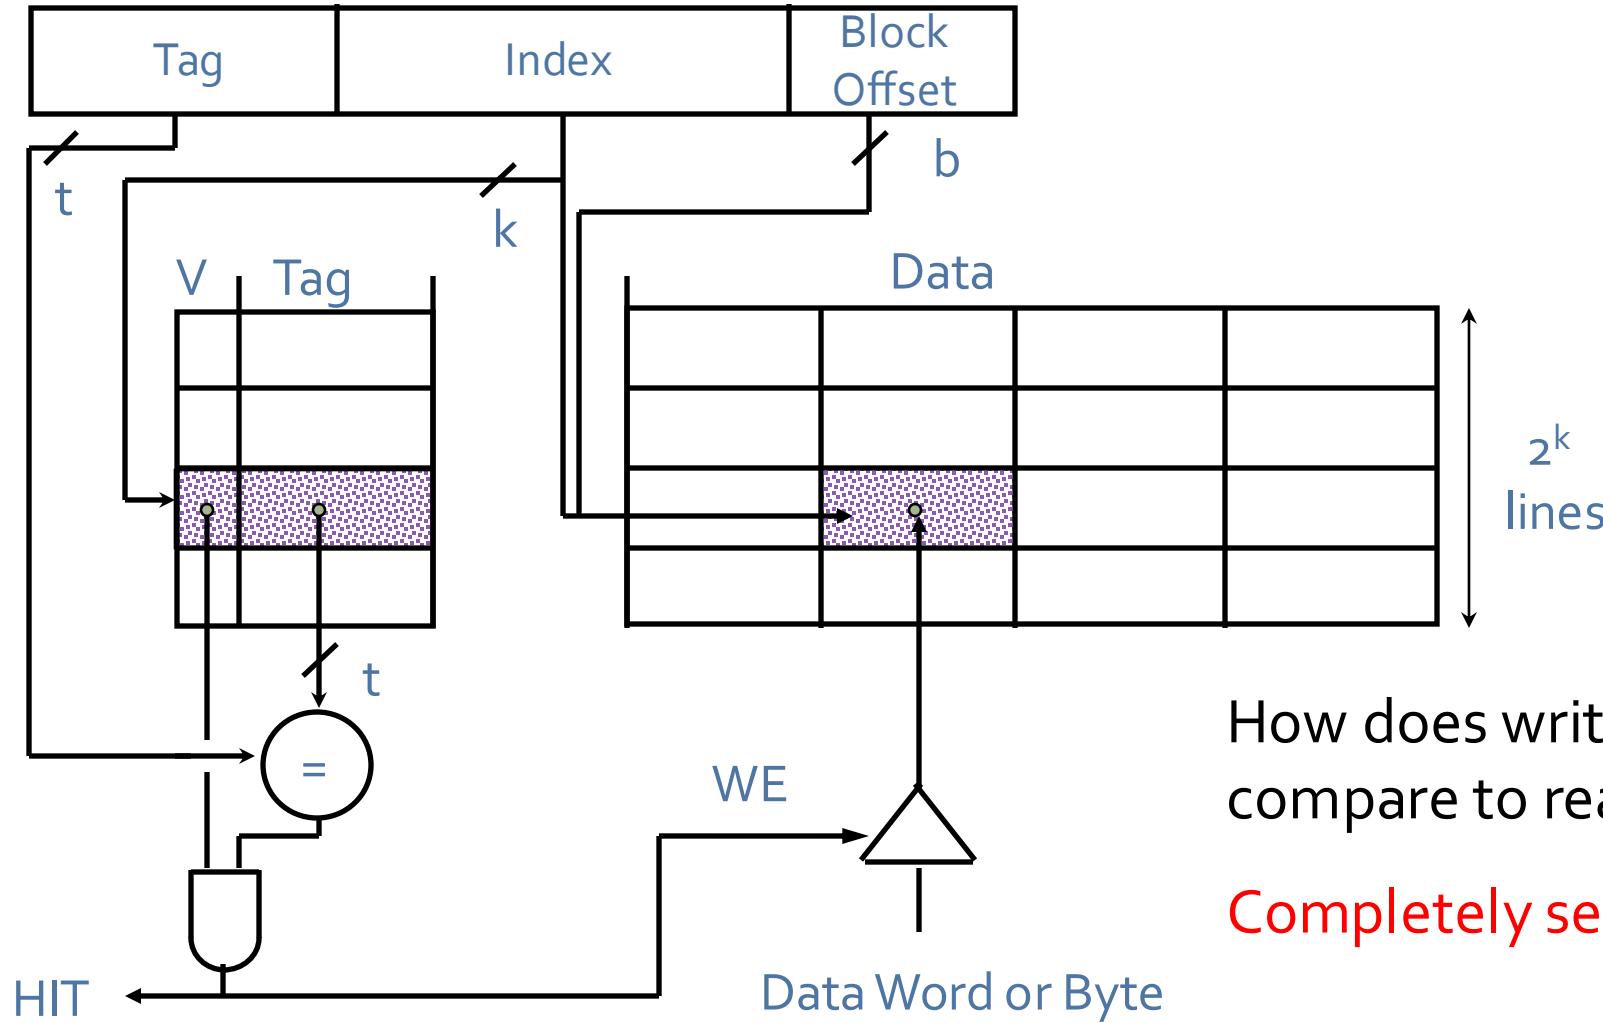
Empirical Rule of Thumb:

Direct-mapped cache of size  $N$  has about the same miss rate as a two-way set-associative cache of size  $N/2$

# Agenda

- Review
  - Three C's
  - Basic Cache Optimizations
- Advanced Cache Optimizations
  - Pipelined Cache Write
  - Write Buffer
  - Multilevel Caches
  - Victim Caches
  - Prefetching
    - Hardware
    - Software
  - Multiporting and Banking
  - Software Optimizations
  - Non-Blocking Cache
  - Critical Word First/Early Restart

# Write Performance



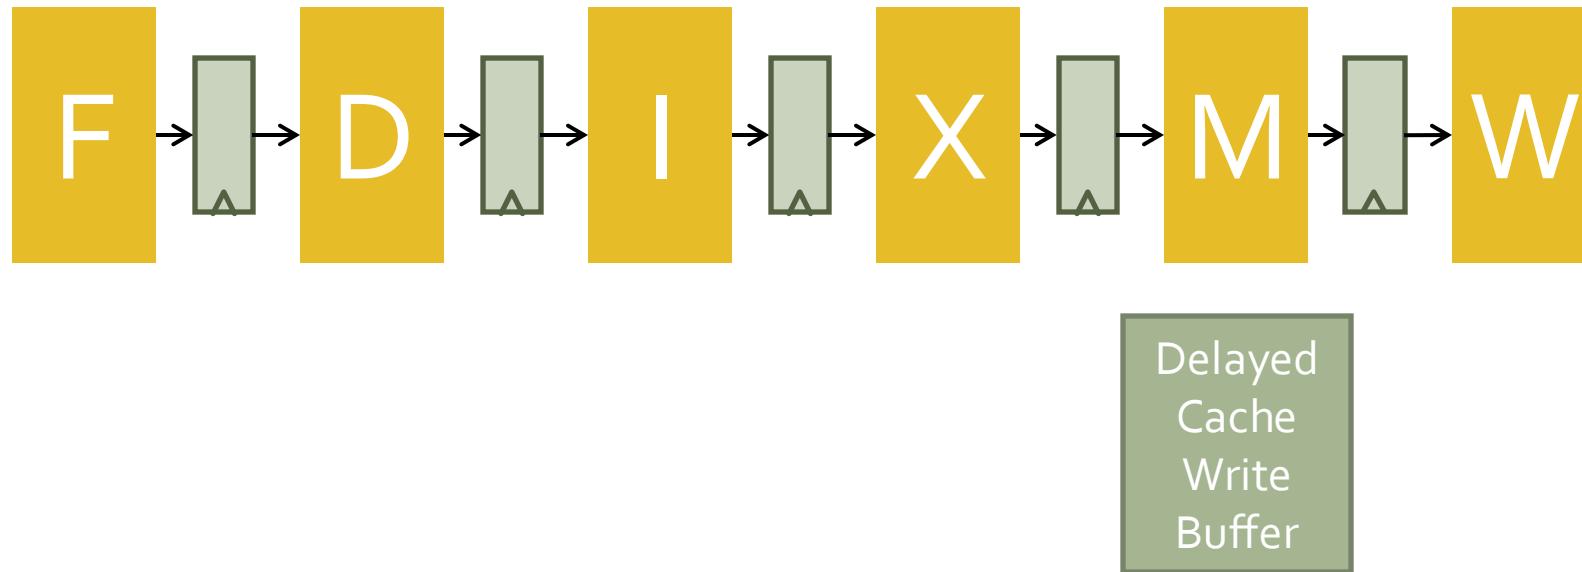
# Reducing Write Hit Time

**Problem:** Writes take two cycles in memory stage, one cycle for tag check plus one cycle for data write if hit

## Solutions:

- Stall - delivering data as fast as possible
  - Circuit-level techniques (CAM tags) may allow one-cycle writes
- Speculate predicting hit with greedy data update:
  - Design data RAM that can perform read and write concurrently
  - Restore old value after tag miss
- Speculate predicting miss with lazy data update:
  - Hold write data for store in single buffer ahead of cache
  - Write cache data during next store's tag check

# Pipelining Cache Writes



Data from a store hit written into data portion of cache during tag access of subsequent store

# Pipelined/Delayed Write Timing

- Problem: Need to commit lazily saved write data
- Solution: Write data during idle data cycle of next store's tag check

Code Sequence:

LD<sub>0</sub>

ST<sub>1</sub>

ST<sub>2</sub>

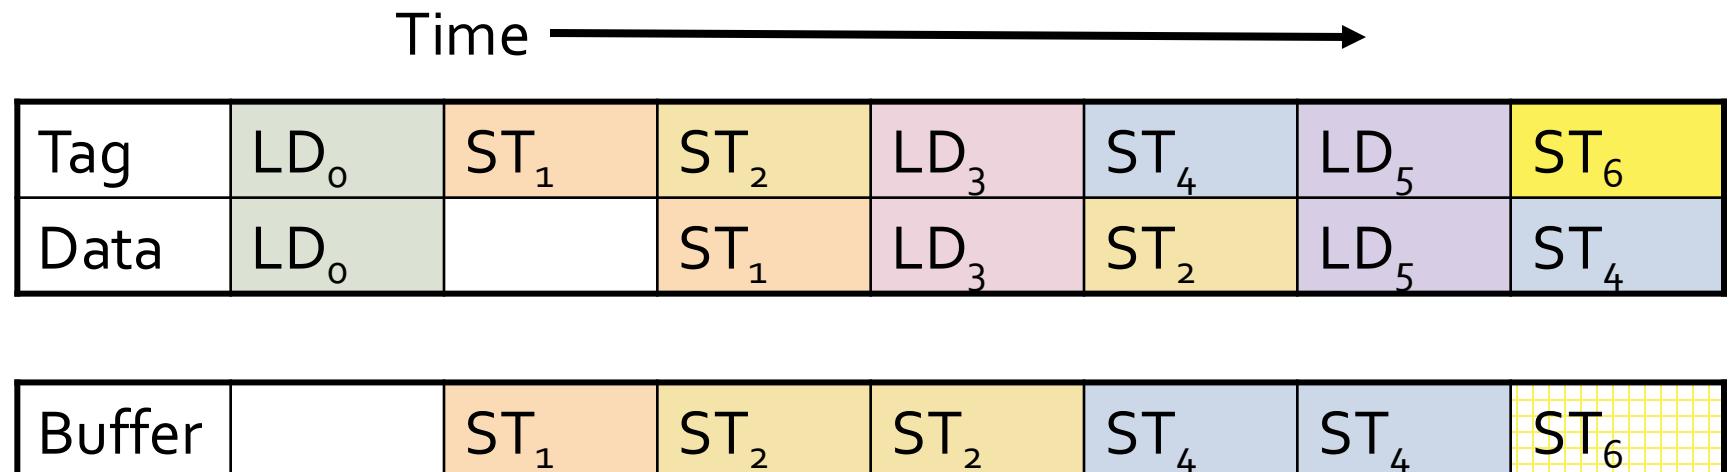
LD<sub>3</sub>

ST<sub>4</sub>

LD<sub>5</sub>

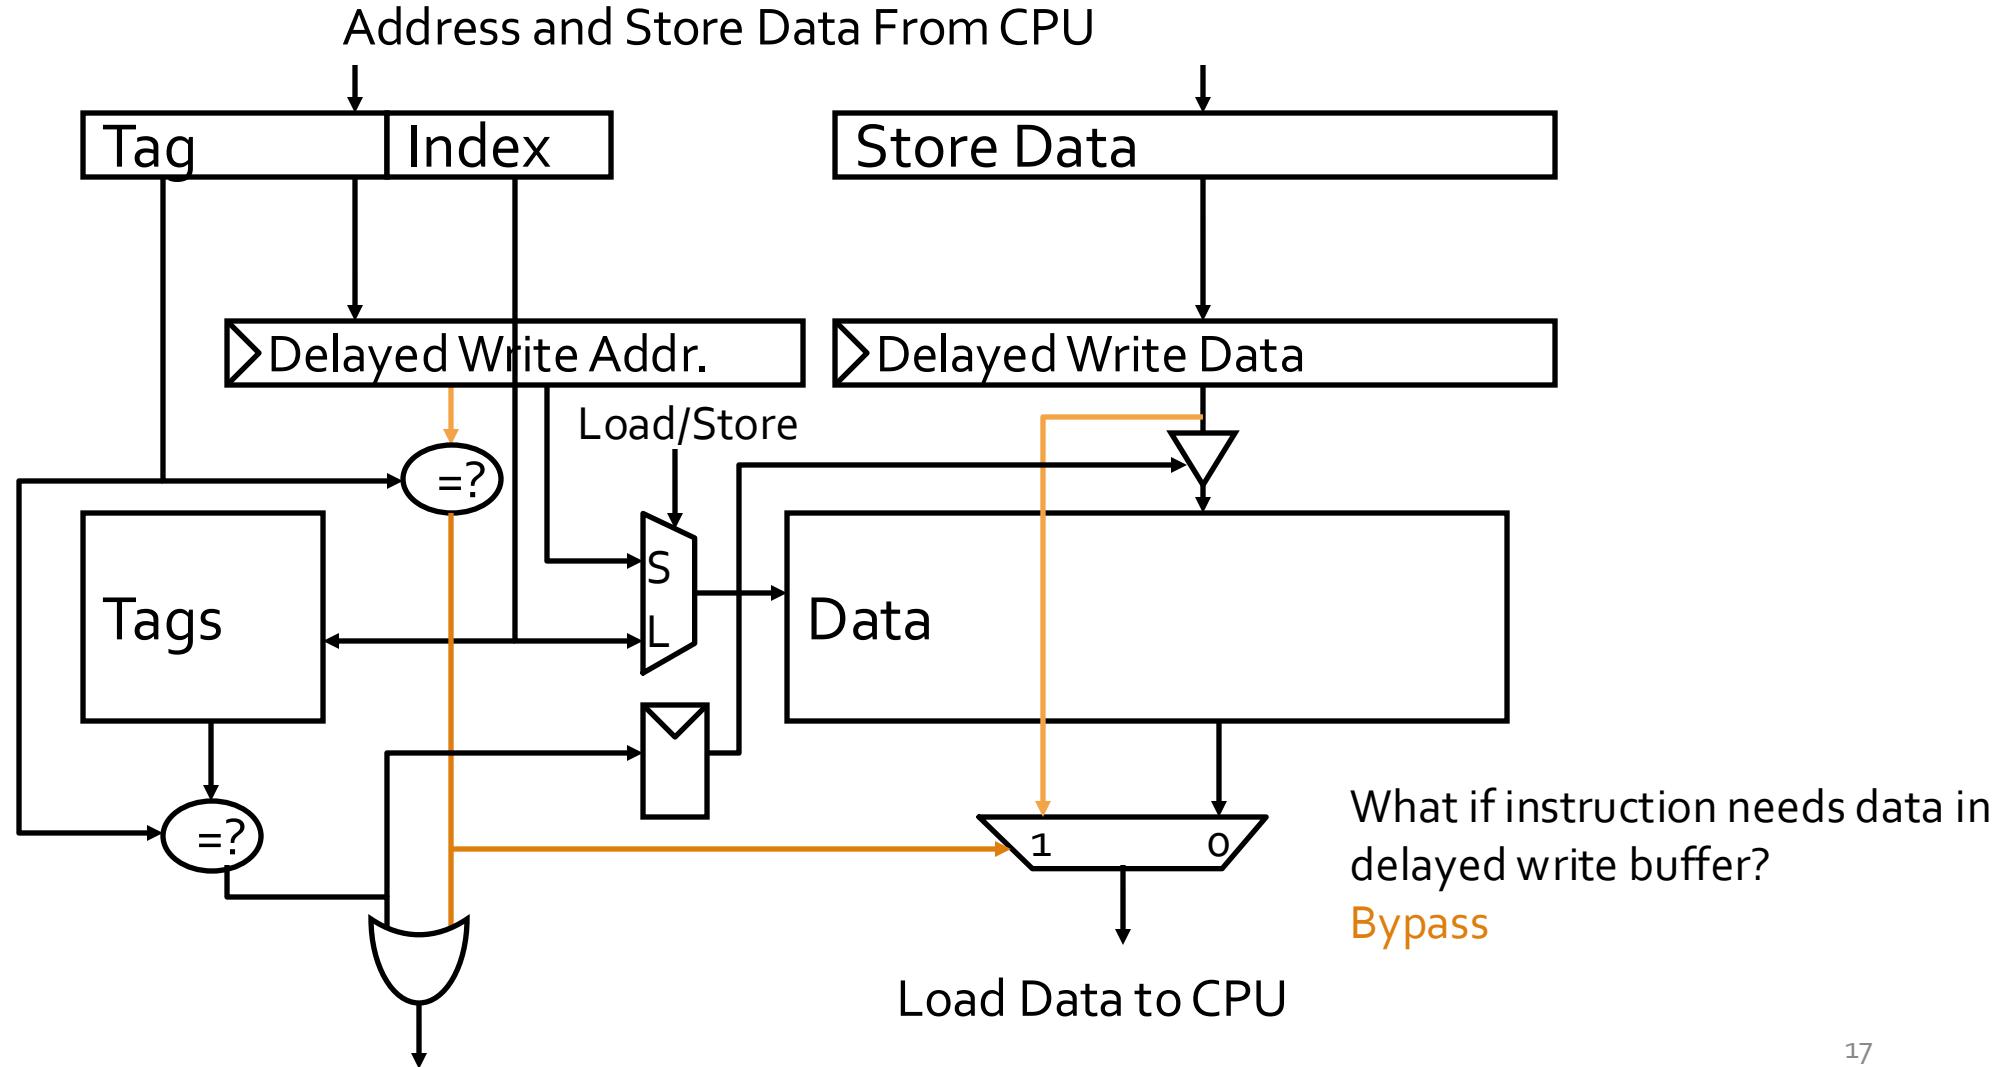
ST<sub>6</sub> miss

Time →



# Pipelining Cache Writes

What is condition for bypass?  
Tag and index match  
“delayed write addr”



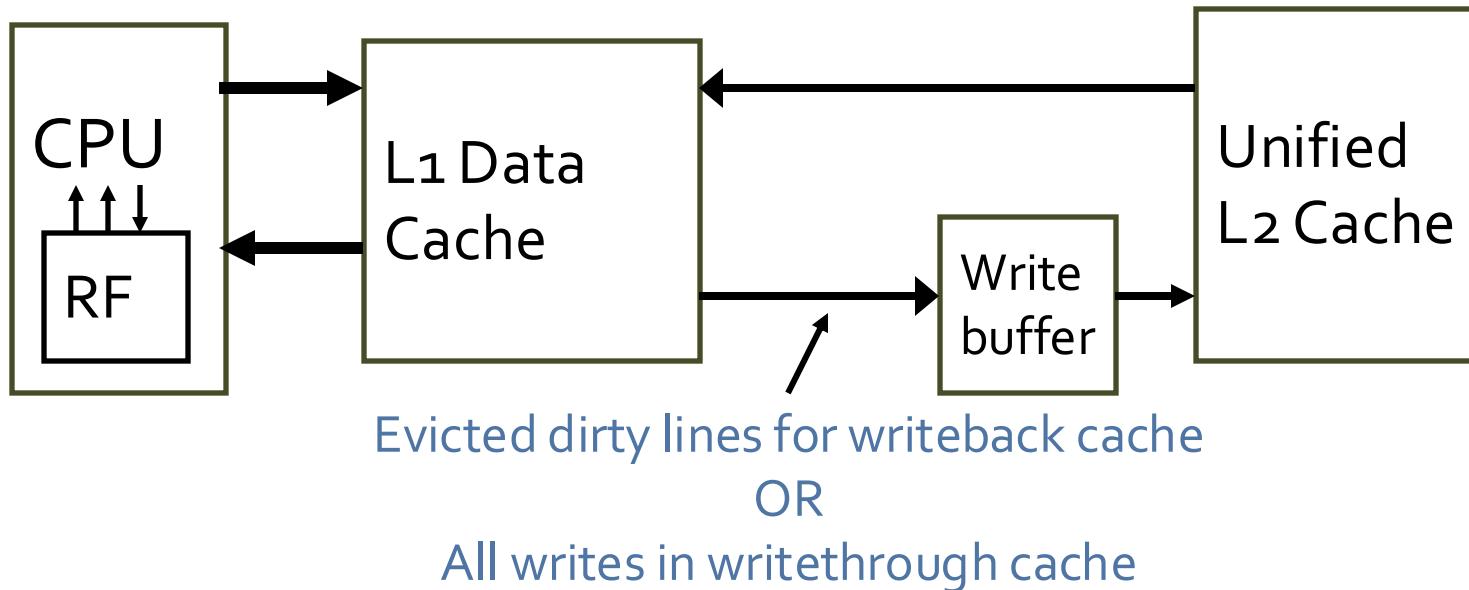
# Pipelined Cache Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Pipelined Writes			—	+

# Write Policy Choices

- Cache Hit
  - **Write Through:** write both cache and memory
    - generally higher traffic but simpler to design
  - **Write Back:** write cache only, memory is written when evicted
    - a dirty bit per block avoids unnecessary write backs, more complicated
- Cache Miss
  - **No Write Allocate:** only write to main memory
  - **Write Allocate:** fetch block into cache, then write
- Common Combinations
  - Write Through & No Write Allocate
  - Write Back & Write Allocate

# Write Buffer to Reduce Read Miss Penalty



Processor is not stalled on writes, and read misses can go ahead of write to main memory

**Problem:** Write buffer may hold updated value of location needed by a read miss (RAW data hazard)

**Simple scheme (stall):** on a read miss, wait for the write buffer to go empty

**Faster scheme (bypass):** Check write buffer addresses against read miss addresses, if no match, allow read miss to go ahead of writes, else, return value in write buffer

# Merging Write Buffer

- When storing to a block that is already pending in write buffer, merge requests to reduce miss penalty

Wr. addr	V	V	V	V
100	1 Mem[100]	0		0
108	1 Mem[108]	0		0
116	1 Mem[116]	0		0
124	1 Mem[124]	0		0

4 separate writes back to lower level memory

Wr. addr	V	V	V	V
100	1 Mem[100]	1 Mem[108]	1 Mem[116]	1 Mem[124]
	0	0	0	0
	0	0	0	0
	0	0	0	0

1 single write back to lower level memory

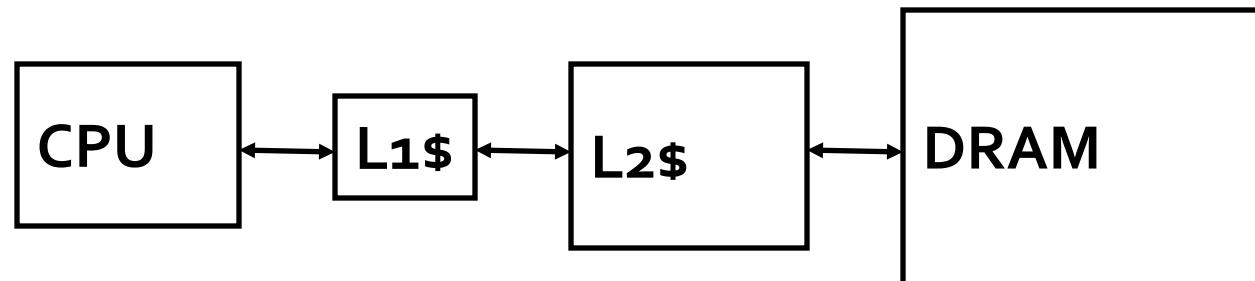
# Write Buffer Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Write Buffer		+		

# Multilevel Caches

**Problem:** A memory cannot be large and fast

**Solution:** Increasing sizes of cache at each level



- Local miss rate = misses in cache / accesses to cache
- Global miss rate = misses in cache / CPU memory accesses
- Misses per instruction = misses in cache / number of instructions

# Presence of L<sub>2</sub> Influences L<sub>1</sub> design

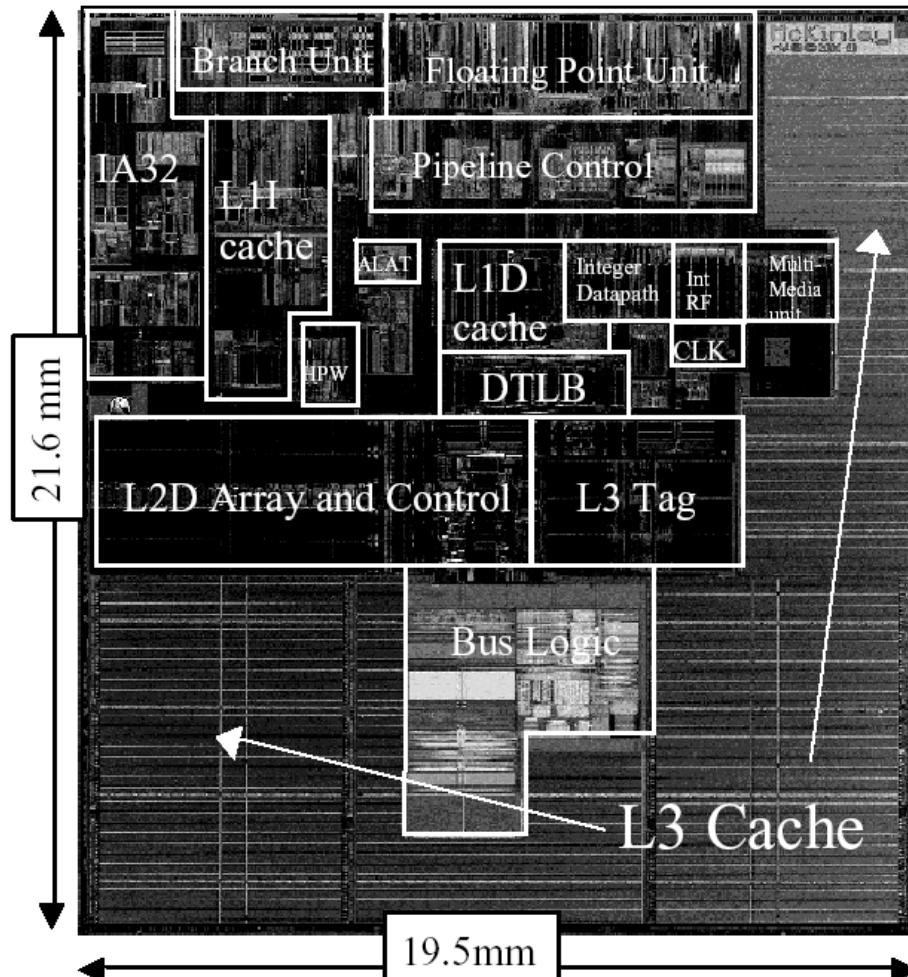
- Use smaller L<sub>1</sub> if there is also L<sub>2</sub>
  - Trade increased L<sub>1</sub> miss rate for reduced L<sub>1</sub> hit time and reduced L<sub>1</sub> miss penalty
  - Reduces average access energy
- Use simpler write-through L<sub>1</sub> with on-chip L<sub>2</sub>
  - Write-back L<sub>2</sub> cache absorbs write traffic, doesn't go off-chip
  - At most one L<sub>1</sub> miss request per L<sub>1</sub> access (no dirty victim write back) simplifies pipeline control
  - Simplifies coherence issues
  - Simplifies error recovery in L<sub>1</sub> (can use just parity bits in L<sub>1</sub> and reload from L<sub>2</sub> when parity error detected on L<sub>1</sub> read)

# Inclusion Policy

- **Inclusive** multilevel cache:
  - Inner cache holds copies of data in outer cache
  - External coherence snoop access need only check outer cache
- **Non-Inclusive** multilevel caches:
  - Inner cache may hold data not in outer cache
  - **Exclusive Caches** Swap lines between inner/outer caches on miss
  - Used in AMD Athlon with 64KB primary and 256KB secondary cache

Why choose one type or the other?

# Itanium-2 On-Chip Caches [Intel/HP, 2002]



**Level 1:** 16KB, 4-way s.a.,  
64B line, quad-port (2  
load+2 store), single cycle  
latency

**Level 2:** 256KB, 4-way s.a.,  
128B line, quad-port (4  
load or 4 store), five cycle  
latency

**Level 3:** 3MB, 12-way s.a.,  
128B line, single 32B port,  
twelve cycle latency

# Power 7 On-Chip Caches [IBM 2009]

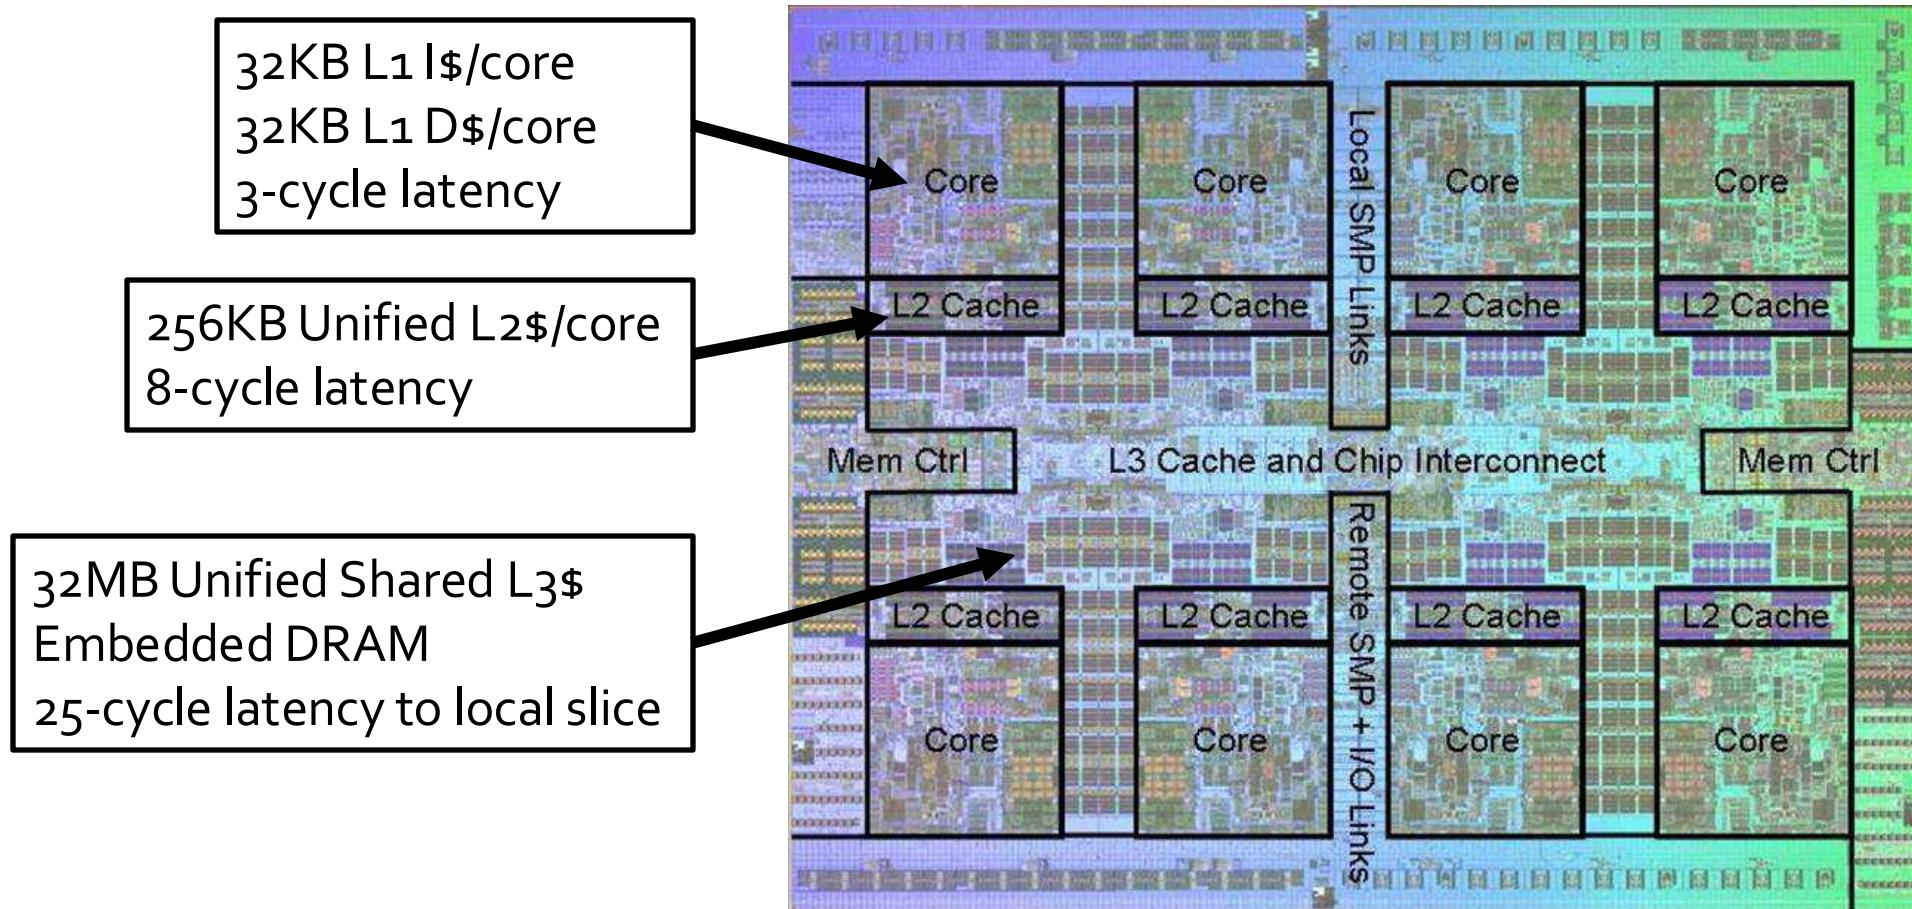


Image Credit: IBM

Courtesy of International Business Machines Corporation,  
© International Business Machines Corporation.

# Apple M3 [2023]

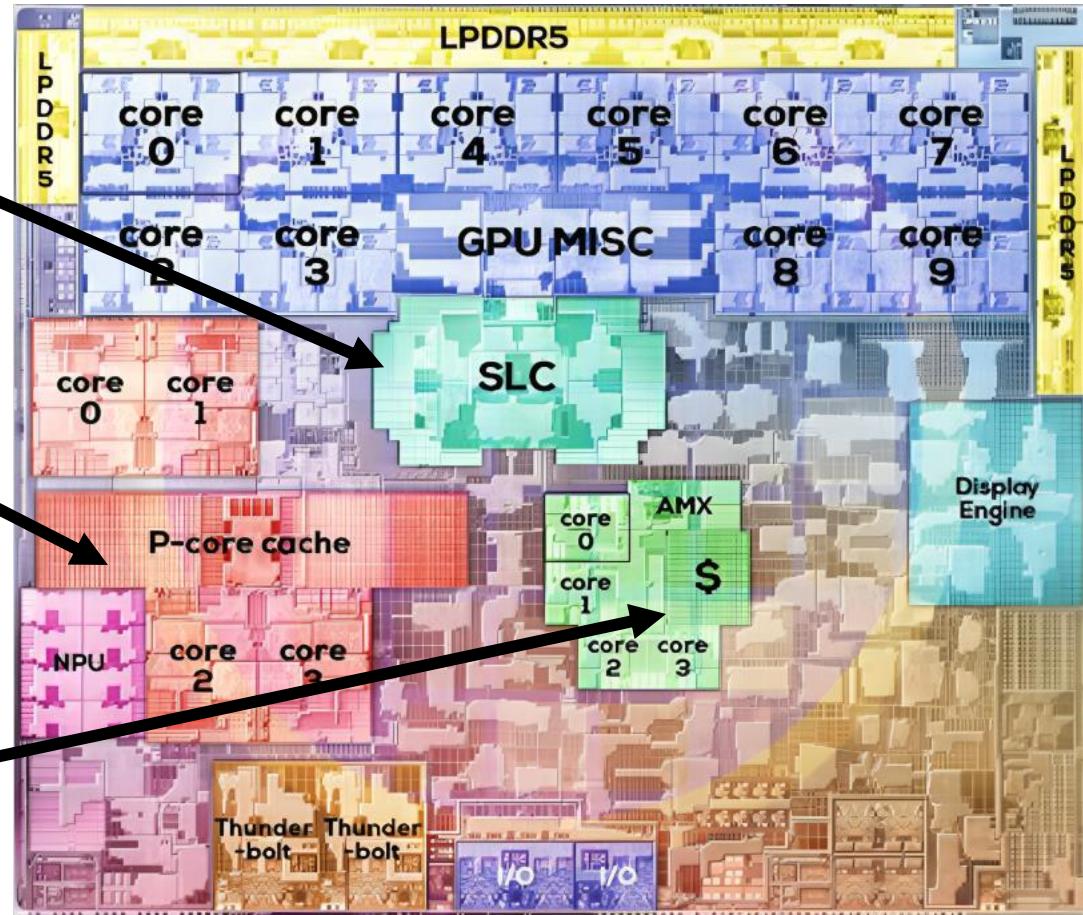
The SoC also has a 8MB System Level Cache (SLC) shared by the GPU and CPU

High-performance cores:

- 192 KB of L1 I-cache
- 128 KB of L1 D-cache
- Shared 16MB L2 cache

Energy-efficient cores:

- 128 KB L1 I-cache
- 64 kB of L1 D-cache
- Shared 4 MB L2 cache



[Image credit: High Yield on Twitter/X]

# Multilevel Cache Efficacy L1

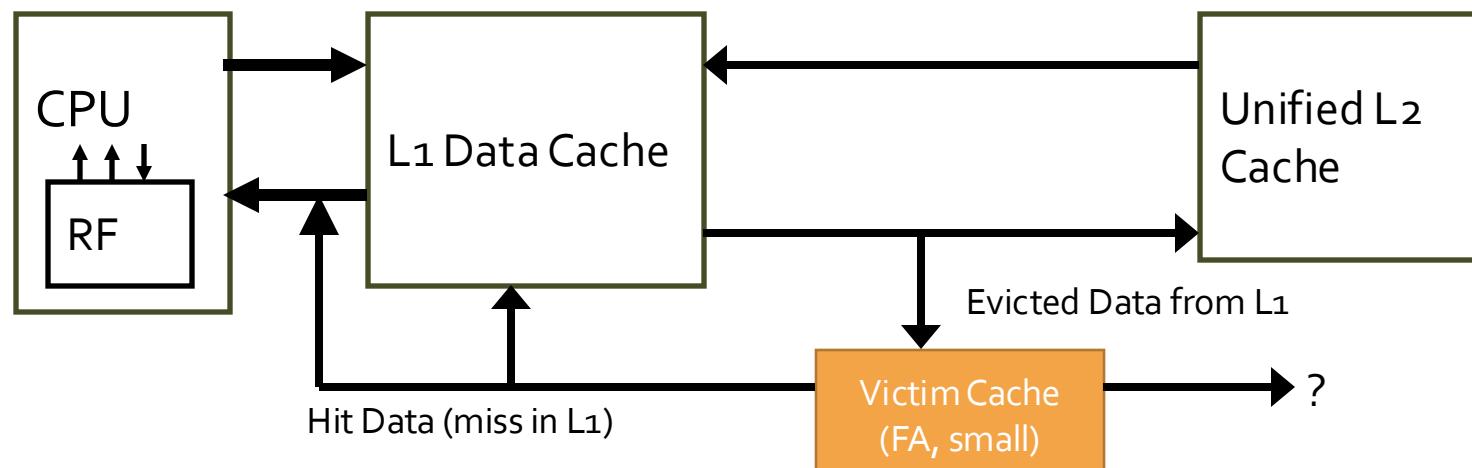
Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Multilevel Cache		+		

# Multilevel Cache Efficacy L<sub>1</sub>, L<sub>2</sub>, L<sub>3</sub>

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Multilevel Cache	+	+		

# Victim Cache

- Small Fully Associative cache for recently evicted lines
  - Usually small (4-16 blocks)
- Reduced conflict misses
  - Higher associativity for small number of lines
- Can be checked in parallel or series with main cache
- On Miss in L<sub>1</sub>, Hit in VC: VC->L<sub>1</sub>, L<sub>1</sub>->VC
- On Miss in L<sub>1</sub>, Miss in VC: L<sub>1</sub>->VC, VC->? (Can always be clean)



# Victim Cache Efficacy L1

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Victim Cache		+		

# Victim Cache Efficacy L1 and VC

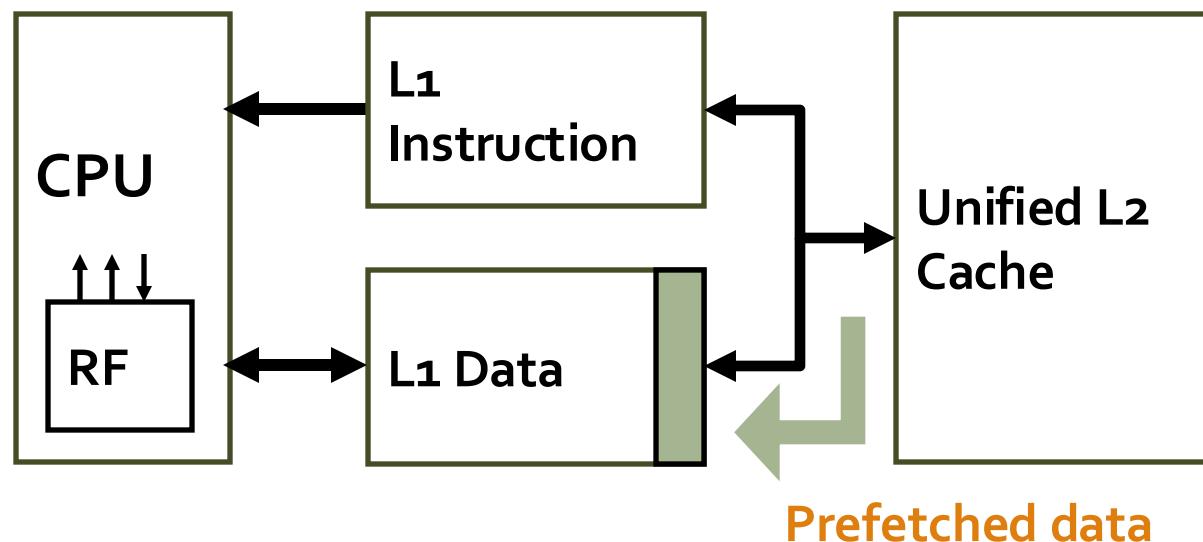
Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Victim Cache	+	+		

# Prefetching

- Speculate on future instruction and data accesses and fetch them into cache(s)
  - Instruction accesses easier to predict than data accesses
- Varieties of prefetching
  - Hardware prefetching
  - Software prefetching
  - Mixed schemes
- **What types of misses does prefetching affect?**

# Issues in Prefetching

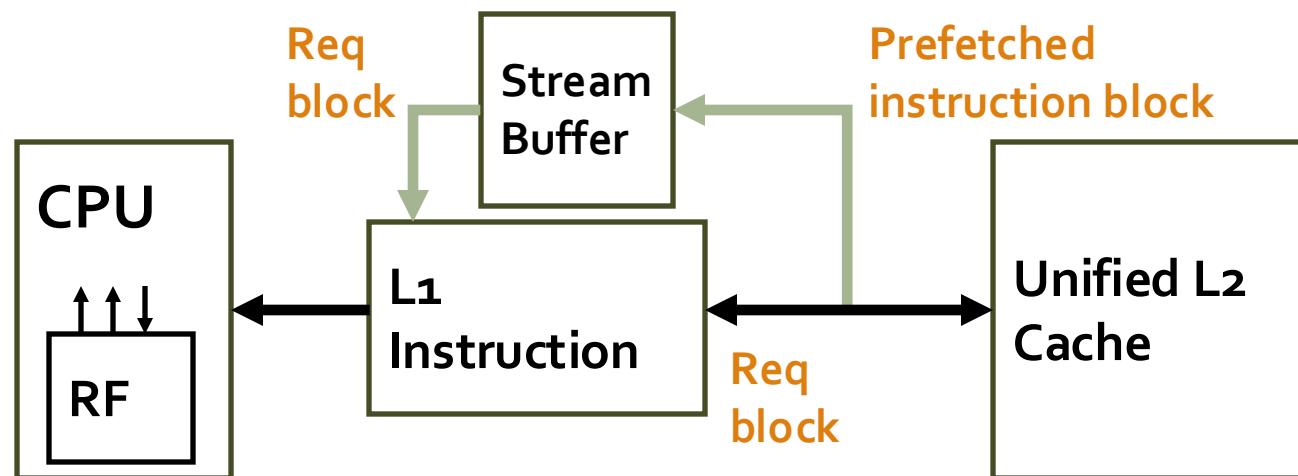
- Usefulness – should produce hits
- Timeliness – not late and not too early
- Cache and bandwidth pollution



# Hardware Instruction Prefetching

## Instruction prefetch in Alpha AXP 21064

- Fetch two blocks on a miss; the requested block (i) and the next consecutive block (i+1)
- Requested block placed in cache, and next block in instruction stream buffer
- If miss in cache but hit in stream buffer, move stream buffer block into cache and prefetch next block (i+2)



# Hardware Data Prefetching

- Prefetch-on-miss:
  - Prefetch  $b + 1$  upon miss on  $b$
- One Block Lookahead (OBL) scheme
  - Initiate prefetch for block  $b + 1$  when block  $b$  is accessed
  - **Why is this different from doubling block size?**
  - Can extend to N-block lookahead
- Strided prefetch
  - If observe sequence of accesses to block  $b$ ,  $b+N$ ,  $b+2N$ , then prefetch  $b+3N$  etc.

**Example:** IBM Power 5 [2003] supports eight independent streams of strided prefetch per processor, prefetching 12 lines ahead of current access

# Software Prefetching

```
for(i=0; i < N; i++) {  
    prefetch( &a[i + 1] );  
    prefetch( &b[i + 1] );  
    SUM = SUM + a[i] * b[i];  
}
```

# Software Prefetching Issues

- Timing is the biggest issue, not predictability
  - If you prefetch very close to when the data is required, you might be too late
  - Prefetch too early, cause pollution
  - Estimate how long it will take for the data to come into L1, so we can set P appropriately
  - **Why is this hard to do?**

```
for(i=0; i < N; i++) {  
    prefetch( &a[i + P] );  
    prefetch( &b[i + P] );  
    SUM = SUM + a[i] * b[i];  
}
```

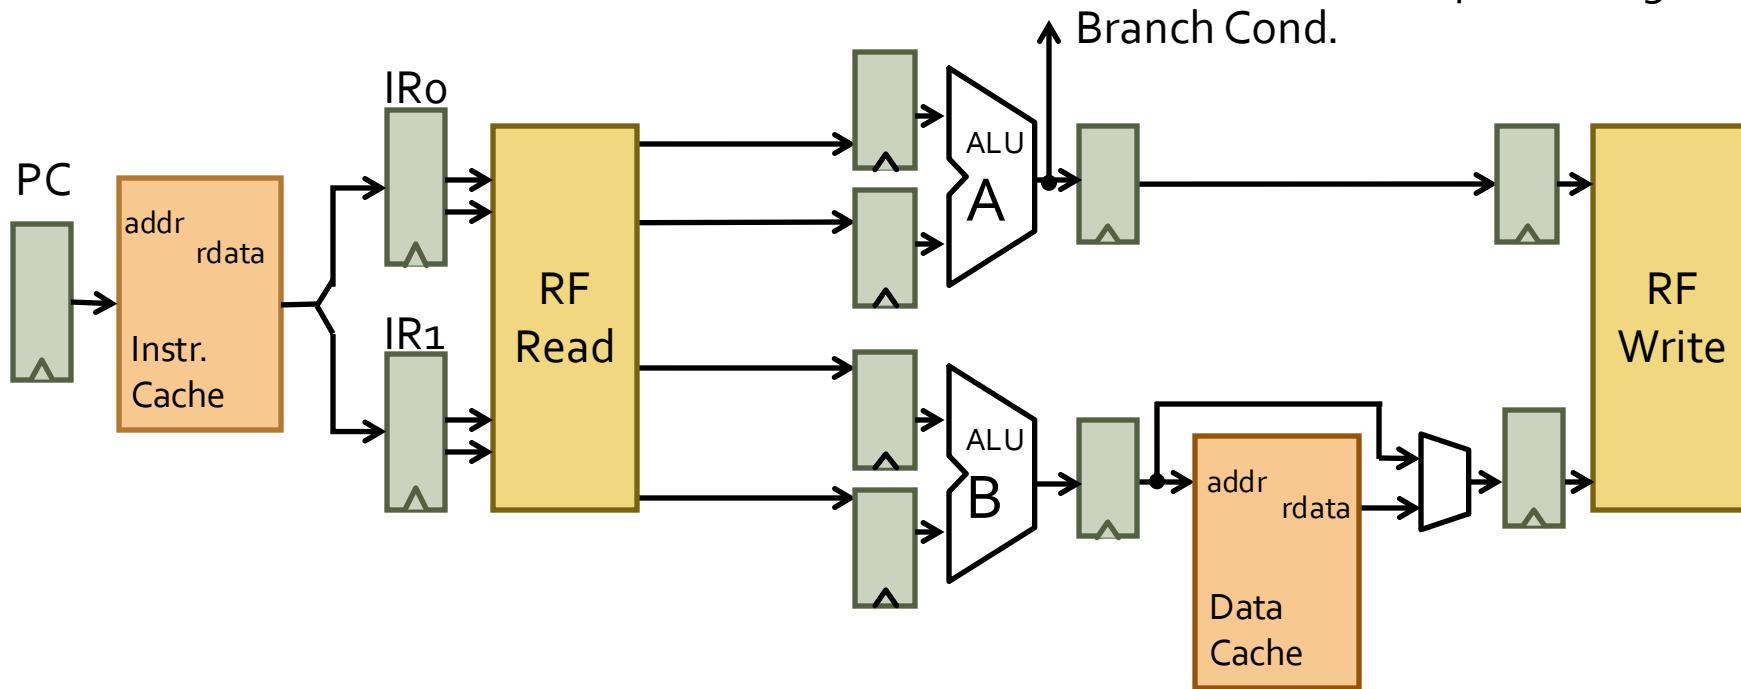
Must consider cost of prefetch instructions

# Prefetching Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Prefetching	+	+		

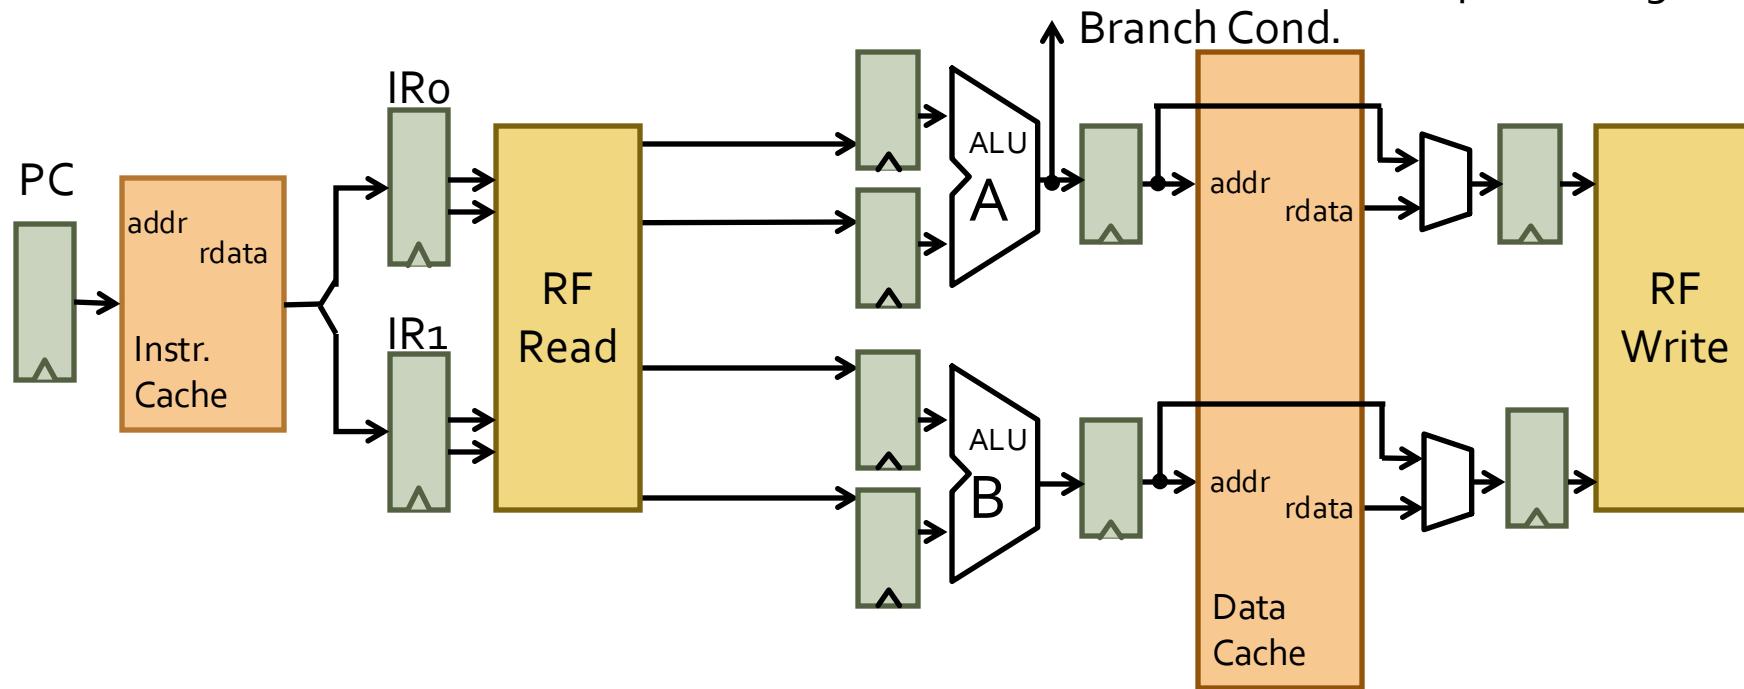
# Increasing Cache Bandwidth Multiporting and Banking

Pipe A: Integer Ops., Branches  
Pipe B: Integer Ops., Memory



# Increasing Cache Bandwidth Multiporting and Banking

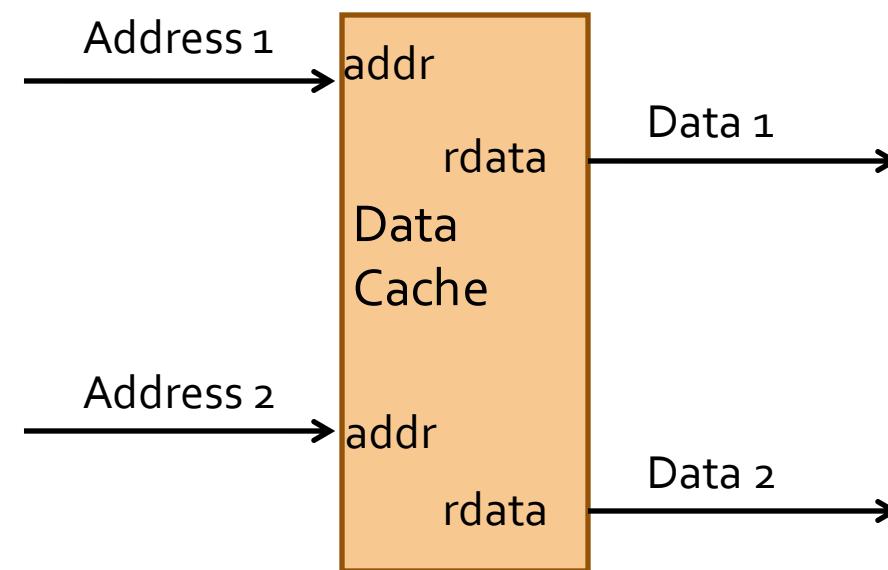
Pipe A: Integer Ops., Branches  
Pipe B: Integer Ops., Memory



**Challenge:** Two stores to the same line, or Load and Store to same line

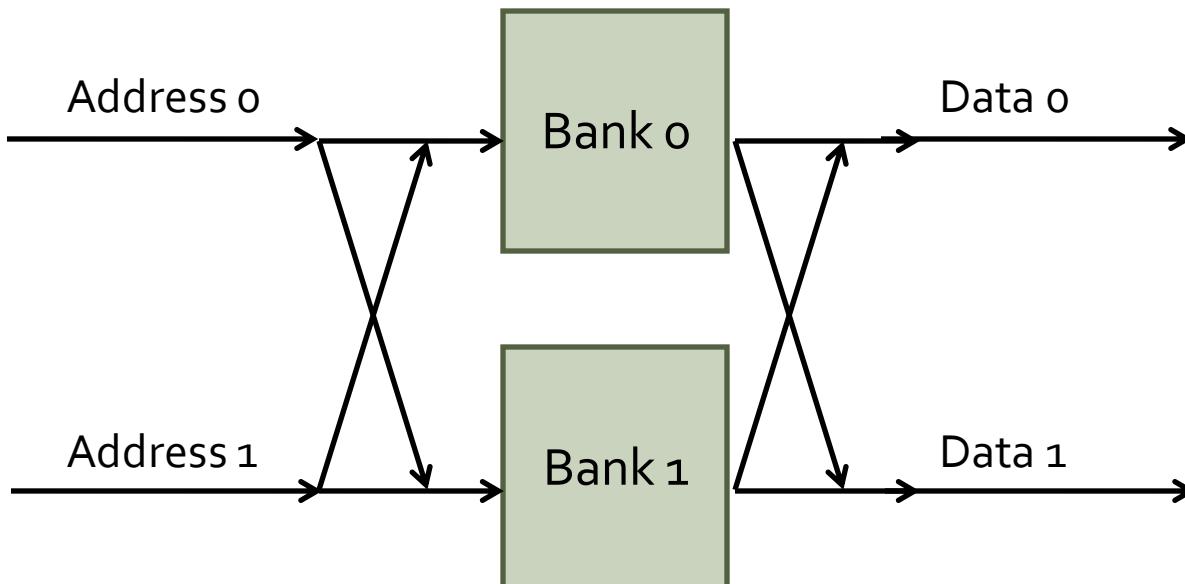
# True Multiport Caches

- Large area increase (could be double for 2-port)
- Hit time increase (can be made small)



# Banked Caches

- Partition address space into multiple banks
  - Use portions of address (low or high order interleaved)
- Benefits:
  - Higher throughput
- Challenges:
  - Bank conflicts
  - Extra wiring



# Cache Banking Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Cache Banking				+

# Compiler Optimizations

- Restructuring code affects the data block access sequence
  - Group data accesses together to improve spatial locality
  - Re-order data accesses to improve temporal locality
- Prevent data from entering the cache
  - Useful for variables that will only be accessed once before being replaced
  - Needs mechanism for software to tell hardware not to cache data (“no-allocate” instruction hints or page table bits)
- Kill data that will never be used again
  - Streaming data exploits spatial locality but not temporal locality
  - Replace into dead cache locations

# Loop Interchange

```
for(j=0; j < N; j++) {  
    for(i=0; i < M; i++) {  
        x[i][j] = 2 * x[i][j];  
    }  
}
```



```
for(i=0; i < M; i++) {  
    for(j=0; j < N; j++) {  
        x[i][j] = 2 * x[i][j];  
    }  
}
```

What type of locality does this improve?

# Loop Fusion

```
for(i=0; i < N; i++)
    a[i] = b[i] * c[i];
```

```
for(i=0; i < N; i++)
    d[i] = a[i] * c[i];
```

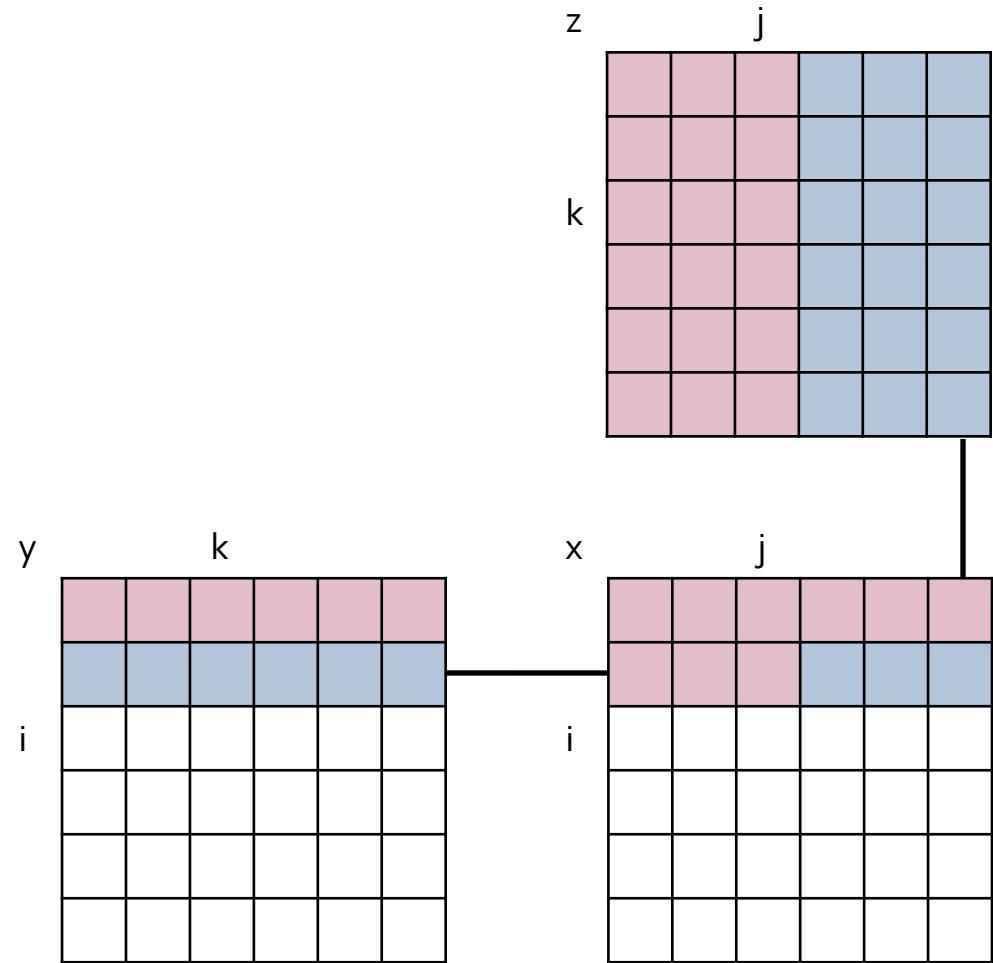


```
for(i=0; i < N; i++)
{
    a[i] = b[i] * c[i];
    d[i] = a[i] * c[i];
}
```

What type of locality does this improve?

# Matrix Multiply, Naïve Code

```
for(i=0; i < N; i++)  
    for(j=0; j < N; j++) {  
        r = 0;  
        for(k=0; k < N; k++)  
            r = r + y[i][k] * z[k][j];  
        x[i][j] = r;  
    }
```



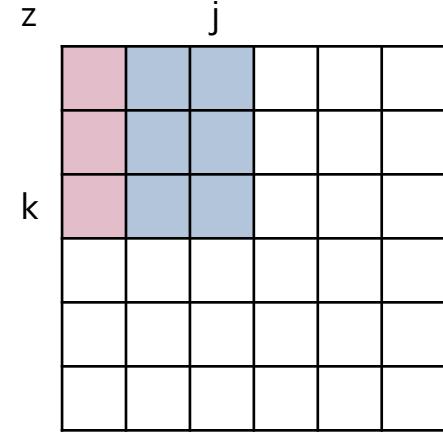
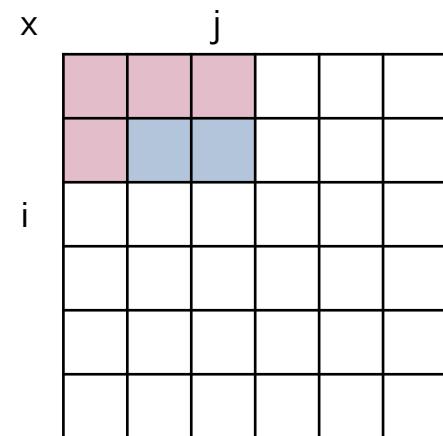
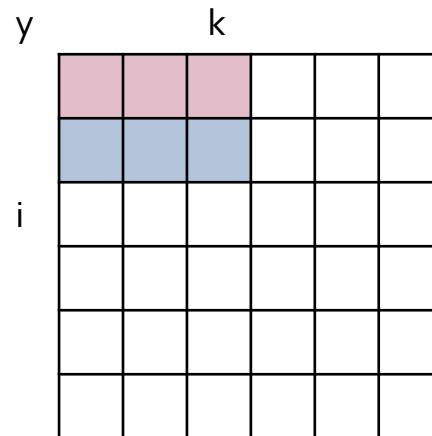
Not touched

Old access

New access

# Matrix Multiply with Cache Tiling/Blocking

```
for(jj=0; jj < N; jj=jj+B)
    for(kk=0; kk < N; kk=kk+B)
        for(i=0; i < N; i++)
            for(j=jj; j < min(jj+B,N); j++) {
                r = 0;
                for(k=kk; k < min(kk+B,N); k++)
                    r = r + y[i][k] * z[k][j];
                x[i][j] = x[i][j] + r;
            }
```



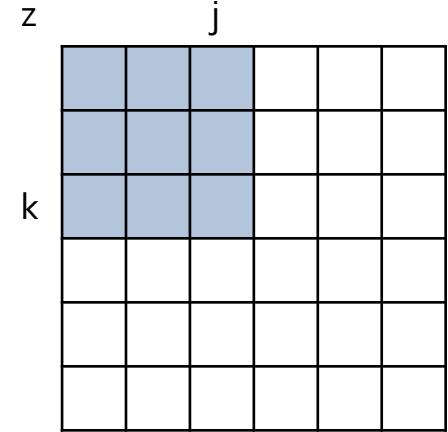
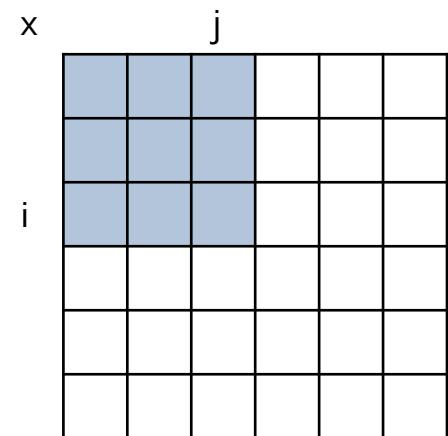
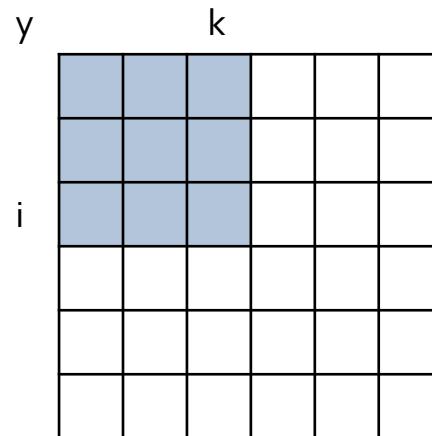
Not touched

Old access

New access

# Matrix Multiply with Cache Tiling/Blocking

```
for(jj=0; jj < N; jj=jj+B)
    for(kk=0; kk < N; kk=kk+B)
        for(i=0; i < N; i++)
            for(j=jj; j < min(jj+B,N); j++) {
                r = 0;
                for(k=kk; k < min(kk+B,N); k++)
                    r = r + y[i][k] * z[k][j];
                x[i][j] = x[i][j] + r;
            }
```



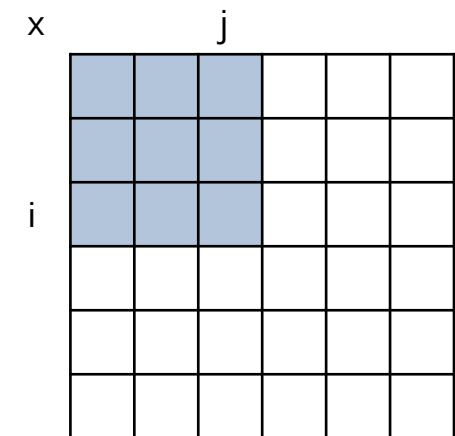
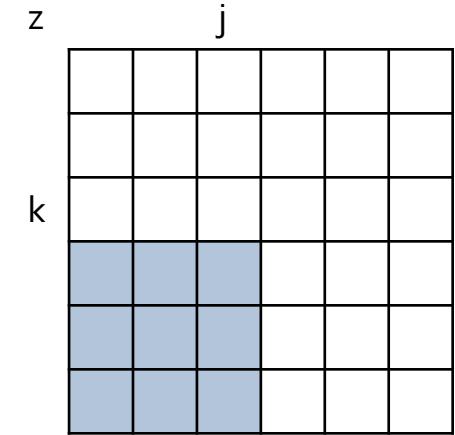
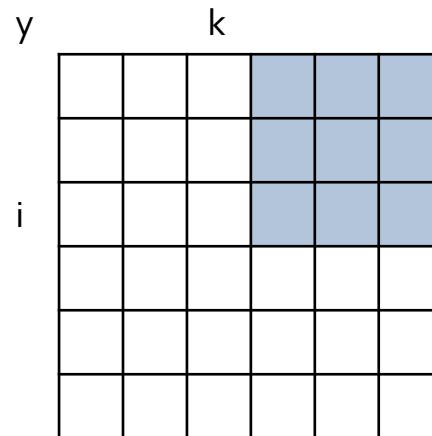
Not touched

Old access

New access

# Matrix Multiply with Cache Tiling/Blocking

```
for(jj=0; jj < N; jj=jj+B)
    for(kk=0; kk < N; kk=kk+B)
        for(i=0; i < N; i++)
            for(j=jj; j < min(jj+B,N); j++) {
                r = 0;
                for(k=kk; k < min(kk+B,N); k++)
                    r = r + y[i][k] * z[k][j];
                x[i][j] = x[i][j] + r;
            }
```



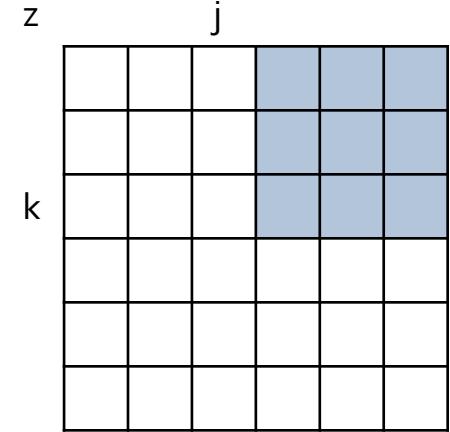
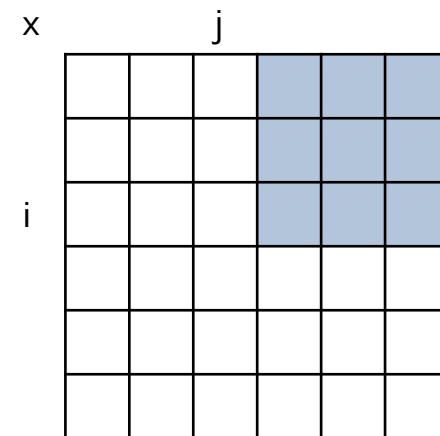
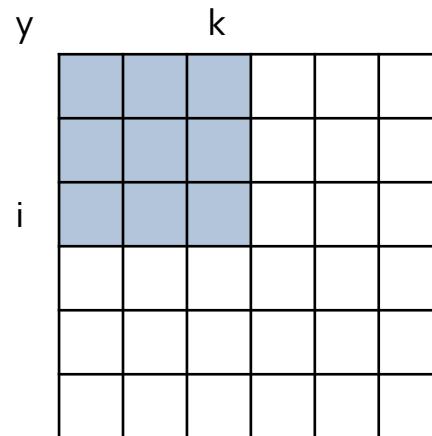
Not touched

Old access

New access

# Matrix Multiply with Cache Tiling/Blocking

```
for(jj=0; jj < N; jj=jj+B)
    for(kk=0; kk < N; kk=kk+B)
        for(i=0; i < N; i++)
            for(j=jj; j < min(jj+B,N); j++) {
                r = 0;
                for(k=kk; k < min(kk+B,N); k++)
                    r = r + y[i][k] * z[k][j];
                x[i][j] = x[i][j] + r;
            }
```



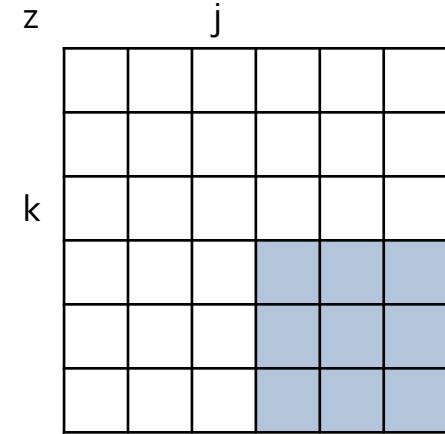
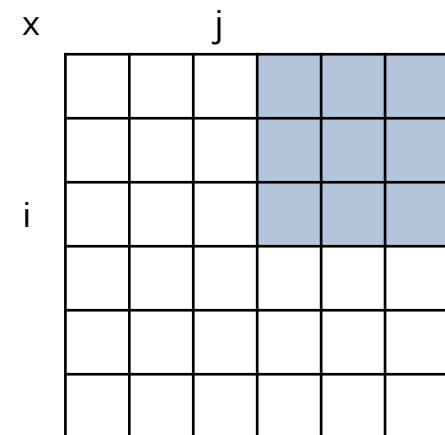
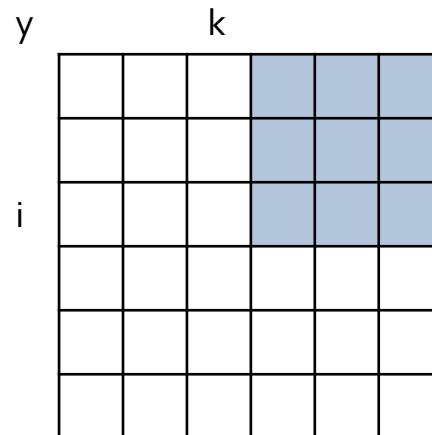
Not touched

Old access

New access

# Matrix Multiply with Cache Tiling/Blocking

```
for(jj=0; jj < N; jj=jj+B)
    for(kk=0; kk < N; kk=kk+B)
        for(i=0; i < N; i++)
            for(j=jj; j < min(jj+B,N); j++) {
                r = 0;
                for(k=kk; k < min(kk+B,N); k++)
                    r = r + y[i][k] * z[k][j];
                x[i][j] = x[i][j] + r;
            }
```



Not touched

Old access

New access

# Compiler Memory Optimizations Efficacy

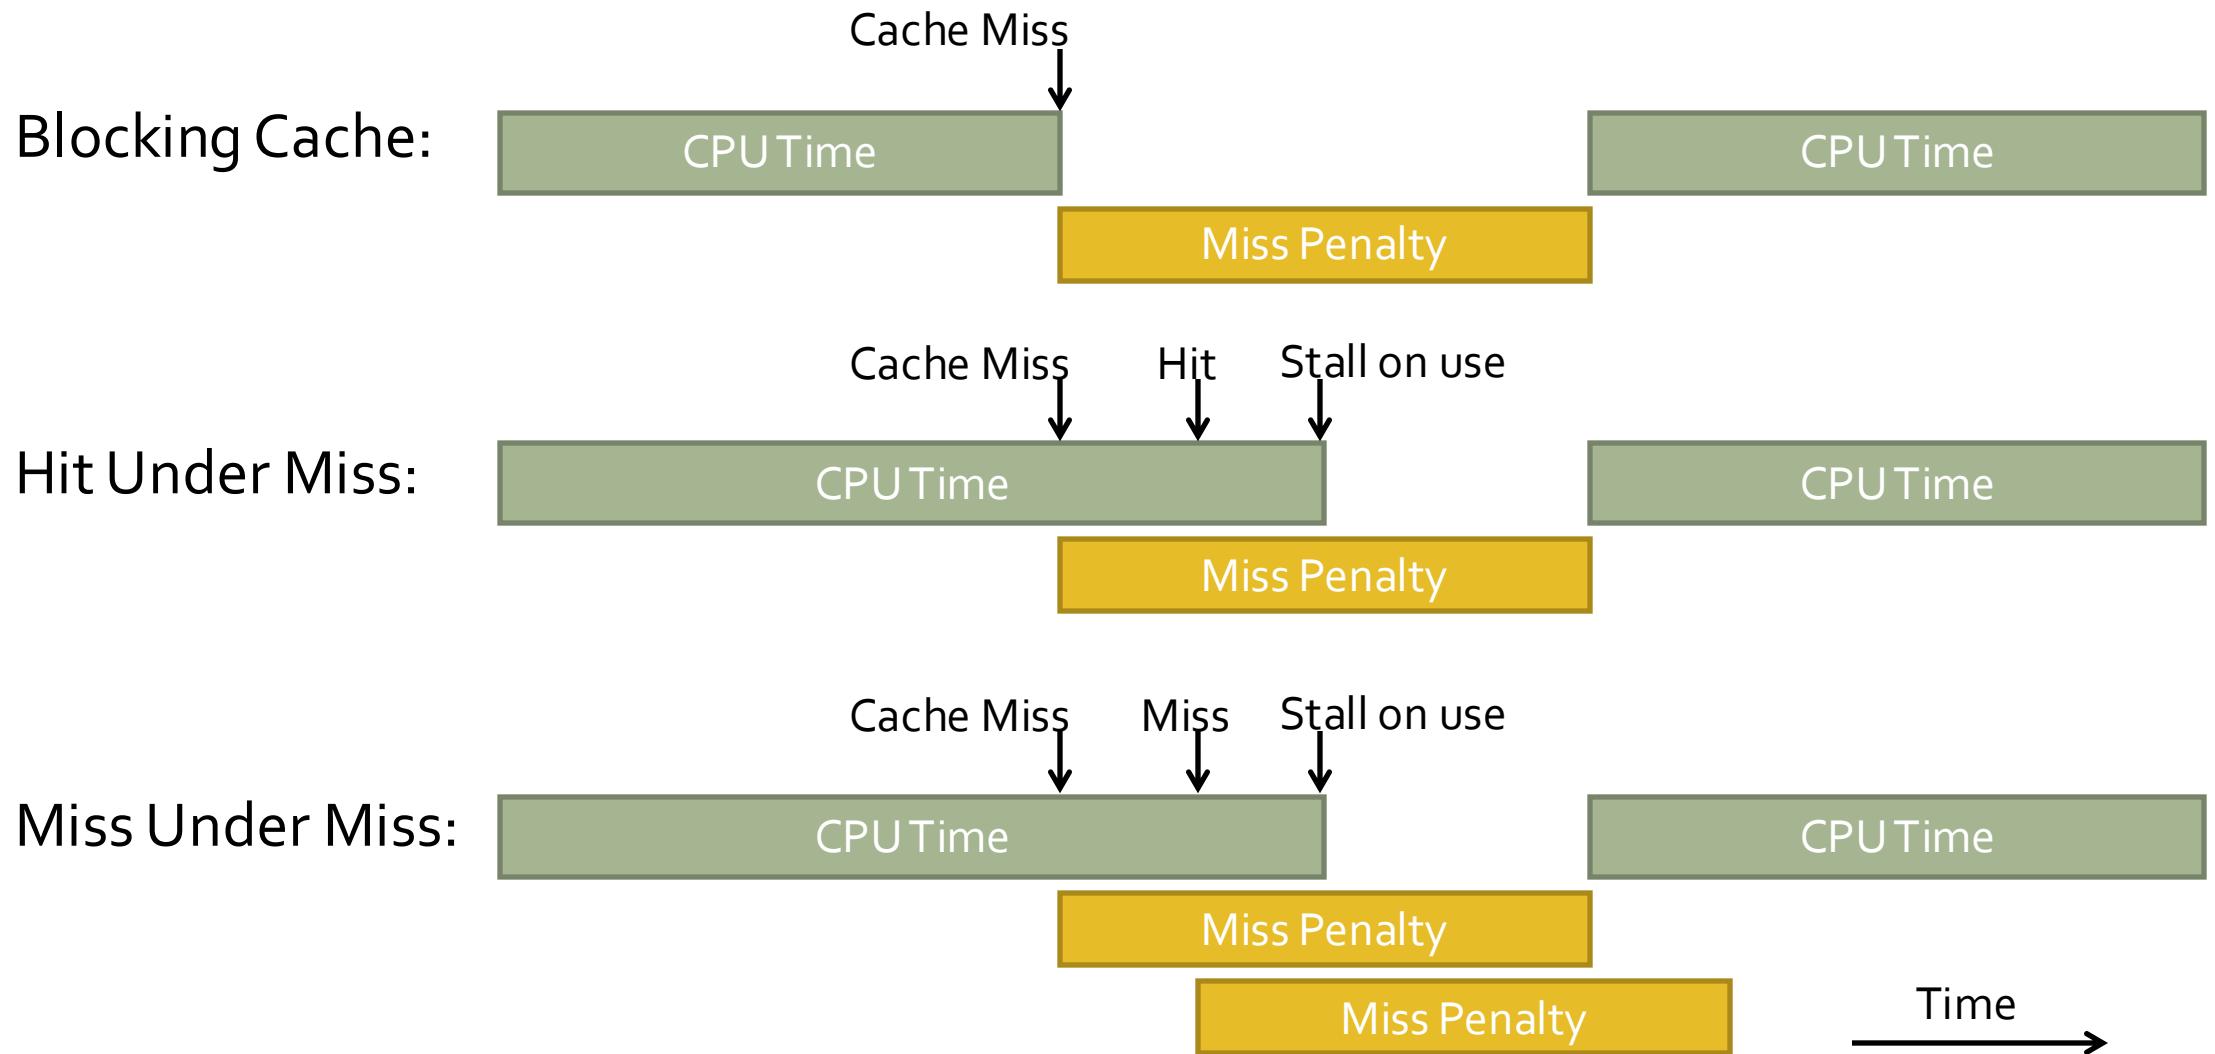
Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Compiler Optimization	+			

# Non-Blocking Caches

(aka Out-Of-Order Memory System)  
(aka Lockup Free Caches)

- Enable subsequent cache accesses after a cache miss has occurred
  - Hit-under-miss
  - Miss-under-miss (concurrent misses)
- Suitable for in-order processor or out-of-order processors
  - **in-order processors** stall when an instruction that uses the load data is the next instruction to be executed
  - **out-of-order processors** can execute instructions after the load consumer
- Challenges
  - Maintaining order when multiple misses that might return out of order
  - Load or Store to an already pending miss address (need merge)

# Non-Blocking Cache Timeline



# Miss Status Handling Register (MSHR)/ Miss Address File (MAF)

MSHR/MAF		
V	Block Address	Issued

**V:** Valid

**Block Address:** Address of cache block in memory system

**Issued:** Issued to Main Memory/Next level of cache

Load/Store Entry				
V	MSHR Entry	Type	Offset	Destination

**V:** Valid

**MSHR Entry:** Entry Number

**Type:** {LW, SW, LH, SH, LB, SB}

**Offset:** Offset within the block

**Destination:** (Loads) Register, (Stores) Store buffer entry

# Non-Blocking Cache Operation

On Cache Miss:

- Check MSHR for matched address
  - If found: allocate new Load/Store entry pointing to MSHR
  - If not found: allocate new MSHR entry and Load/Store entry
  - If all entries full in MSHR or Load/Store entry table, stall or prevent new LDs/STs

On Data Return from Memory:

- Find Load or Store waiting for it
  - Forward Load data to processor/Clear Store Buffer
  - Could be multiple Loads and Stores
- Write Data to cache

When Cache Lines is Completely Returned:

- De-allocate MSHR entry

# Non-Blocking Cache with In-Order Pipelines

- Need Scoreboard for Individual Registers

On Load Miss:

- Mark Destination Register as Busy

On Load Data Return:

- Mark Destination Register as Available

On Use of Busy Register:

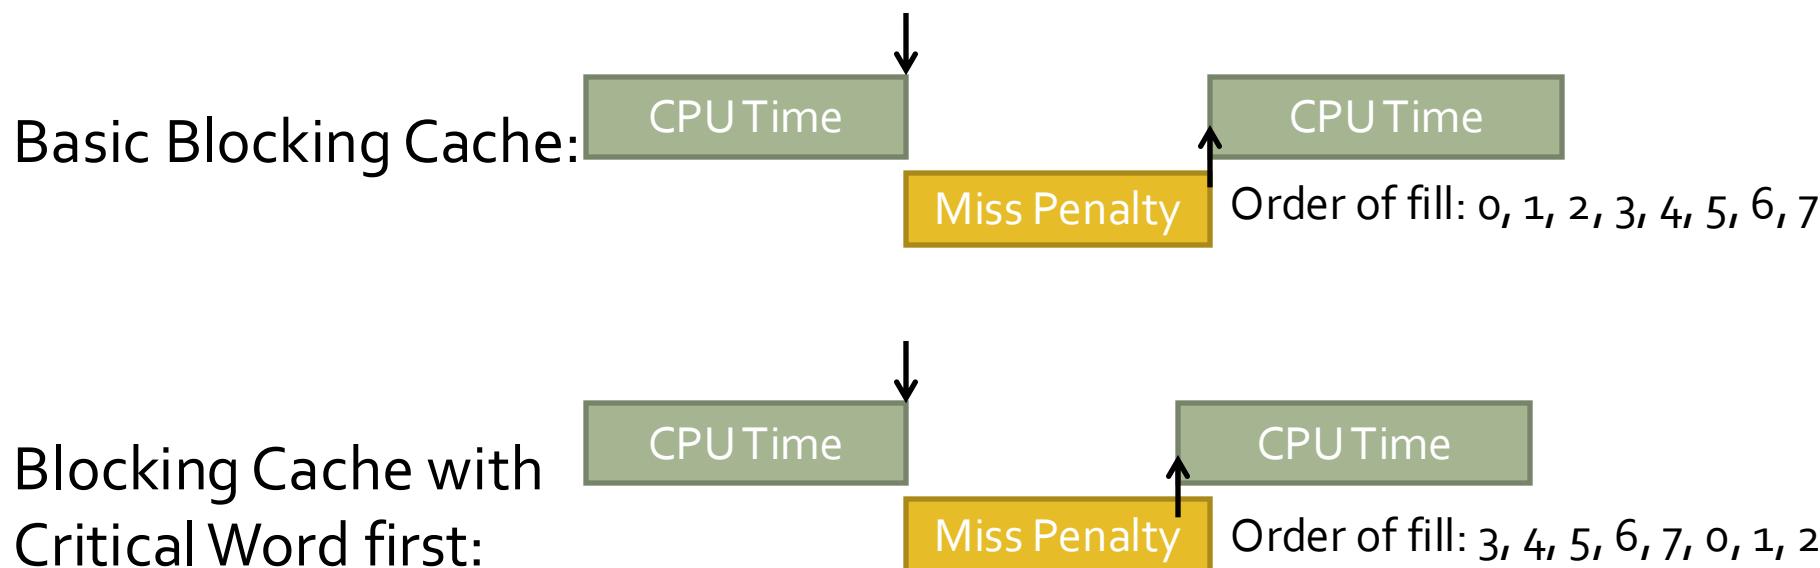
- Stall Processor

# Non-Blocking Cache Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Non-blocking Cache		+		+

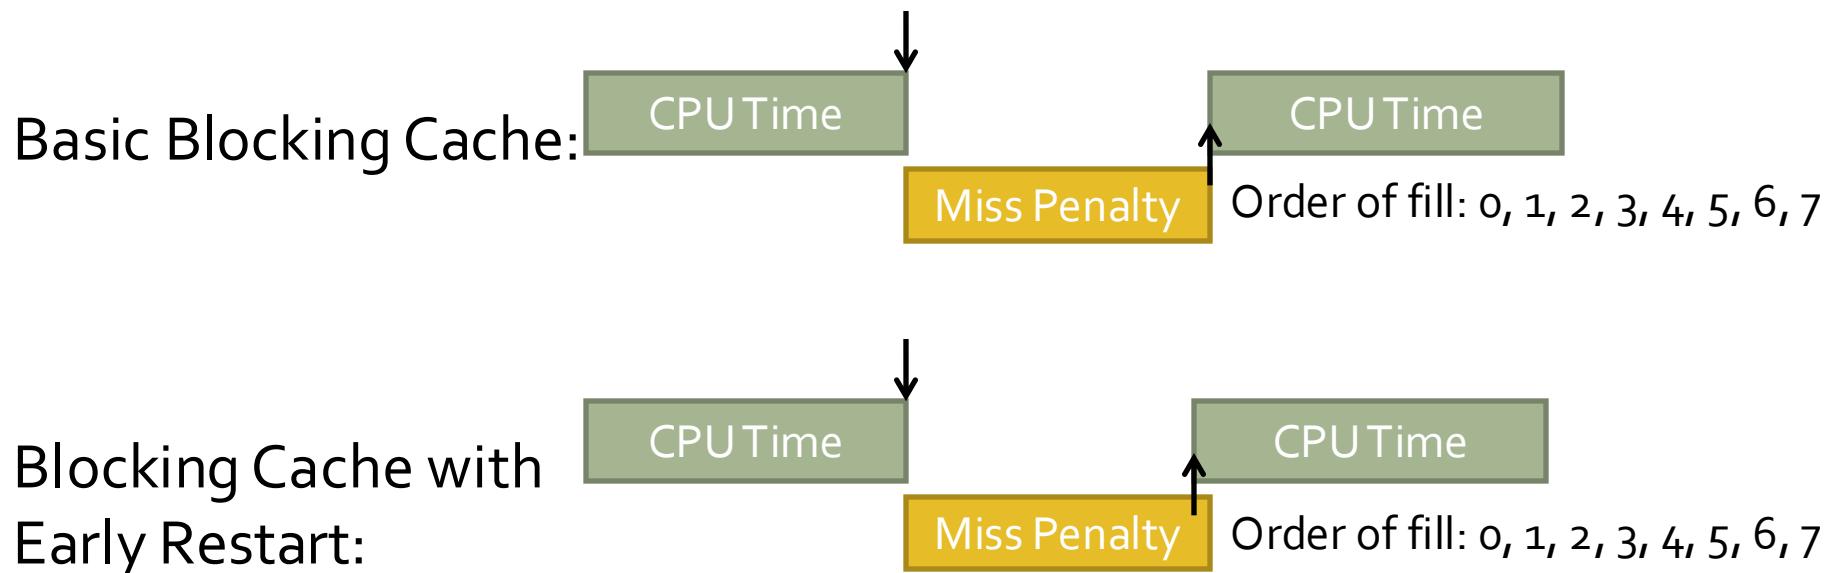
# Critical Word First

- Request the missed word from memory first.
- Rest of cache line comes after “critical word”
  - Commonly words come back in rotated order



# Early Restart

- Data returns from memory in order
- Processor Restarts when needed word is returned



# Critical Word First and Early Restart Efficacy

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Critical Word First/Early Restart		+		

# Summary: Cache Optimizations

Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Pipelined & banked caches			-	+
Writer Buffer		+		
Multilevel Cache	+	+		
Victim Cache	+	+		
Prefetching	+	+		
Compiler Optimization	+			
Non-blocking Cache		+		+
Critical Word First/Early Restart		+		

# Summary: Cache Optimizations

Technique	Hit time	Band-width	Miss penalty	Miss rate	Power consumption	Hardware cost/complexity	Comment
Small and simple caches	+			-	+	0	Trivial; widely used
Way-predicting caches	+				+	1	Used in Pentium 4
Pipelined & banked caches	-	+				1	Widely used
Nonblocking caches	+	+				3	Widely used
Critical word first and early restart			+			2	Widely used
Merging write buffer			+			1	Widely used with write through
Compiler techniques to reduce cache misses				+		0	Software is a challenge, but many compilers handle common linear algebra calculations
Hardware prefetching of instructions and data		+	+	-	2 instr., 3 data		Most provide prefetch instructions; modern high-end processors also automatically prefetch in hardware
Compiler-controlled prefetching		+	+			3	Needs nonblocking cache; possible instruction overhead; in many CPUs
HBM as additional level of cache	+/-	-	+	+		3	Depends on new packaging technology. Effects depend heavily on hit rate improvements

# Acknowledgements

- These slides contain material developed and copyright by:
  - Arvind (MIT)
  - Krste Asanovic (MIT/UCB)
  - Joel Emer (Intel/MIT)
  - James Hoe (CMU)
  - John Kubiatowicz (UCB)
  - David Patterson (UCB)
  - Christopher Batten (Cornell)
  - David Wentzlaff (Princeton)
- MIT material derived from course 6.823
- UCB material derived from course CS252
- Cornell material derived from course ECE 4750
- Princeton material derived from course ECE 475