

# **Computer Architecture**

## **Memory Management**

Ting-Jung Chang

NYCU CS

# Course Administration

- Lab<sub>4</sub> and PS<sub>4</sub> released on e3

# Recap: Cache Optimizations

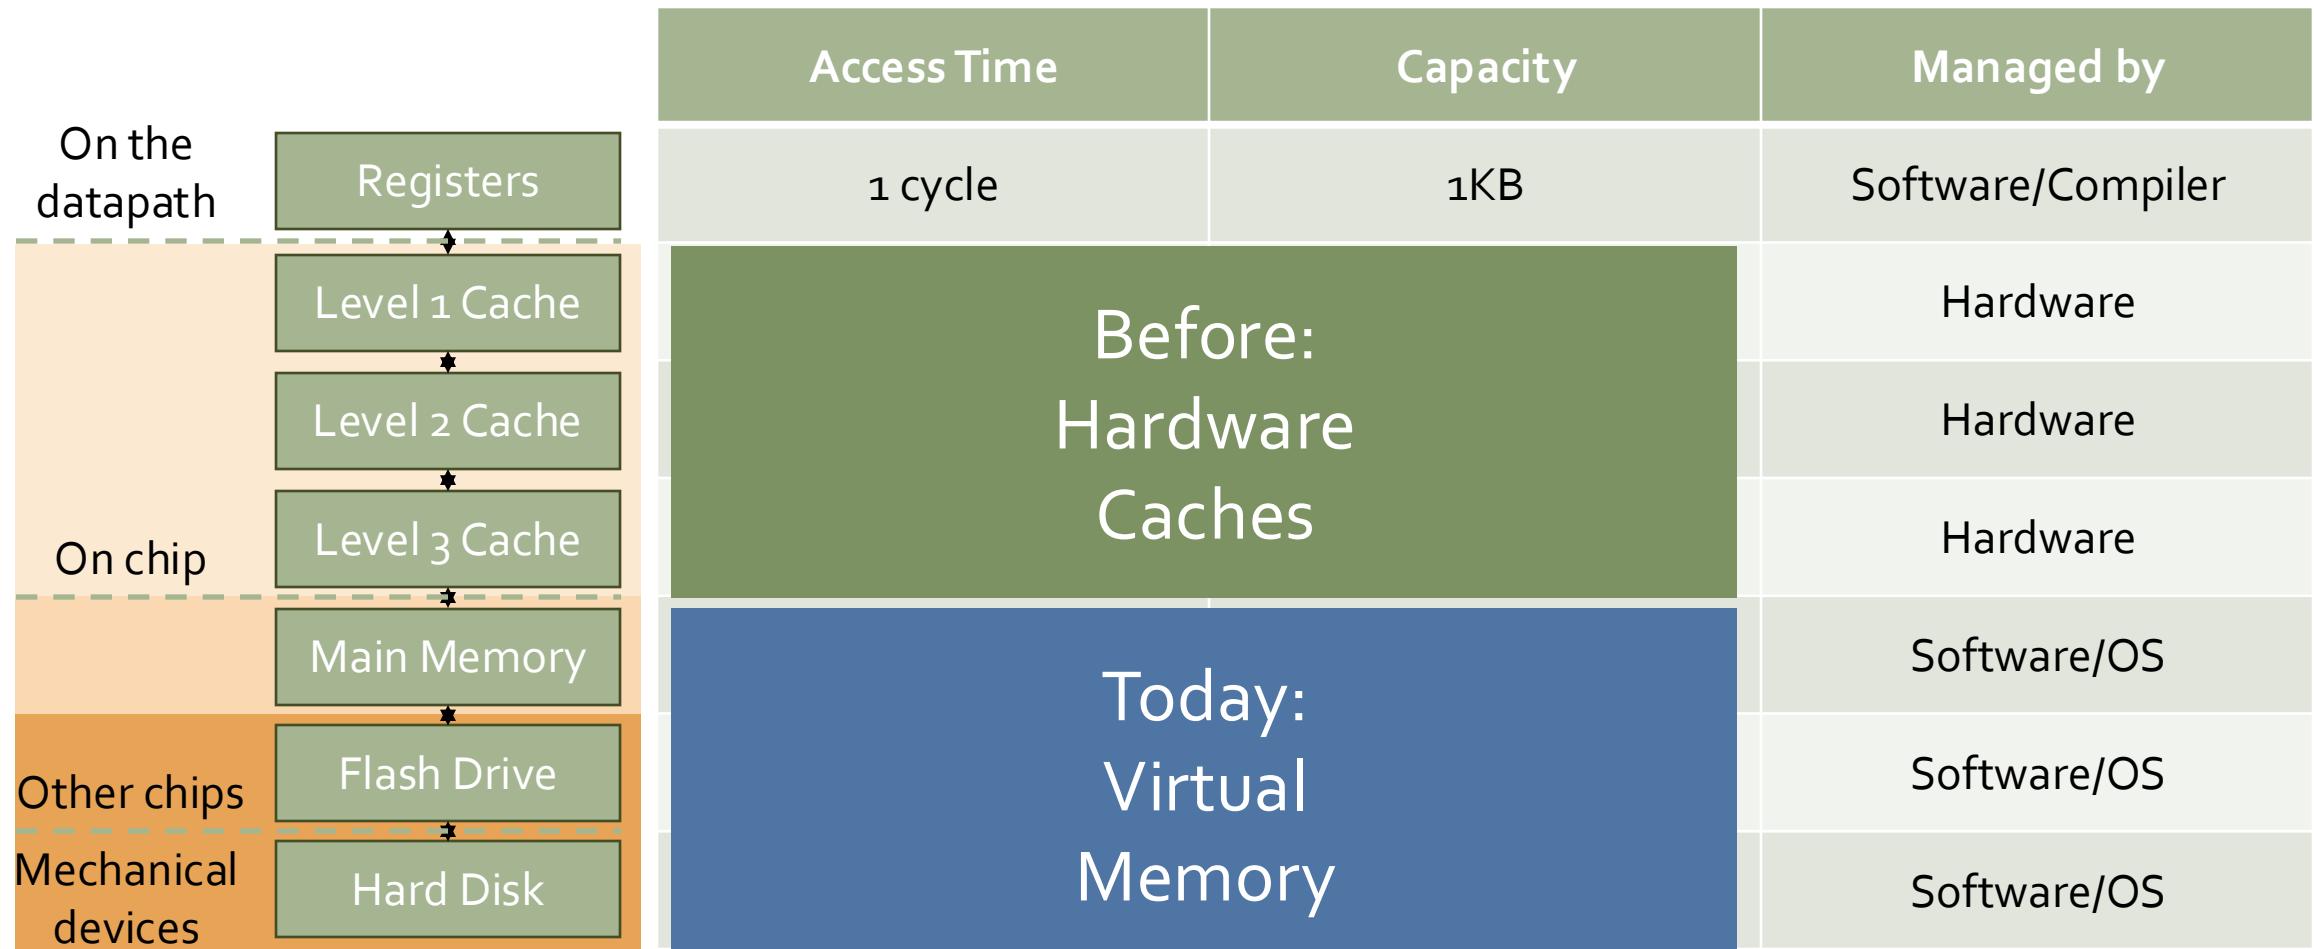
Cache Optimization	Miss Rate	Miss Penalty	Hit Time	Bandwidth
Pipelined & banked caches			-	+
Writer Buffer		+		
Multilevel Cache	+	+		
Victim Cache	+	+		
Prefetching	+	+		
Compiler Optimization	+			
Non-blocking Cache		+		+
Critical Word First/Early Restart		+		

# Memory Hierarchy

	Access Time	Capacity	Managed by
On the datapath	Registers	1 cycle	1KB Software/Compiler
On chip	Level 1 Cache	*	Hardware
	Level 2 Cache	*	Hardware
	Level 3 Cache	*	Hardware
	Main Memory	*	Software/OS
Other chips	Flash Drive	*	Software/OS
Mechanical devices	Hard Disk	*	Software/OS

**Before:**  
**Hardware Caches**

**Today:**  
**Virtual Memory**



# Extending the Memory Hierarchy

- Problem: DRAM vs disk has much more extreme differences than SRAM vs DRAM
  - Access latencies:
    - DRAM ~10-100x slower than SRAM
    - Disk **~100,000x** slower than DRAM
  - Importance of sequential accesses
    - DRAM: Fetching successive words ~5x faster than first word
    - Disk: Fetching successive words **~100,000x** faster than first word
- Result: Design decisions driven by enormous cost of misses
  - Associativity: High, minimize miss rate
  - Block size: Large, amortize cost of a miss over multiple words
  - Write policy: Write back, minimize number of writes

# Memory Management

- From early absolute addressing schemes, to modern virtual memory systems with support for virtual machine monitors
- Can separate into orthogonal functions:
  - Translation (mapping of virtual address to physical address)
  - Protection (permission to access word in memory)
  - Virtual Memory (transparent extension of memory space using slower disk storage)
- But most modern systems provide support for all the above functions with a single page-based system

# Agenda

- Evolution of Virtual Memory
- Modern VM Implementation
- TLB & Cache Organization
- Modern Usage

# Memory Management

- The Fifties
  - Absolute Addresses
  - Dynamic address translation
- The Sixties
  - Atlas and Demand Paging
  - Paged memory systems and TLBs
- Modern Virtual Memory Systems

# Names for Memory Locations



- Machine language address
  - as specified in machine code
- Virtual address
  - ISA specifies translation of machine code address into virtual address of program variable (sometimes called effective address)
- Physical address
  - Operating system specifies mapping of virtual address into name for a physical memory location

# Absolute Addresses

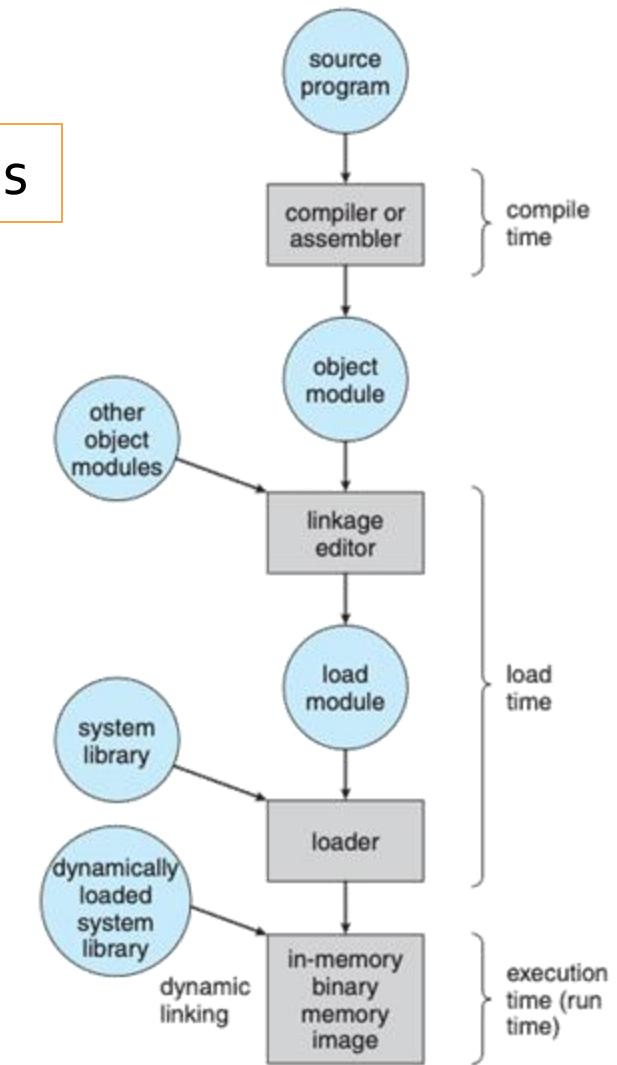
EDSAC, early 50's

virtual address = physical memory address

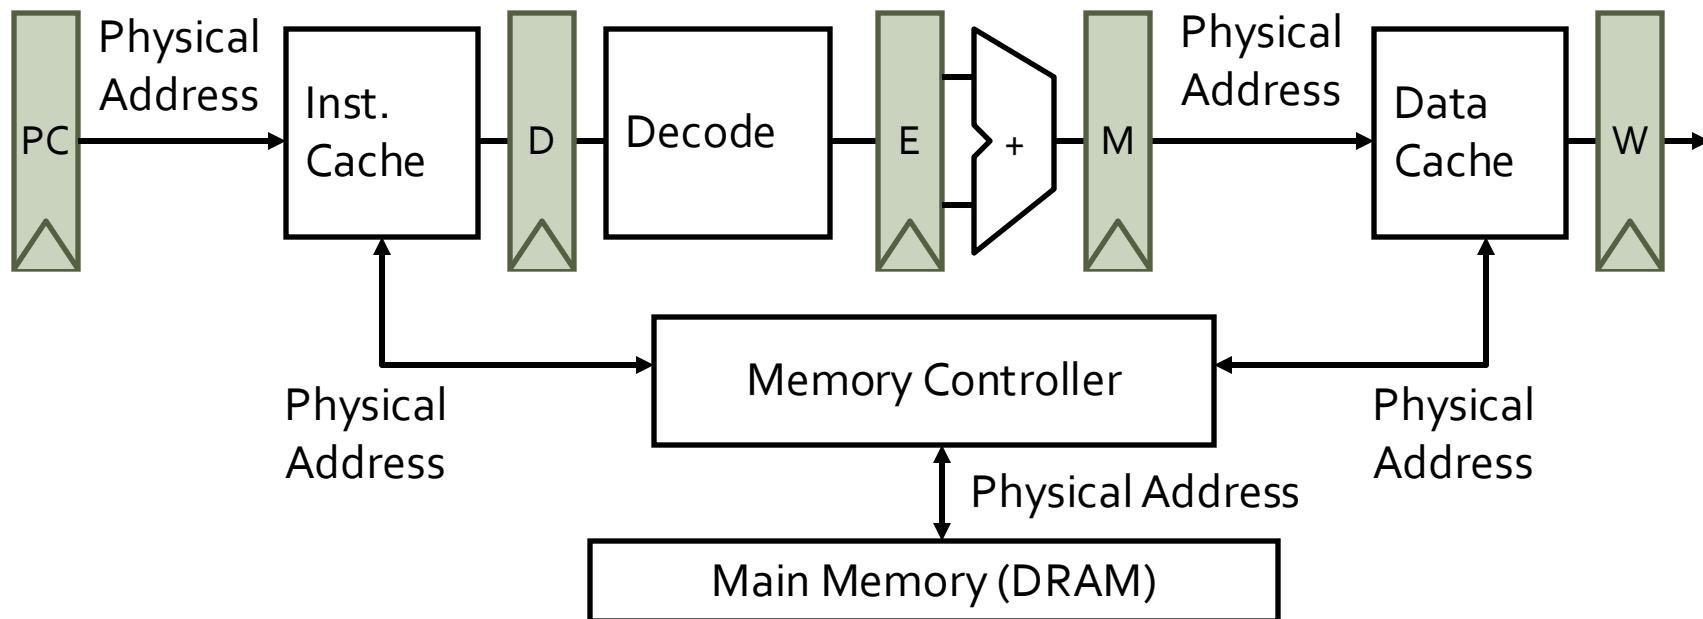
- Only one program ran at a time, with unrestricted access to entire machine (RAM + I/O devices)
- Addresses in a program depended upon where the program was to be loaded in memory
- **But** it was more convenient for programmers to write location-independent subroutines
  - Different programs use different combination of routines

How could location independence be achieved?

Linker and/or loader modify addresses of subroutines and callers when building a program memory image



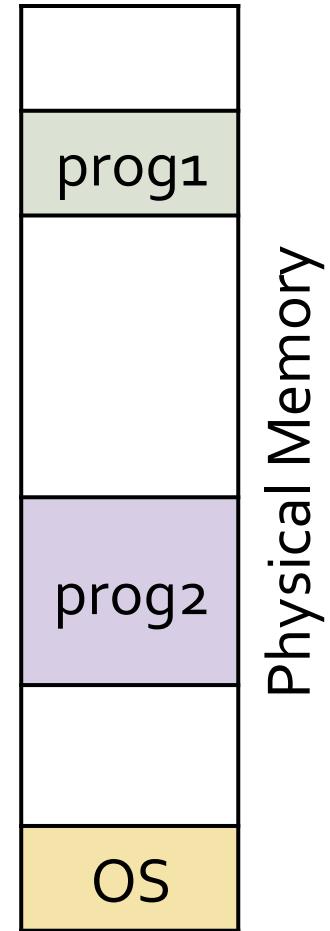
# Bare Machine



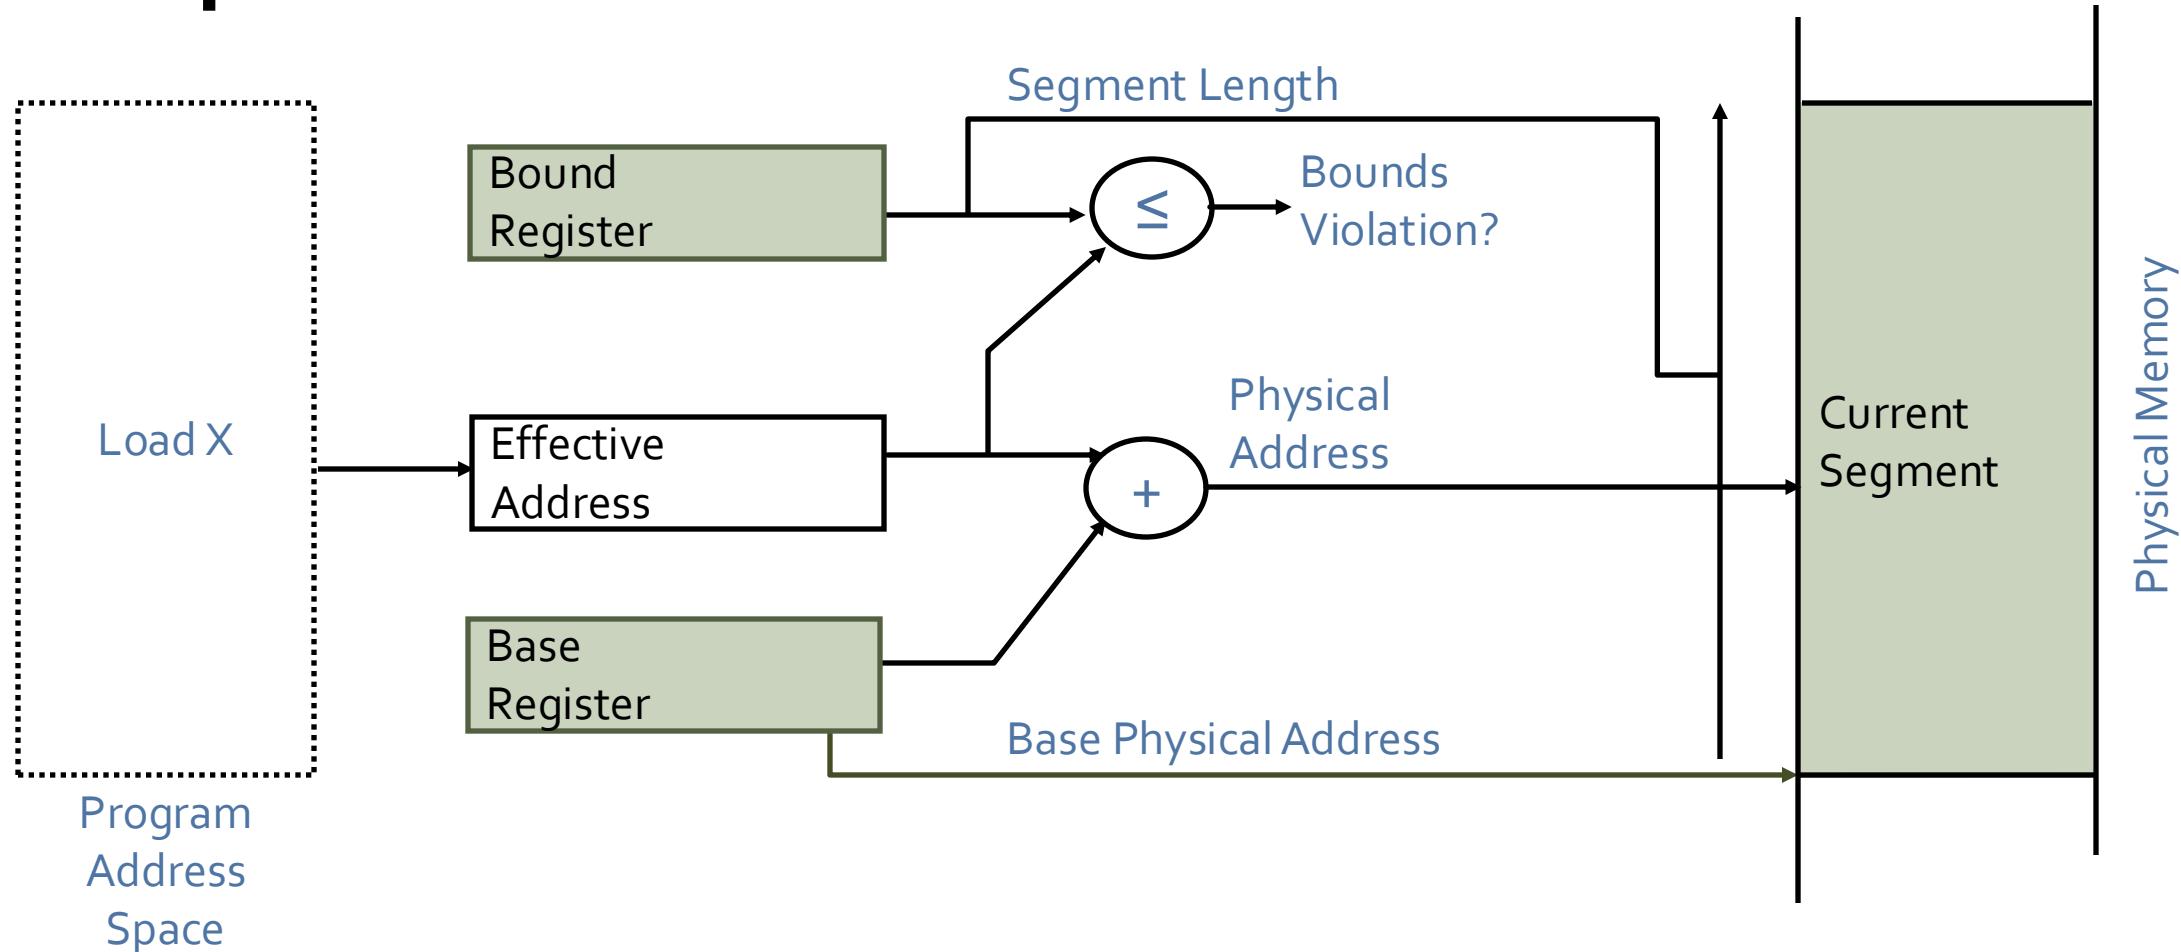
- In a bare machine, the only kind of address is a physical address

# Dynamic Address Translation

- Motivation
  - In early machines, I/O was slow and each I/O transfer involved the CPU (programmed I/O)
  - Higher throughput possible if CPU and I/O of 2 or more programs were overlapped, how?  
    ⇒ multiprogramming with DMA I/O devices, interrupts
- Location-independent programs
  - Programming and storage management ease  
    ⇒ need for a **base** register
- Protection
  - Independent programs should not affect each other inadvertently  
    ⇒ need for a **bound** register
- Multiprogramming drives requirement for resident **supervisor** software to manage context switches between multiple programs

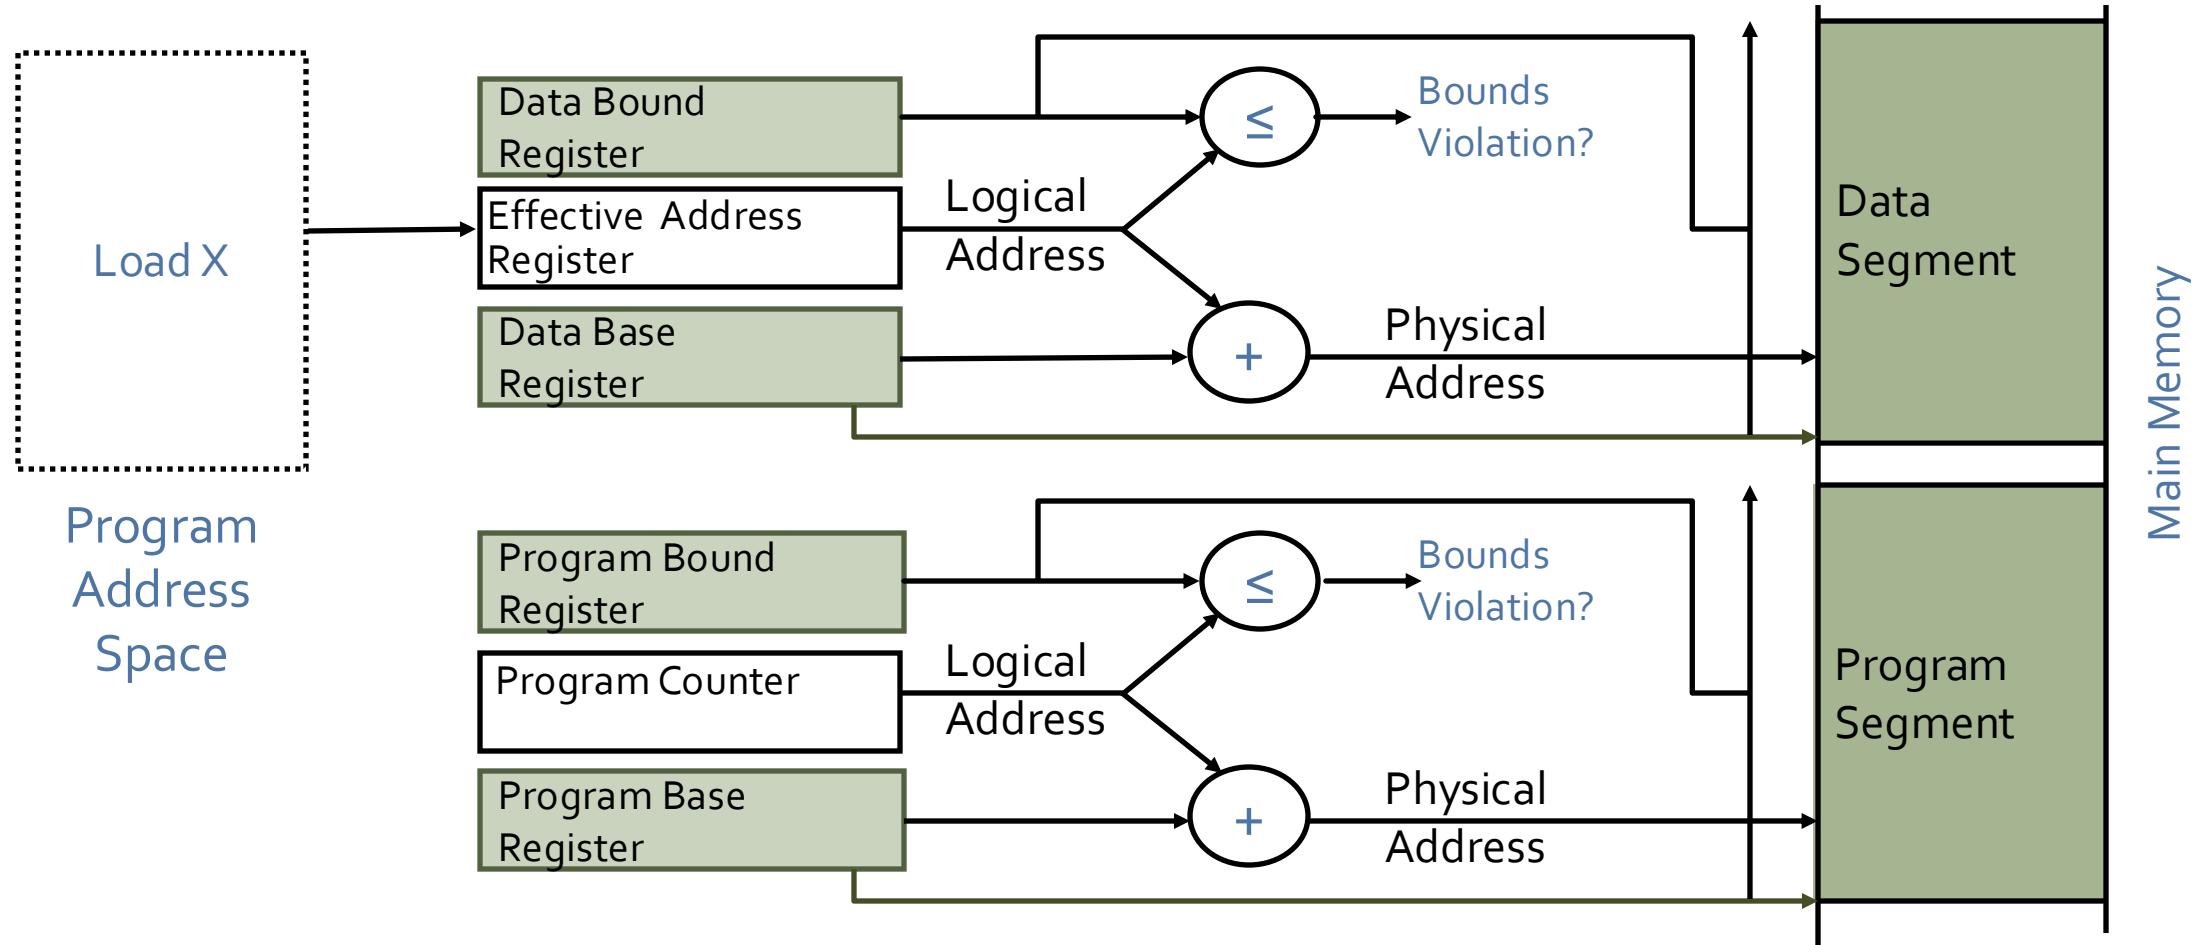


# Simple Base and Bound Translation



Base and bounds registers are visible/accessible only when processor is running in the *supervisor mode*

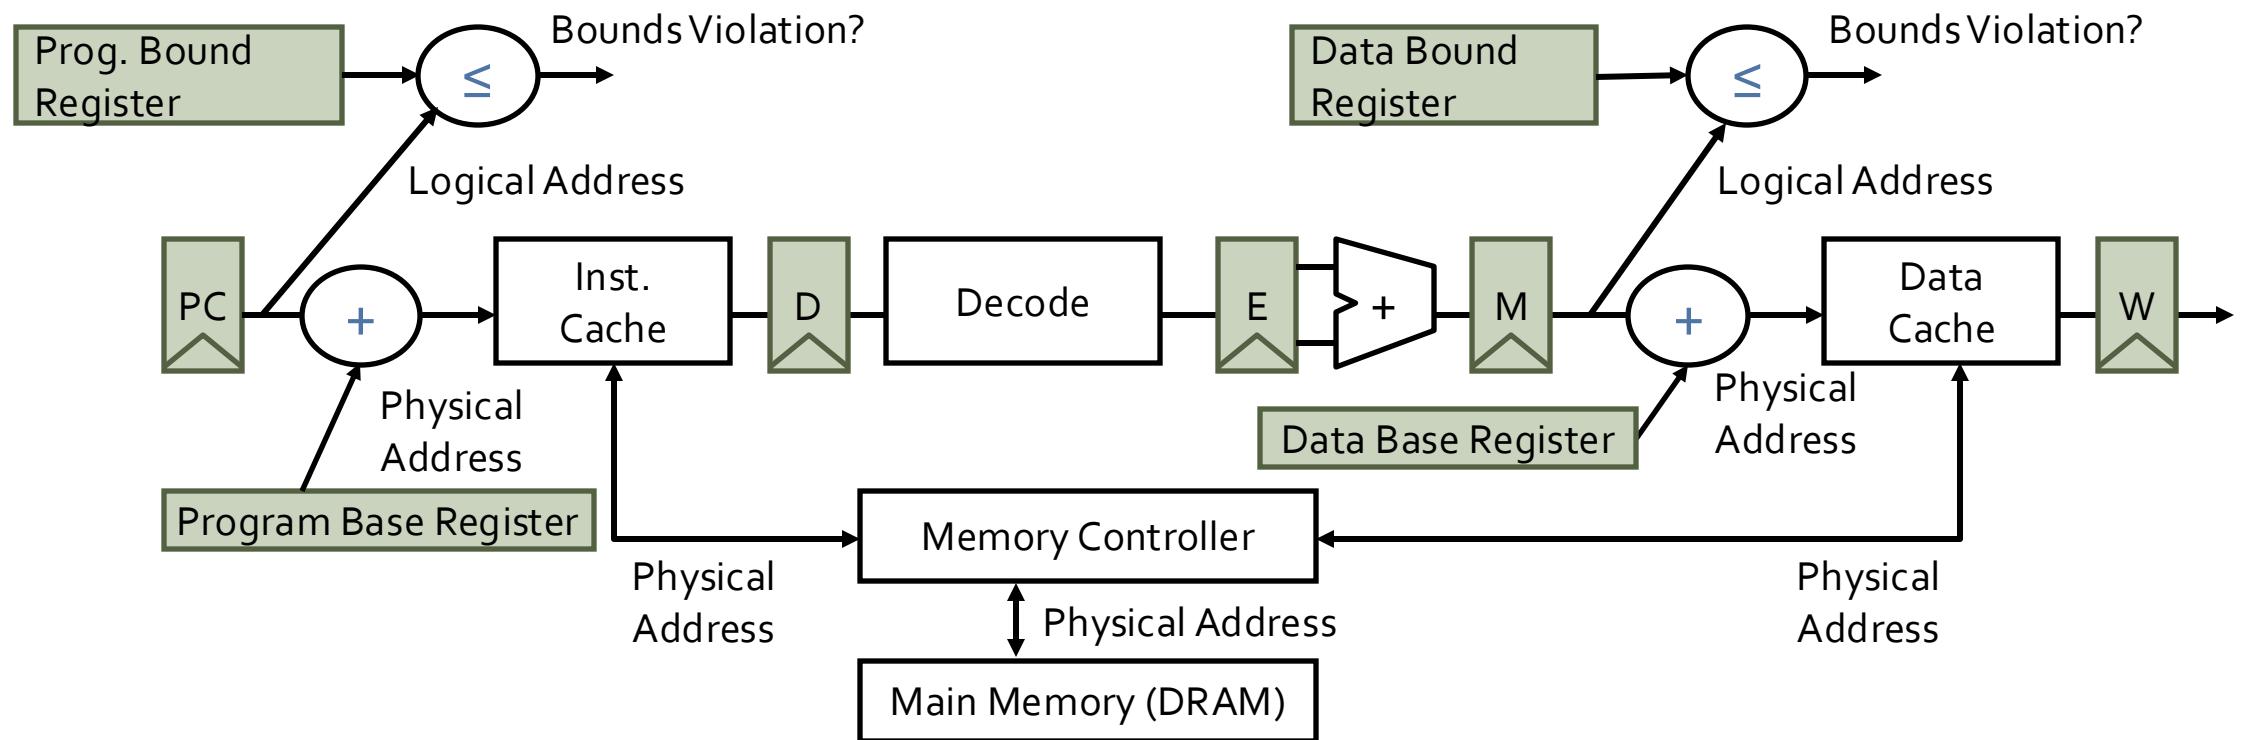
# Separate Areas for Program and Data



What is an advantage of this separation?

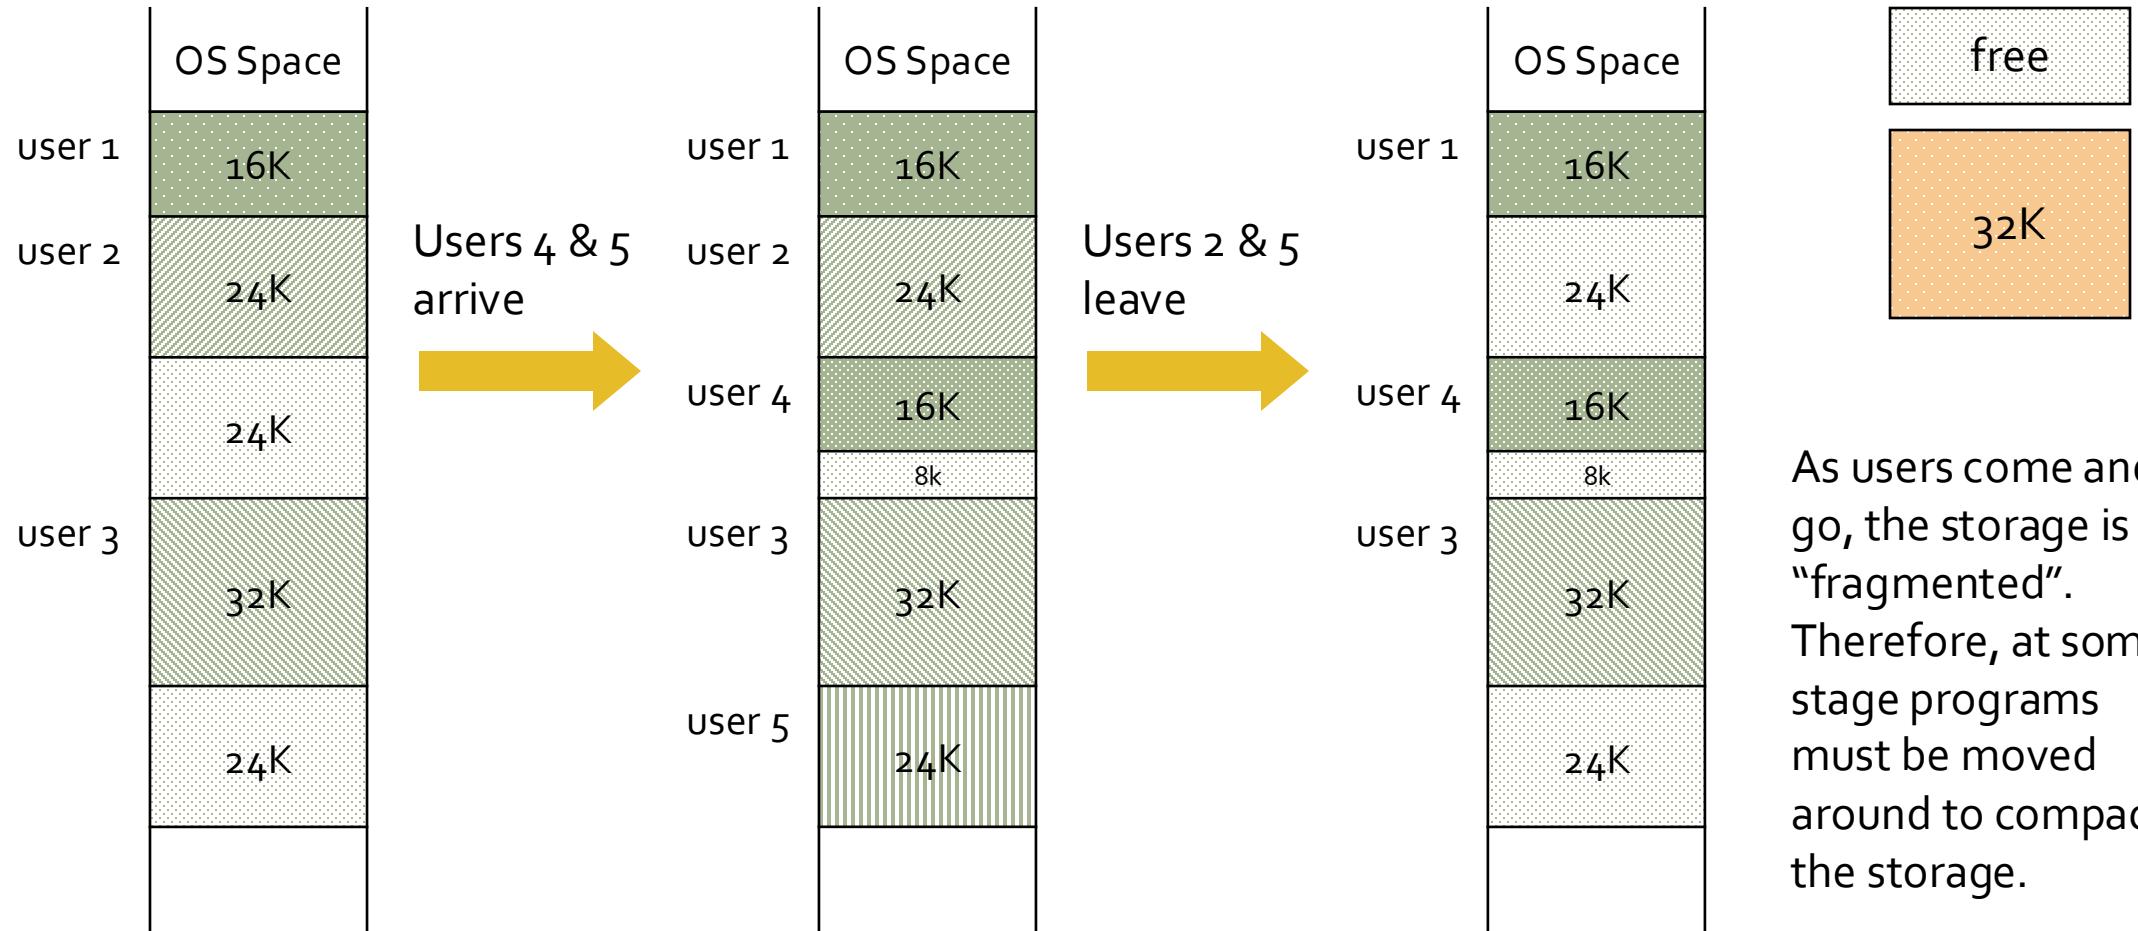
(Scheme used on all Cray vector supercomputers prior to X1, 2002)

# Base and Bound Machine



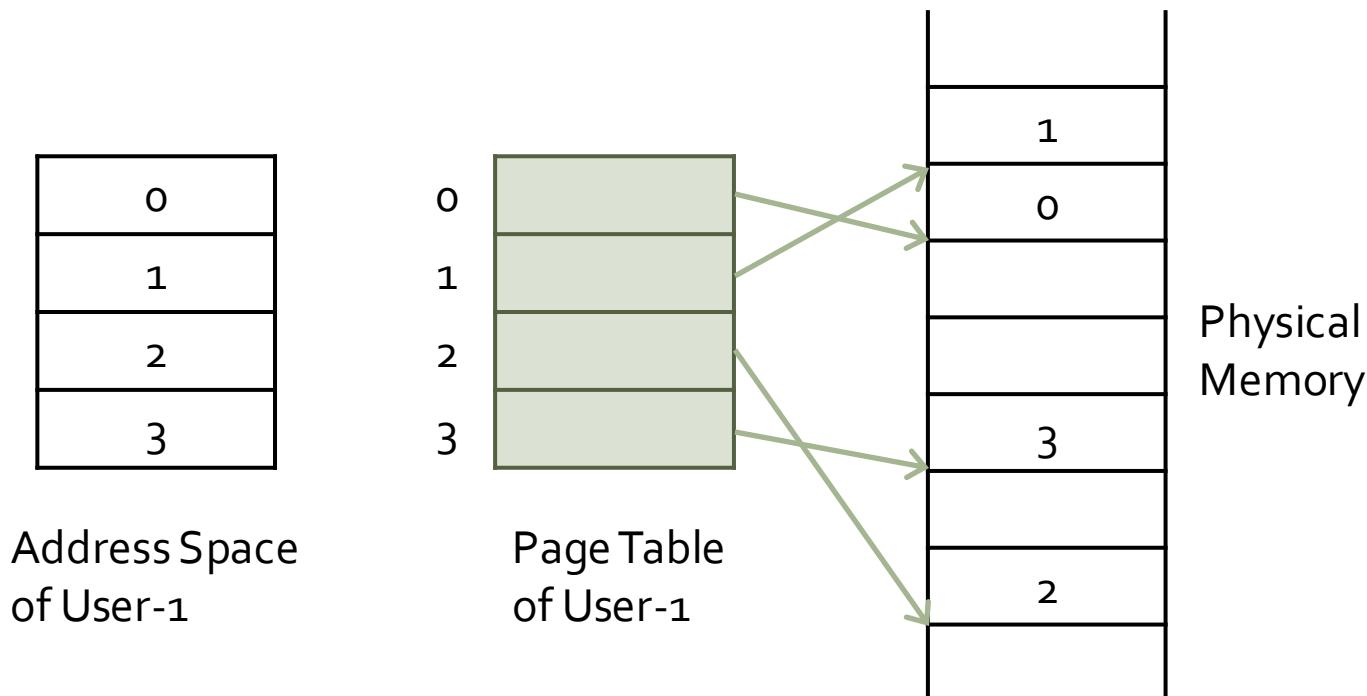
Can fold addition of base register into (base+offset) calculation using a carry-save adder  
(sums three numbers with only a few gate delays more than adding two numbers)

# Memory Fragmentation



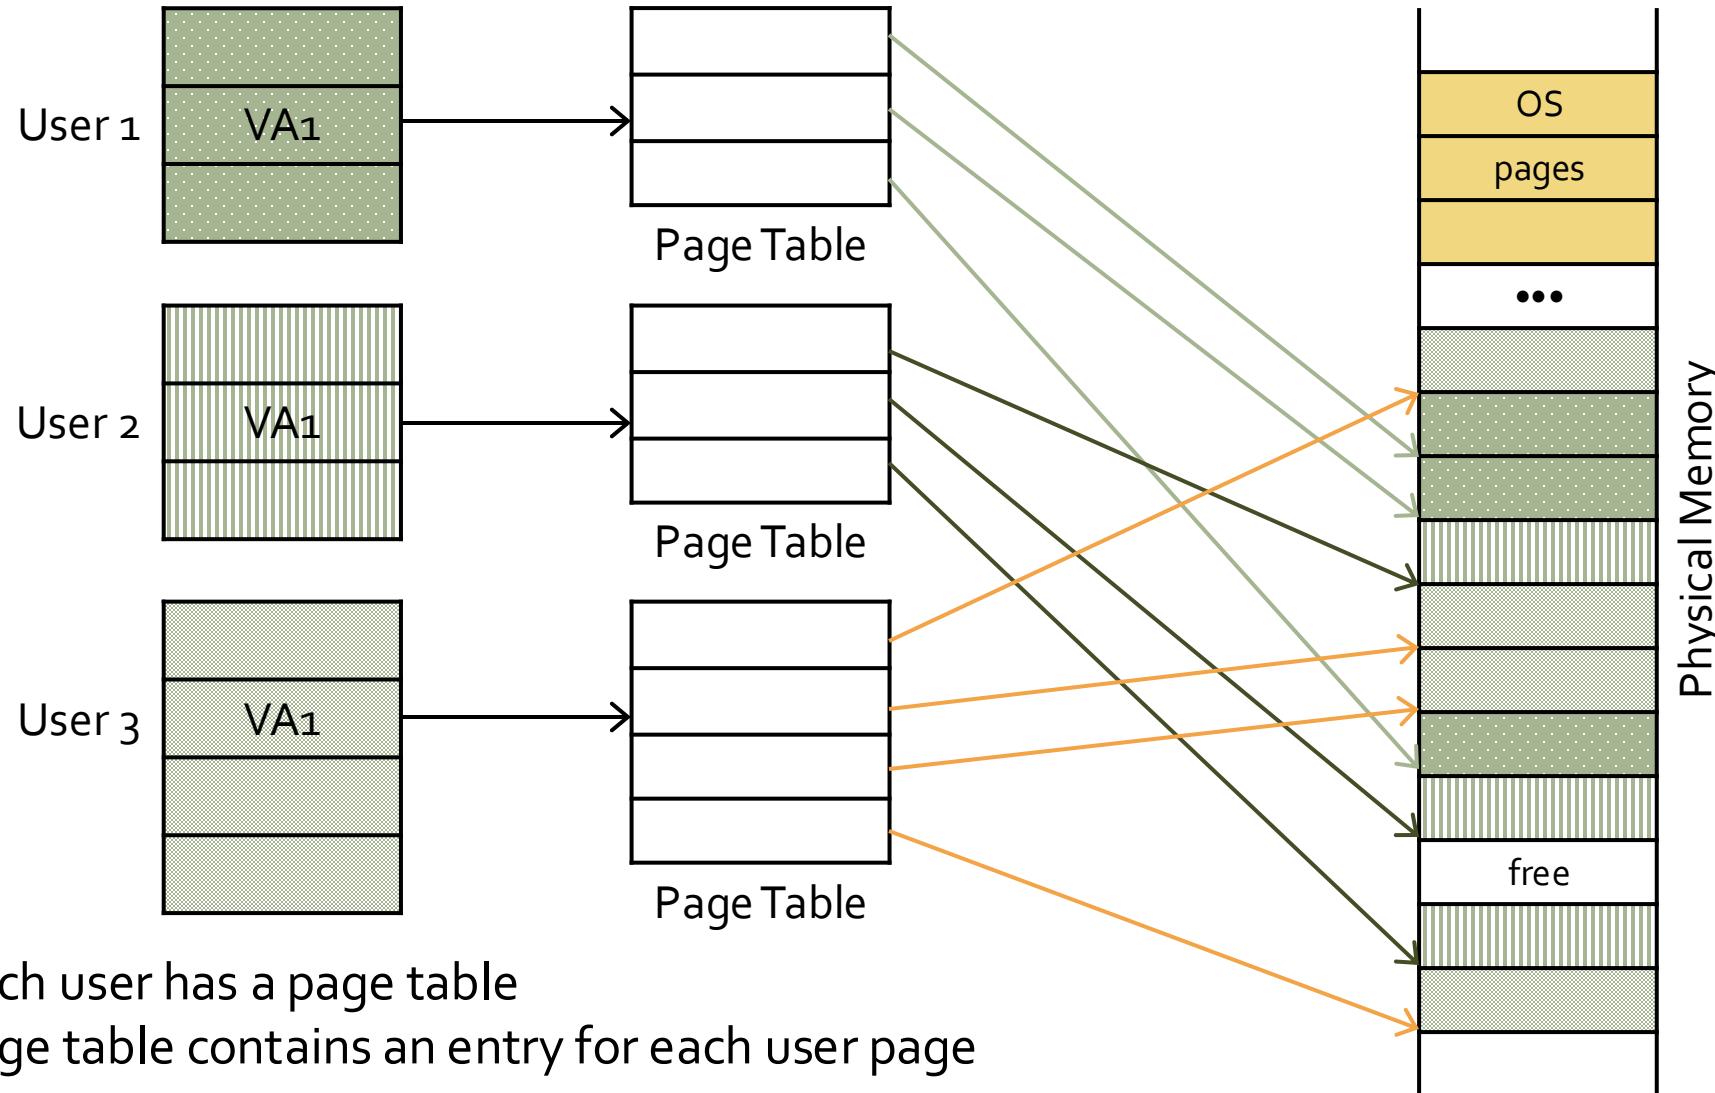
# Paged Memory Systems

- Processor-generated address can be interpreted as a pair <page number, offset>:  page number | offset
- A page table contains the physical address of the base of each page:



Page tables make it possible to store the pages of a program non-contiguously.

# Private Address Space per User



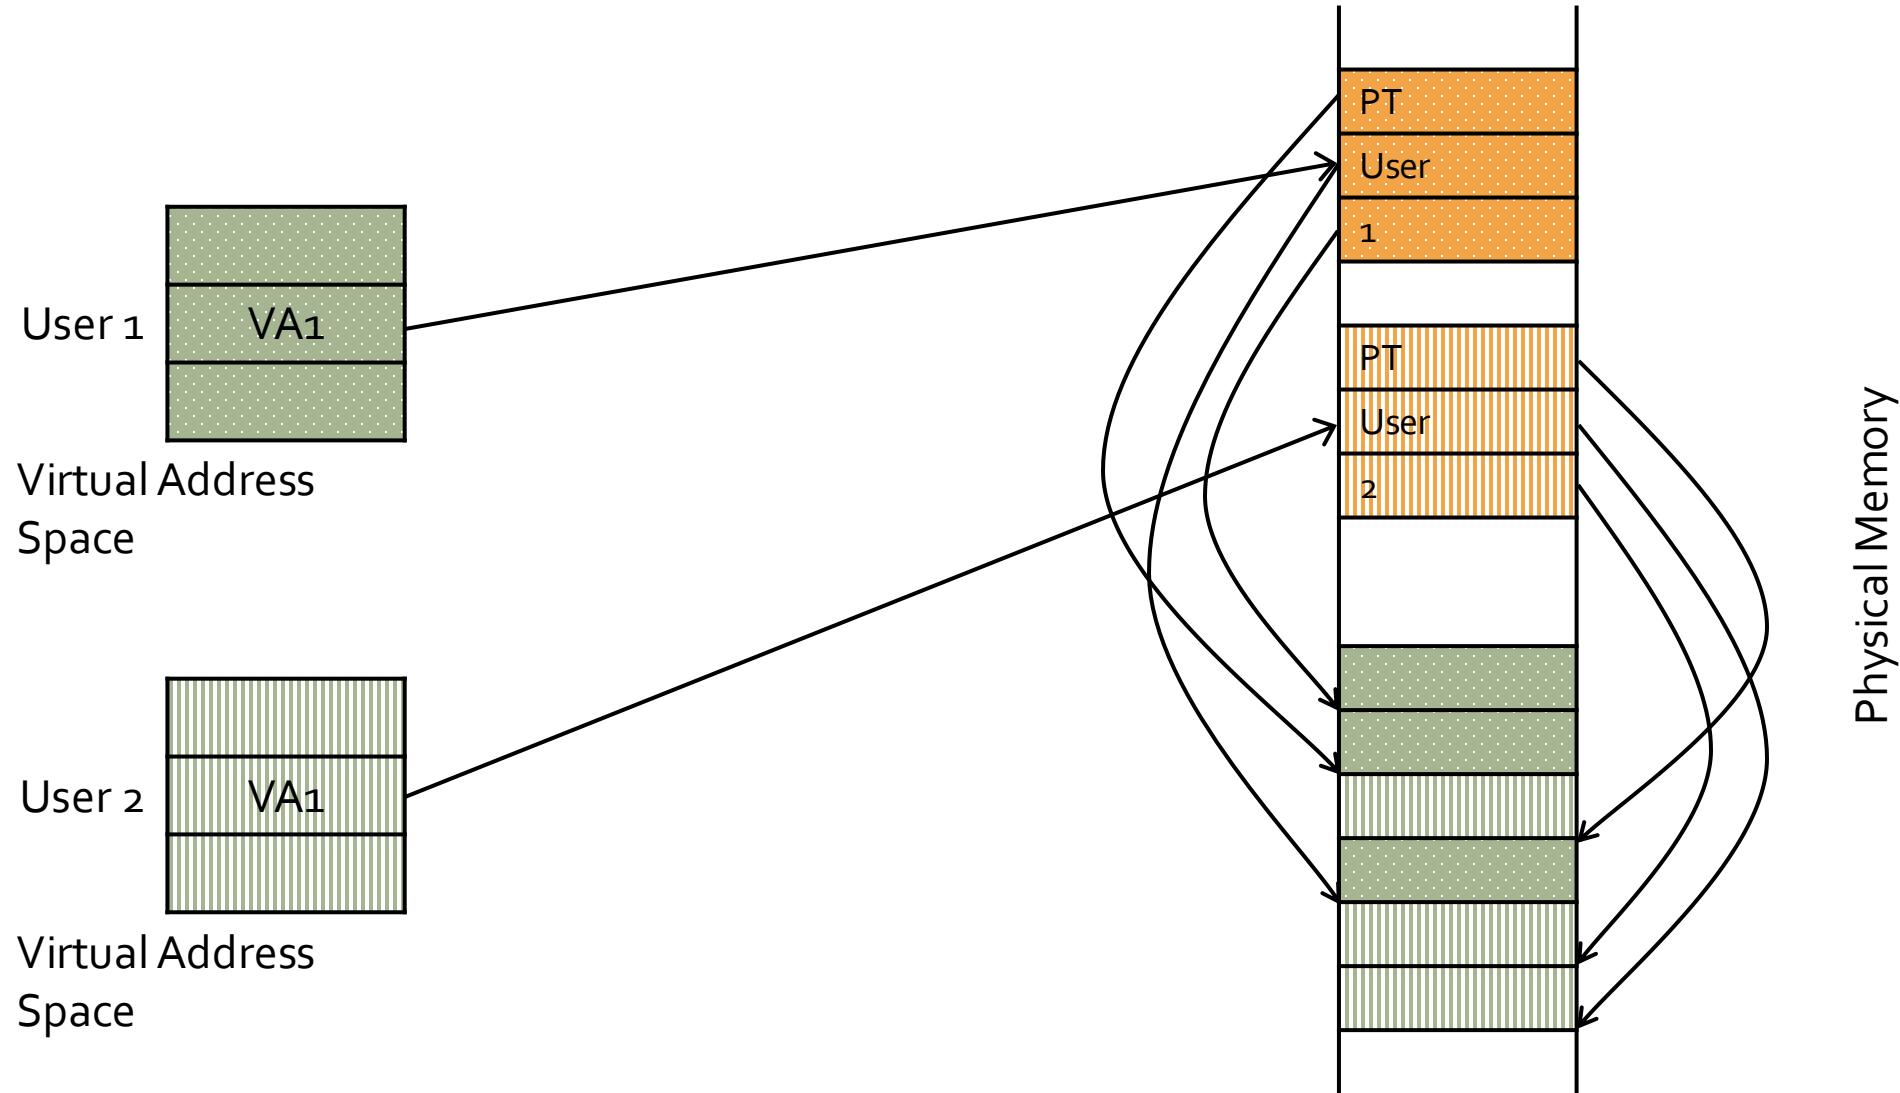
# Paging Simplifies Allocation

- Fixed-size pages can be kept on OS free list and allocated as needed to any process
- Process memory usage can easily grow and shrink dynamically
- Paging suffers from **internal fragmentation** where not all bytes on a page are used
  - Much less of an issue than external fragmentation or compaction for common page sizes (4-8KB)
  - But one reason that many oppose move to larger page sizes

# Where Should Page Tables Reside?

- Space required by the page tables (PT) is proportional to the address space, number of users, (inverse to) size of each page, ...
  - Space requirement is large
  - Too expensive to keep in registers
- Idea: Keep PT of the current user in special registers
  - may not be feasible for large page tables
  - Increases the cost of context swap
- Idea: Keep PTs in the main memory
  - needs one reference to retrieve the page base address and another to access the data word
    - **doubles the number of memory references!**
  - Storage space to store PT grows with size of memory

# Page Tables in Physical Memory

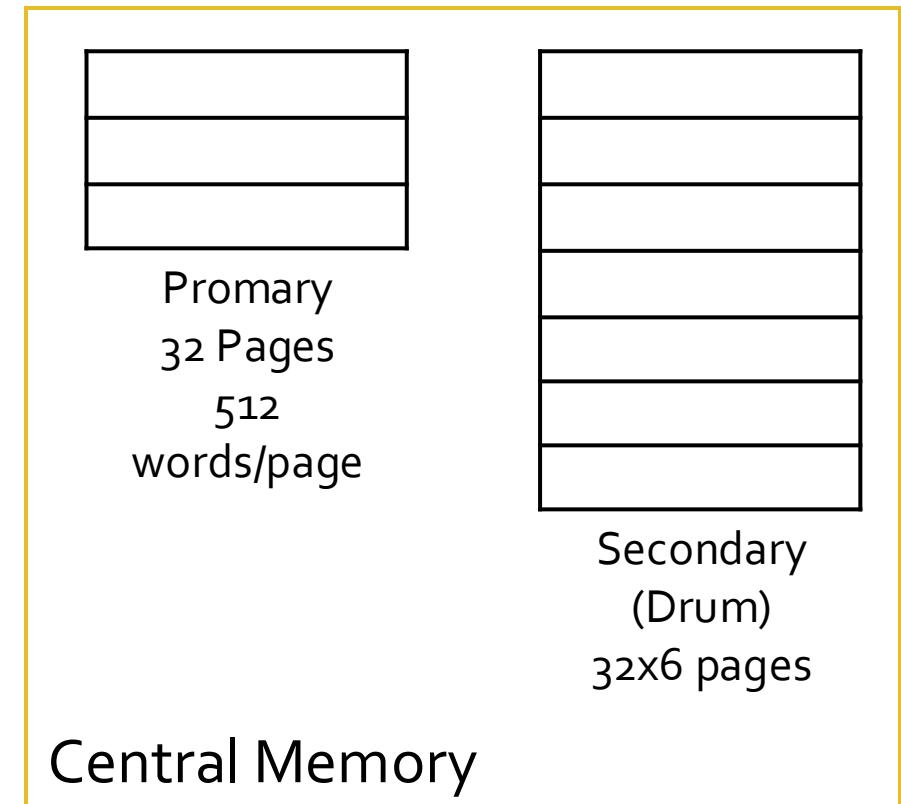


# A Problem in Early Sixties

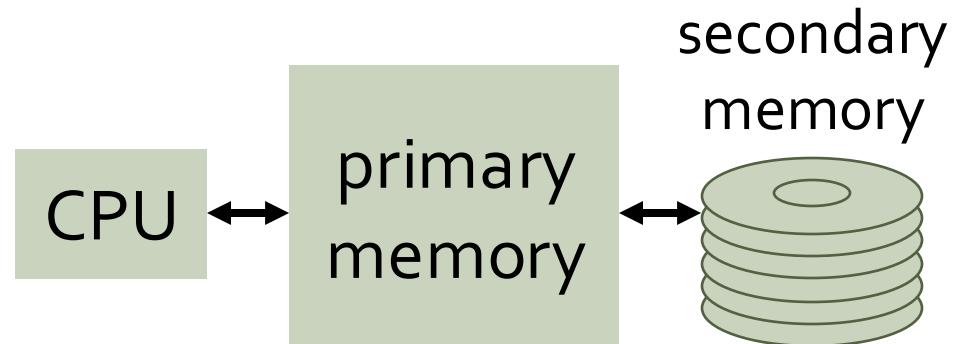
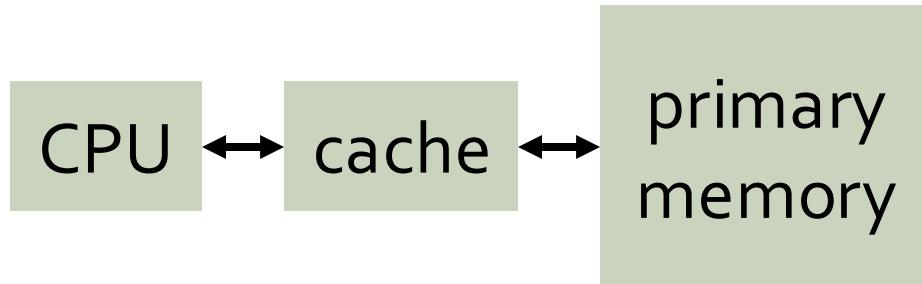
- There were many applications whose data could not fit in the main memory, e.g., payroll
- Paged memory system reduced fragmentation but still required the whole program to be resident in the main memory
- Programmers moved the data back and forth from the secondary store by overlaying it repeatedly on the primary store

# Demand Paging in Atlas (1962)

- A page from secondary storage is brought into the primary storage whenever it is (**implicitly**) demanded by the processor
- User sees the storage size of the secondary storage, since data transfer happens automatically
- Primary memory as a cache for secondary memory



# Caching vs. Demand Paging



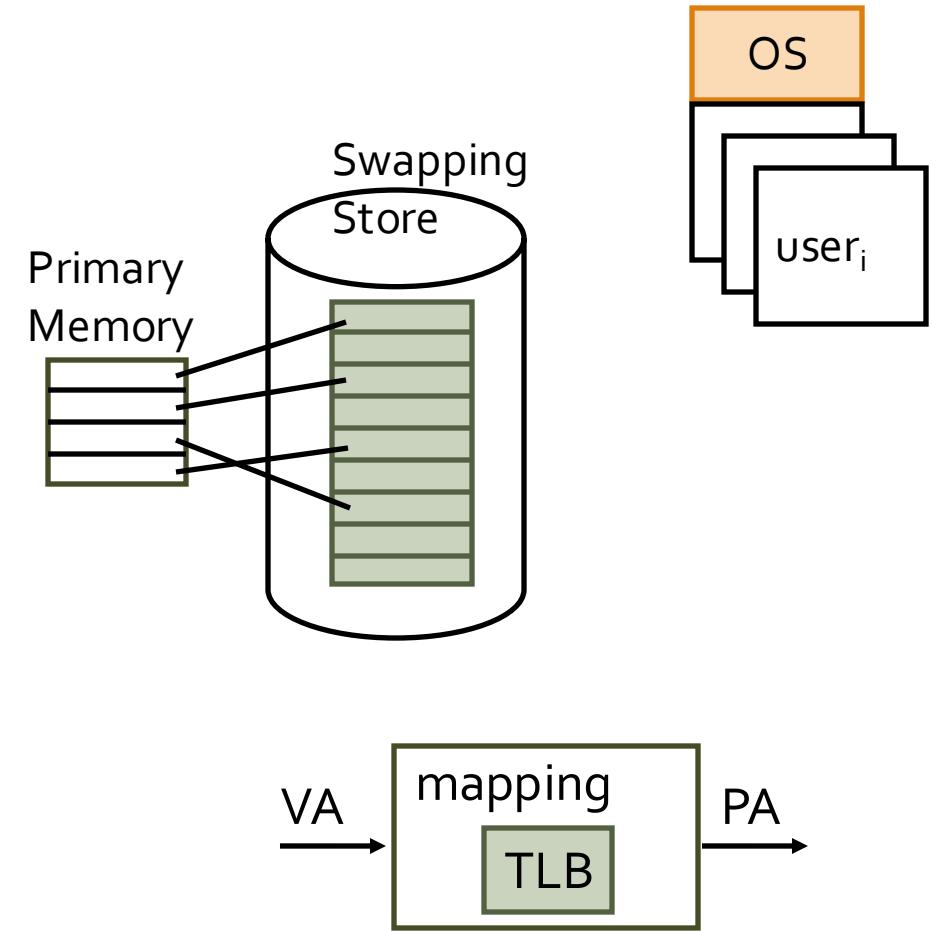
- Caching
- cache entry
- cache block (~32 bytes)
- cache miss rate (1% to 20%)
- cache hit (~1 cycle)
- cache miss (~100 cycles)
- a miss is handled in hardware

- Demand paging
- page frame
- page (~4K bytes)
- page miss rate (<0.001%)
- page hit (~100 cycles)
- page miss (~5M cycles)
- a miss is handled mostly in software

# Modern Virtual Memory Systems

Illusion of a large, private, uniform store

- Protection & Privacy
  - Several users, each with their private address space and one or more shared address spaces  
page table  $\equiv$  name space
- Demand Paging
  - Provides the ability to run programs larger than the primary memory
  - Hides differences in machine configurations
- The price is address translation on each memory reference

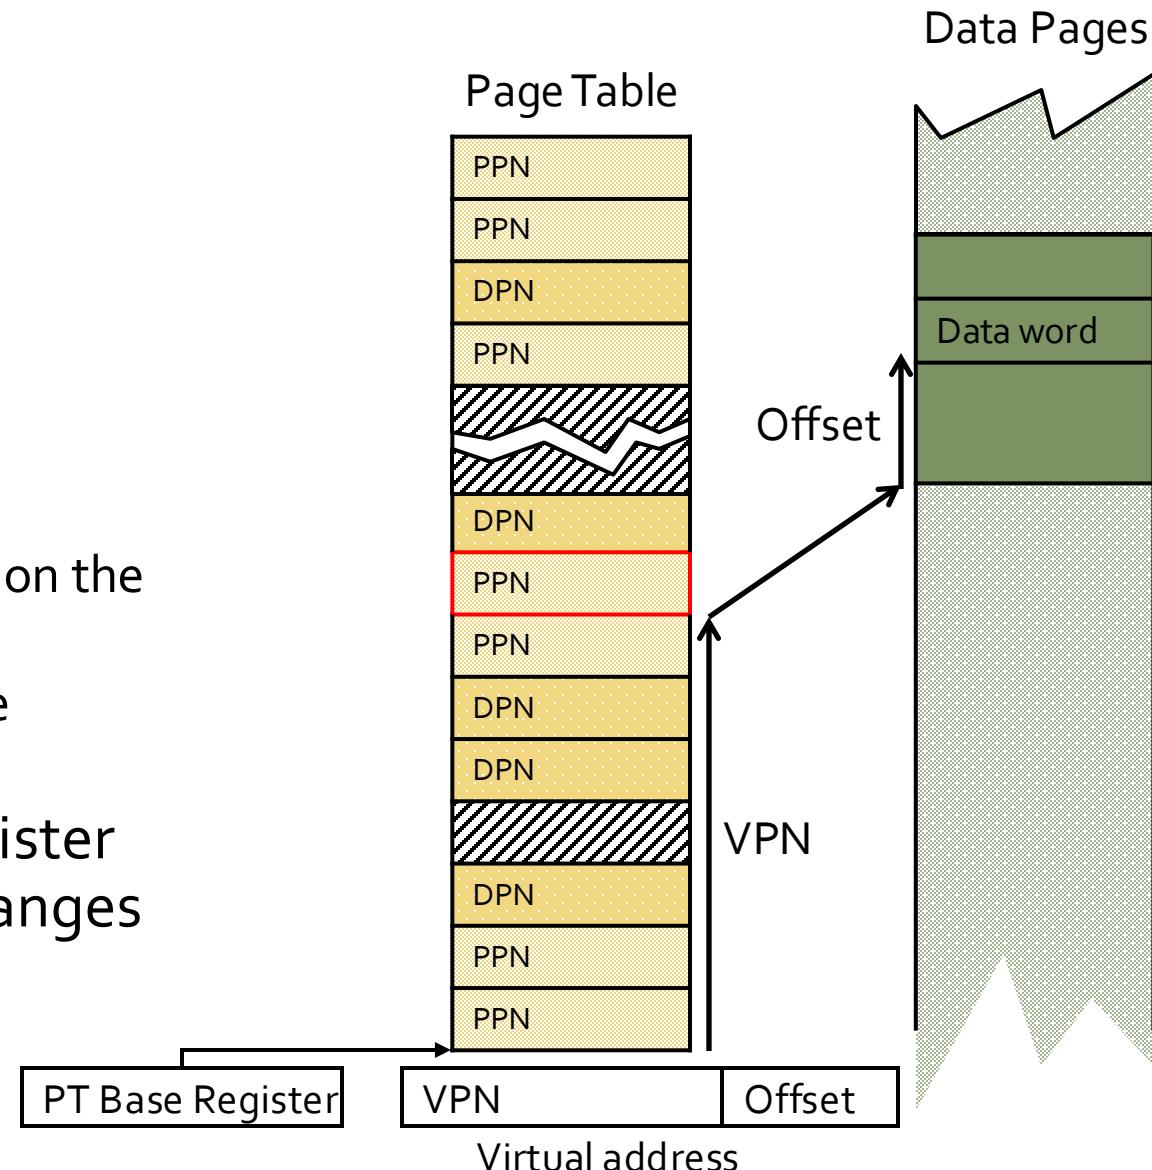


# Agenda

- Evolution of Virtual Memory
- Modern VM Implementation
- TLB & Cache Organization
- Modern Usage

# Linear Page Table

- Page Table Entry (PTE) contains:
  - A bit to indicate if a page exists
  - PPN (physical page number) for a memory-resident page
  - DPN (disk page number) for a page on the disk
  - Status bits for protection and usage
- OS sets the Page Table Base Register whenever active user process changes



# Size of Linear Page Table

With 32-bit addresses, 4-KB pages & 4-byte PTEs:

- ⇒  $2^{20}$  PTEs, i.e., 4 MB page table per user per process
- ⇒ 4 GB of swap needed to back up full virtual address space

Larger pages?

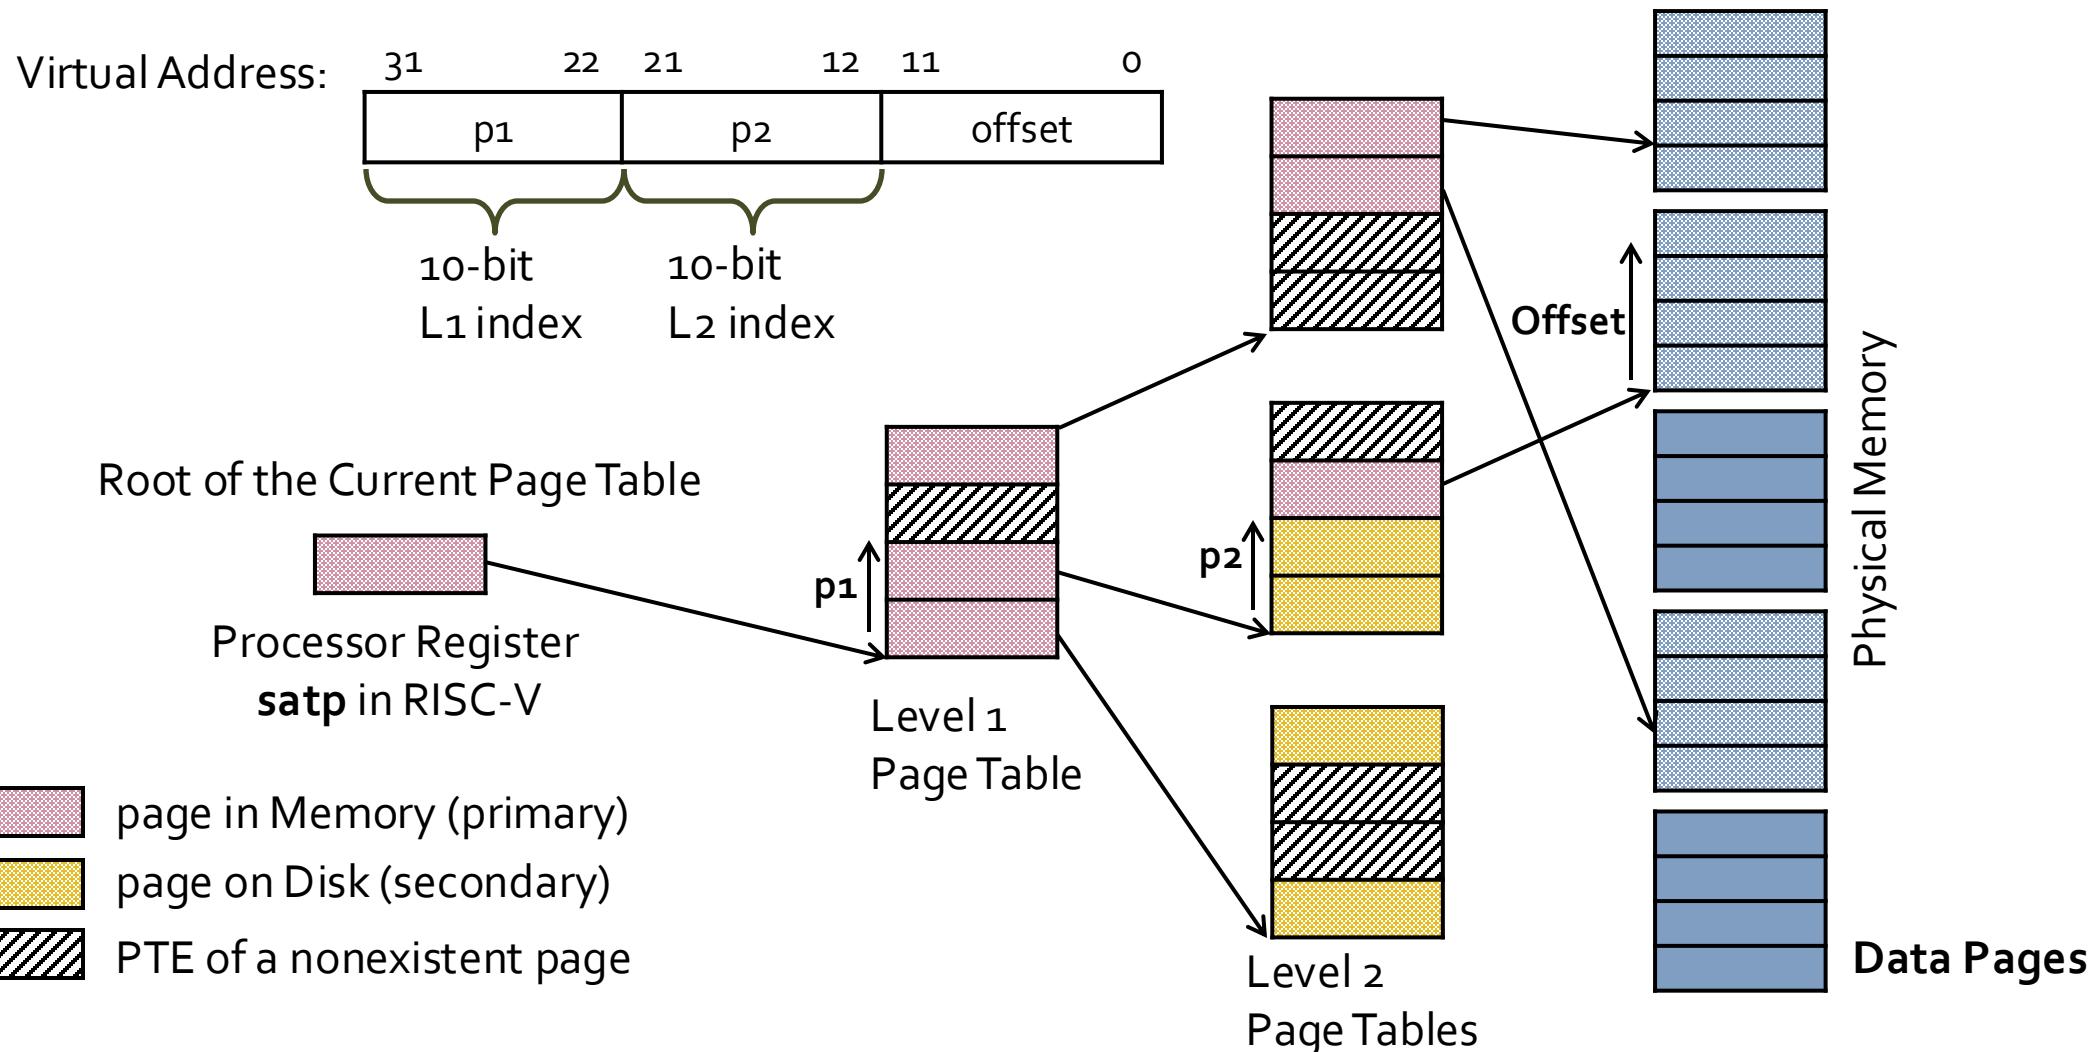
- Internal fragmentation (Not all memory in page is used)
- Larger page fault penalty (more time to read from disk)

What about 64-bit virtual address space???

- Even 1MB pages would require  $2^{44}$  8-byte PTEs (TB!)

**What is the “saving grace” ?**

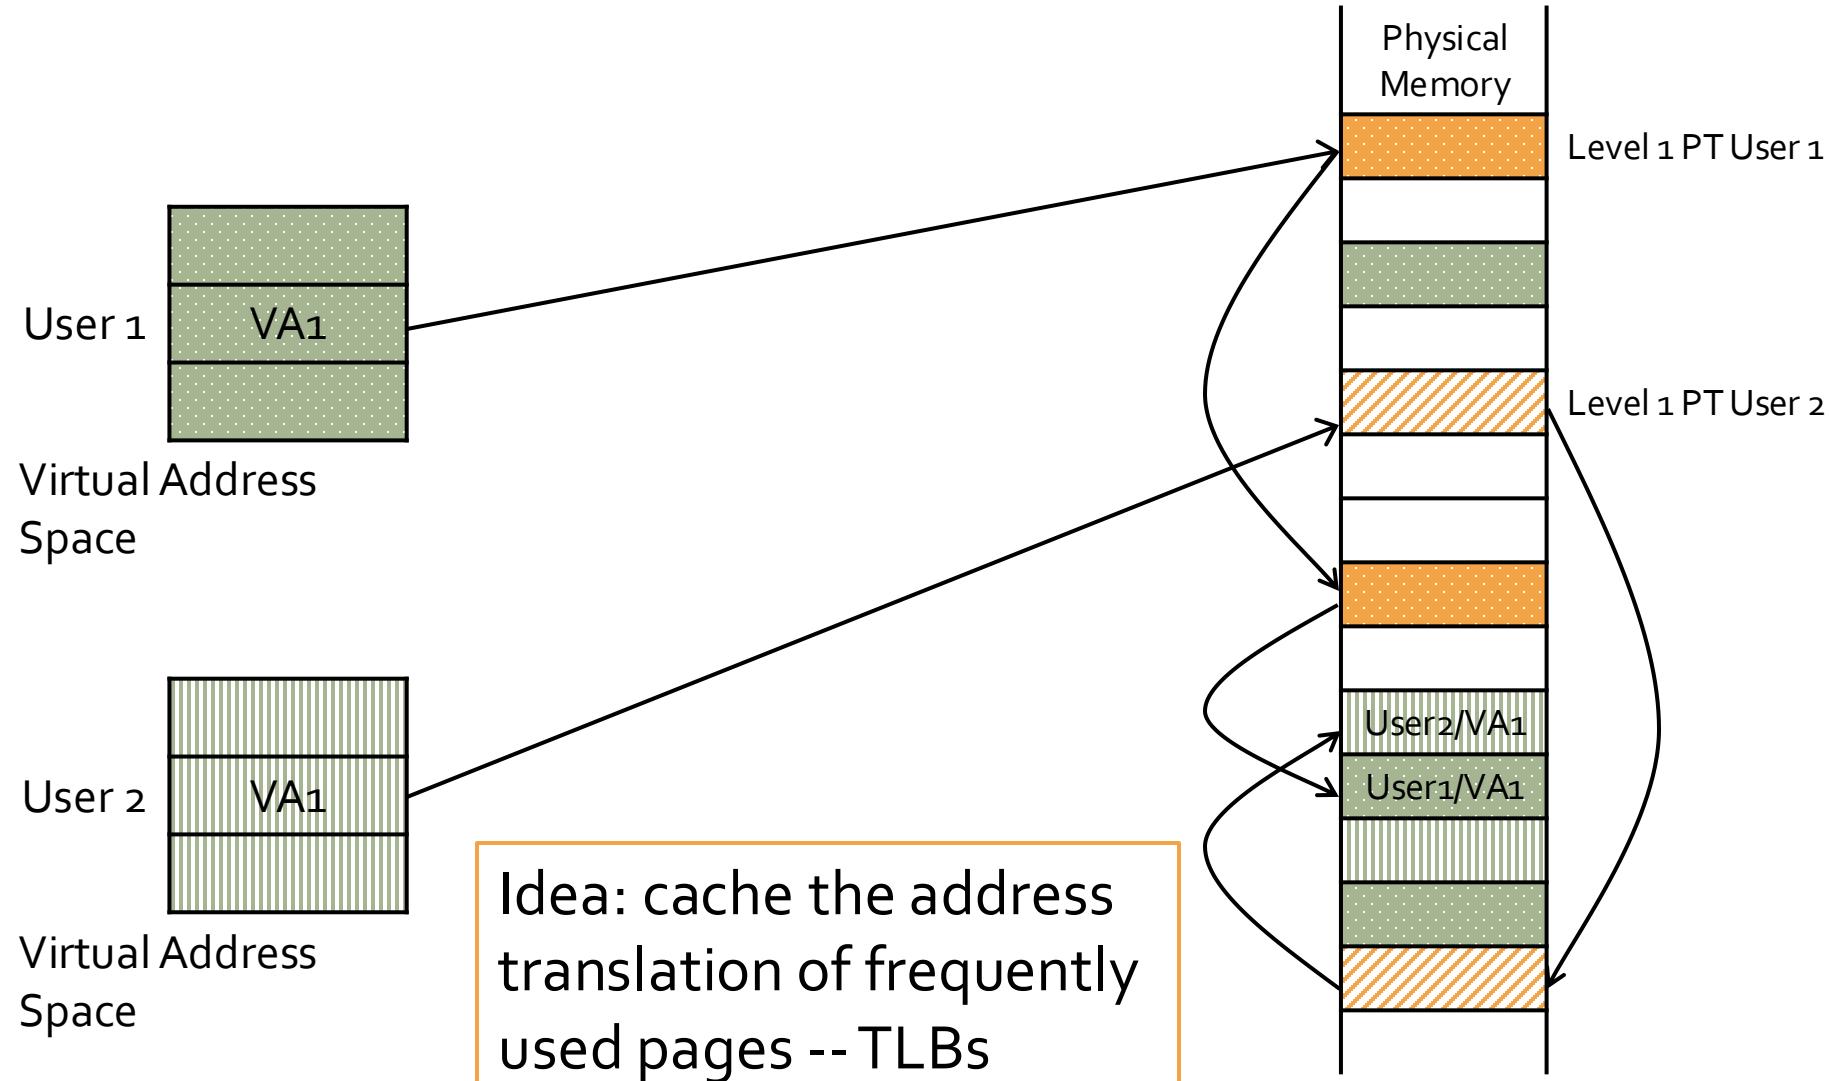
# Hierarchical Page Table



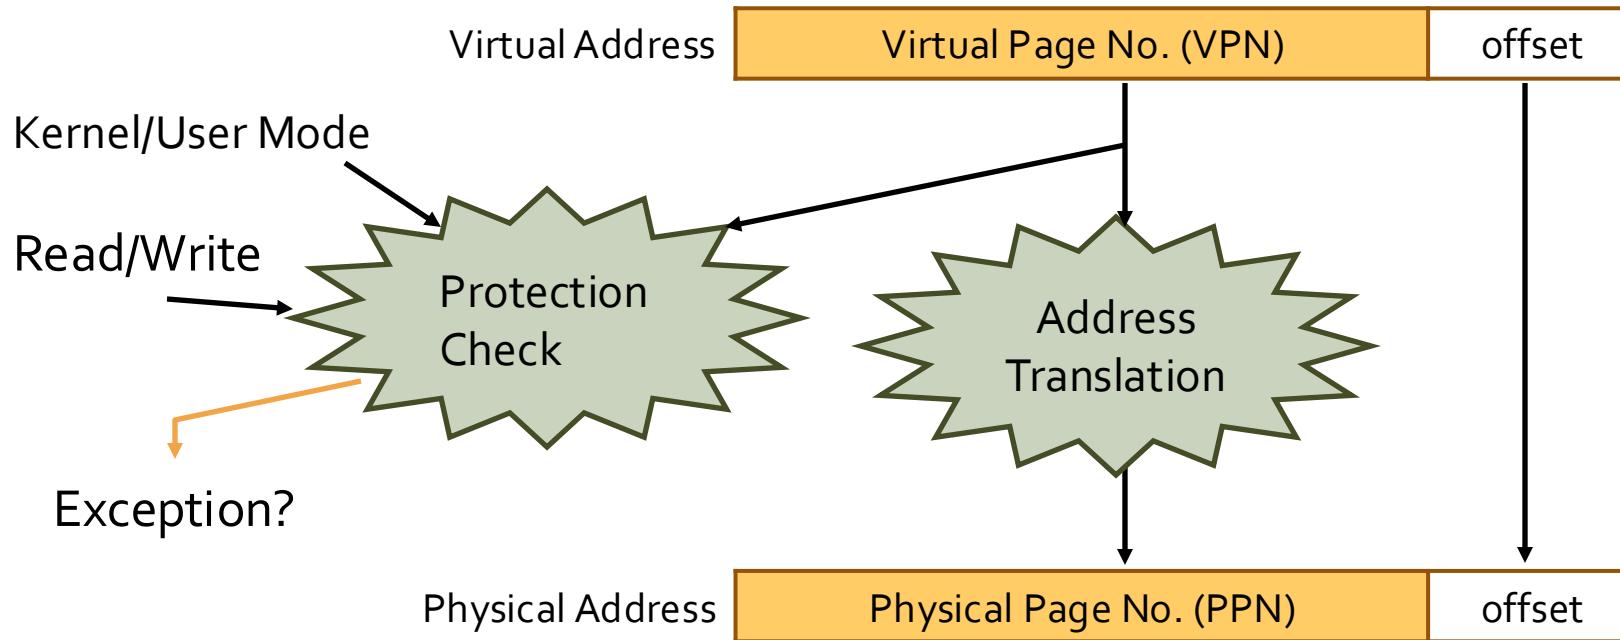
# Real-world ISAs – RISC-V

- RISC-V defines three different virtual memory specifications
  - Sv32: Two level hierarchy, 32-bit virtual, 34-bit physical address
  - Sv39: Three level, 39-bit virtual, 56-bit physical address
  - Sv48: Four level, 48-bit virtual, 56-bit physical address
  - Up to chip designers to choose what to support
  - OS can learn which is implemented by reading the CSR “**satp**” (Supervisor Address Translation and Protection)
- Some specs have different PTE sizes
  - Sv32 has 32 bit PTEs
  - Sv39 and Sv48 has 64-bit PTEs

# Two-Level Page Tables in Physical Memory



# Address Translation & Protection



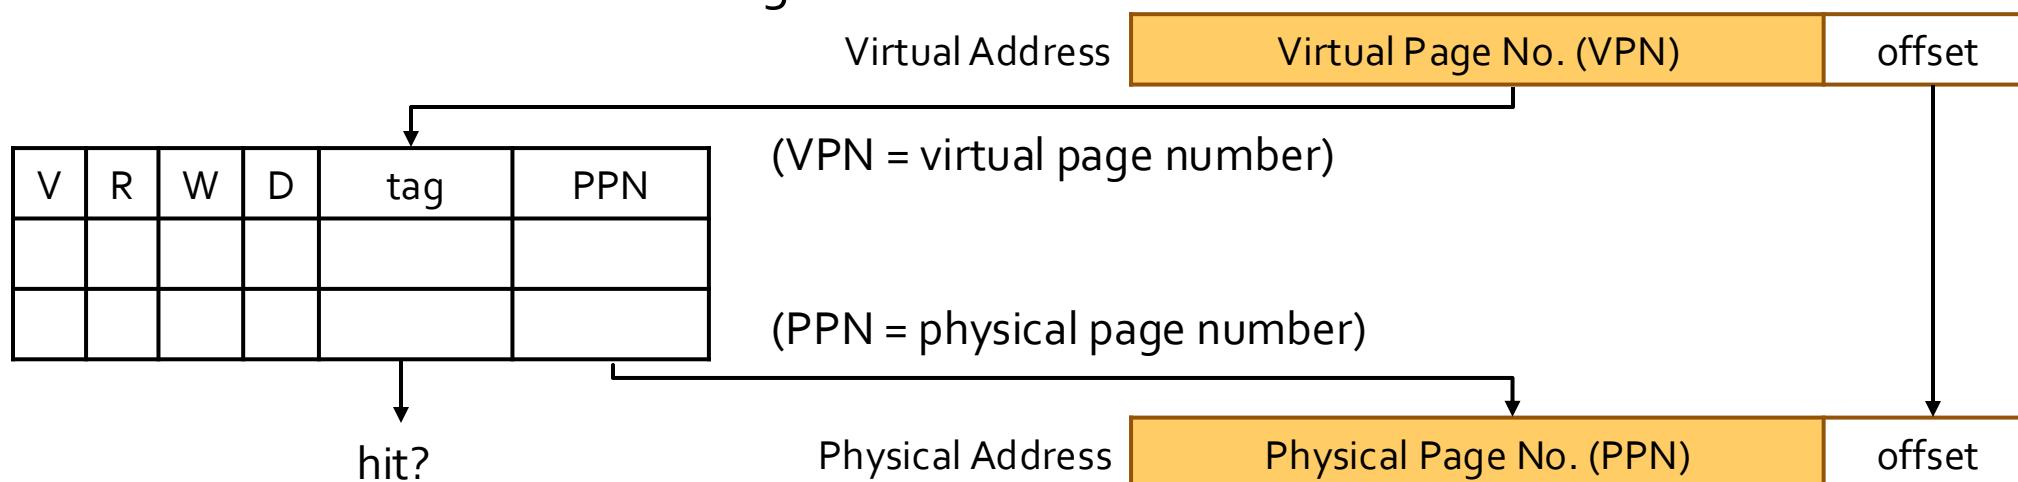
- Every instruction and data access needs address translation and protection checks
- A good Virtual Memory (VM) design needs to be fast (~ one cycle) and space efficient

# Translation Lookaside Buffers (TLB)

- Problem: Address translation is very expensive!
  - In a hierarchical page table, each reference becomes several memory accesses
- Solution: Cache translations in TLB

TLB hit  $\Rightarrow$  Single-Cycle Translation

TLB miss  $\Rightarrow$  Page-Table Walk to refill



# TLB Designs

- Typically 32-128 entries, usually fully associative
  - Each entry maps a large page, hence less spatial locality across pages → more likely that two entries conflict
  - Sometimes larger TLBs (256-512 entries) are 4-8 way set-associative
  - Larger systems sometimes have multi-level (L<sub>1</sub> and L<sub>2</sub>) TLBs
- Random (Clock Algorithm) or FIFO replacement policy
- Keep process information in TLB?
  - No process id → Must flush on context switch
  - Tag each entry with process id → No flush, but costlier
- TLB Reach: Size of largest virtual address space that can be simultaneously mapped by TLB
  - Example: 64 TLB entries, 4KB pages, one page per entry

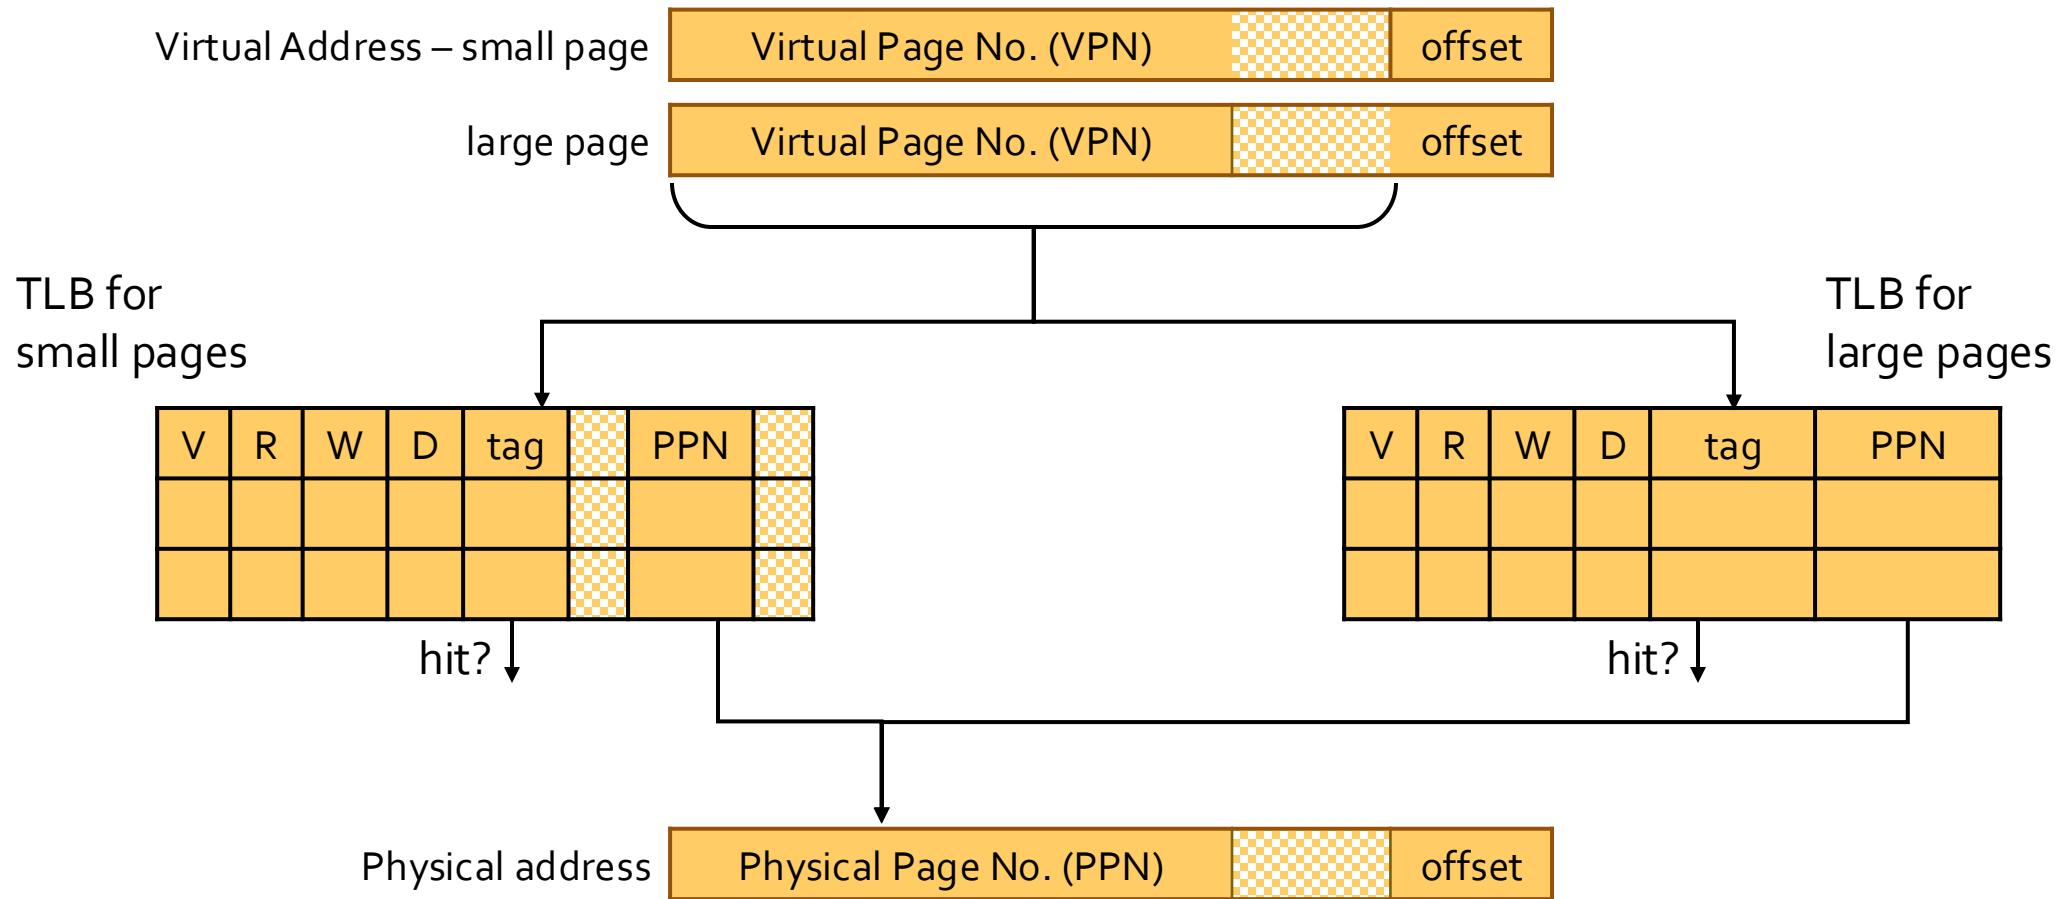
TLB Reach = 64 entries \* 4 KB = 256 KB (if contiguous) ?

# TLB Extensions

- Address Space Identifier (ASID)
  - Allow TLB Entries from multiple processes to be in TLB at same time. ID of address space (Process) is matched on.
  - Global Bit (G) can match on all ASIDs
- Variable Page Size (PS)
  - Can increase reach on a per page basis
- Kernel/User (K)

K	V	R	W	D	tag	PPN	PS	G	ASID

# TLB Supporting Multiple Page Sizes

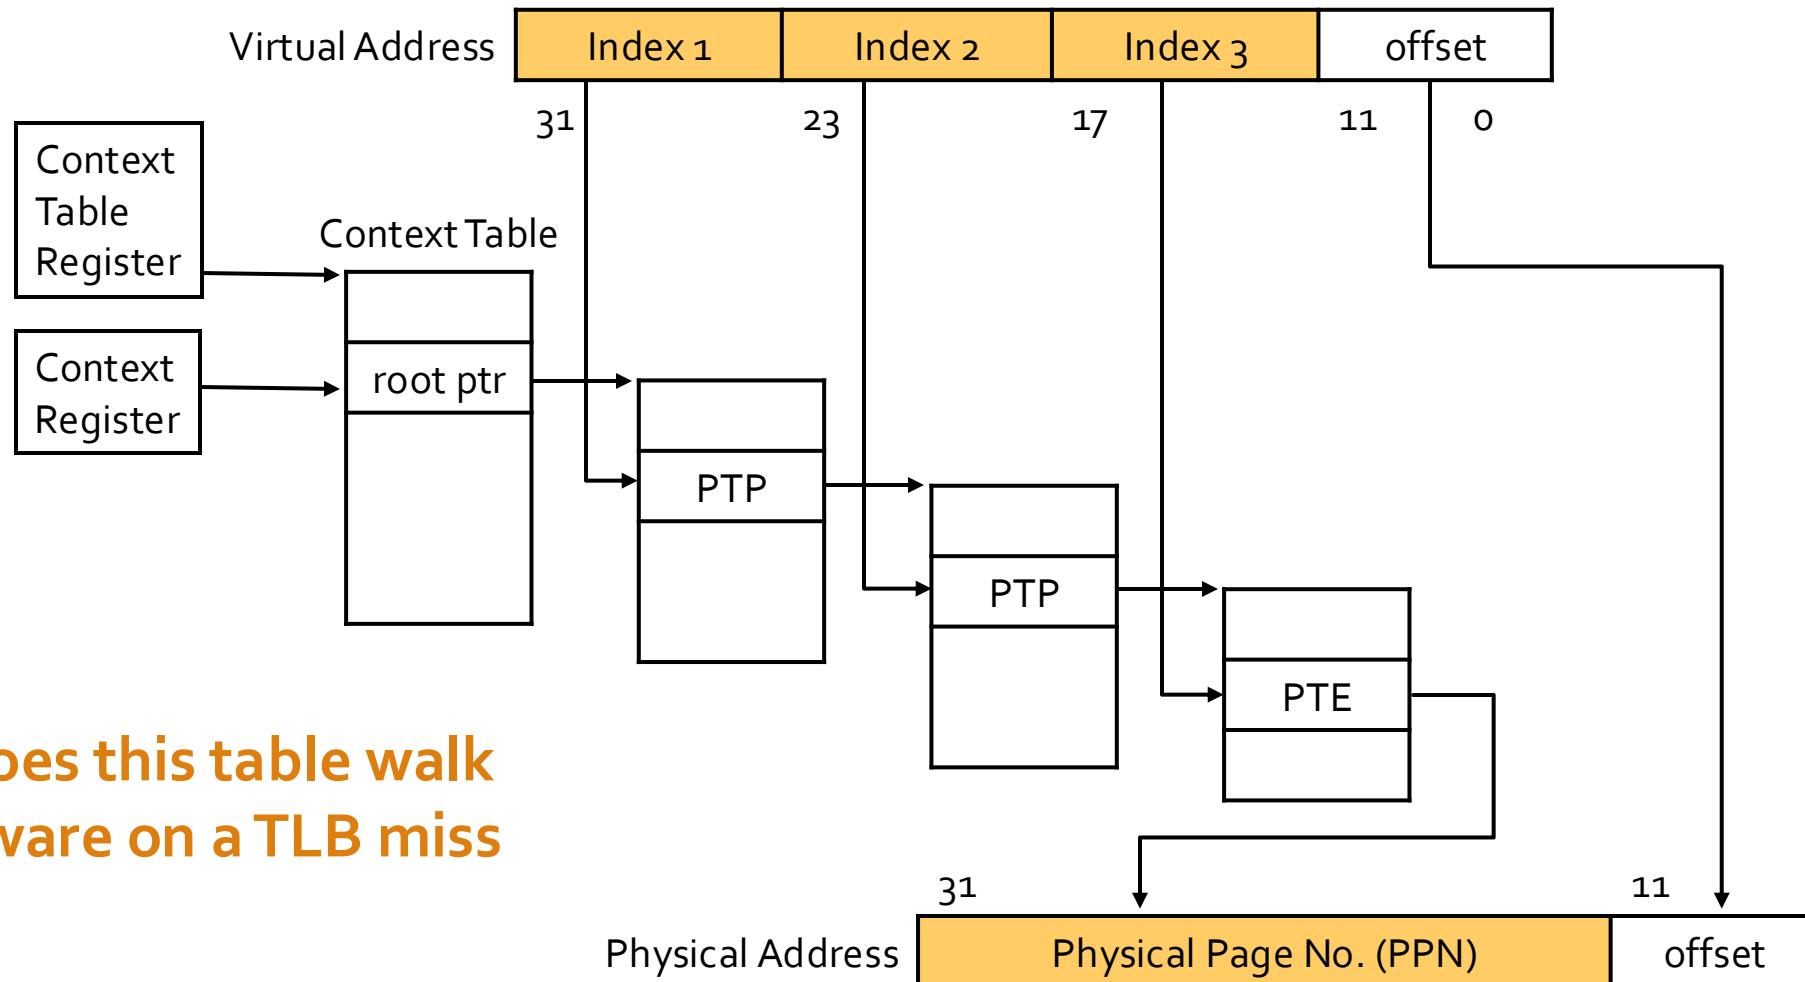


Pros/cons compared to unified TLB?

# Handling a TLB Miss

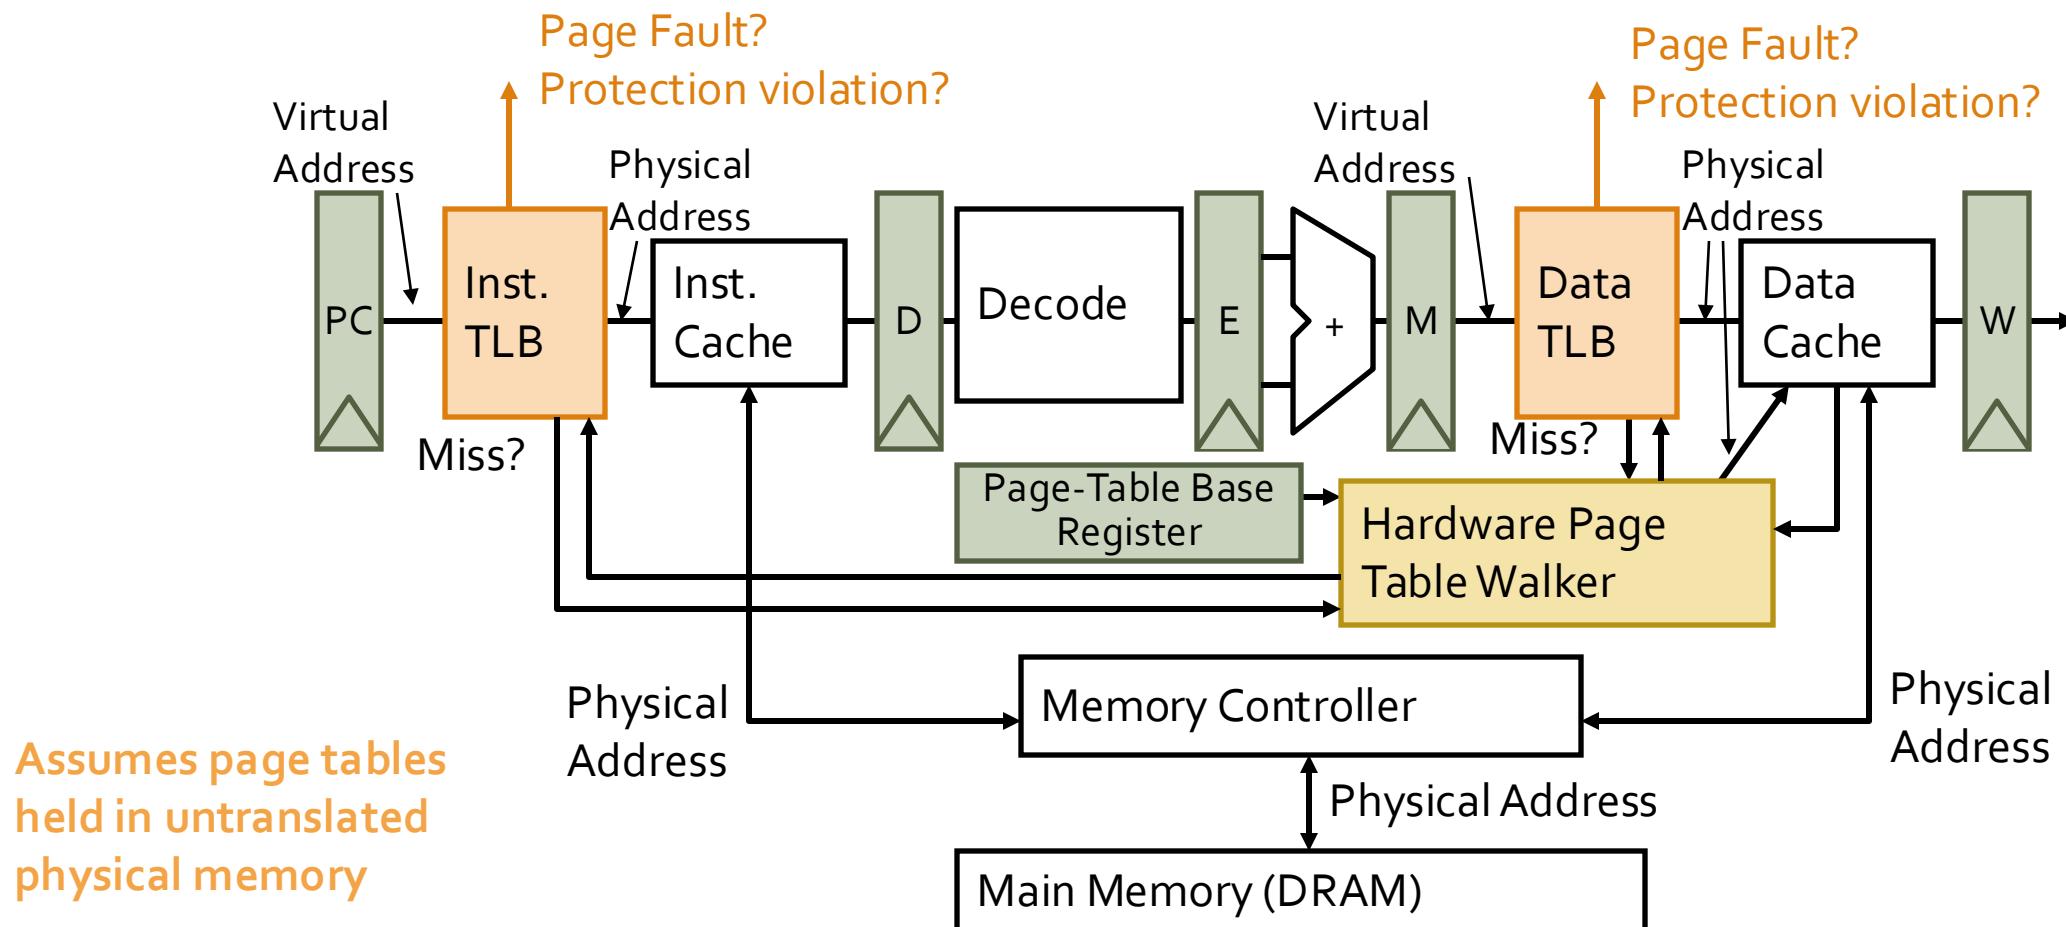
- Software (MIPS, Alpha)
  - TLB miss causes an exception and the operating system walks the page tables and reloads TLB. A privileged “untranslated” addressing mode used for walk
  - Software TLB miss can be very expensive on out-of-order superscalar processor as requires a flush of pipeline to jump to trap handler
- Hardware (SPARC v8, x86, PowerPC)
  - A memory management unit (MMU) walks the page tables and reloads the TLB
  - If a missing (data or PT) page is encountered during the TLB reloading, MMU gives up and signals a Page-Fault exception for the original instruction
- NOTE: A given ISA can use either TLB miss strategy

# Hierarchical Page Table Walk: SPARC v8

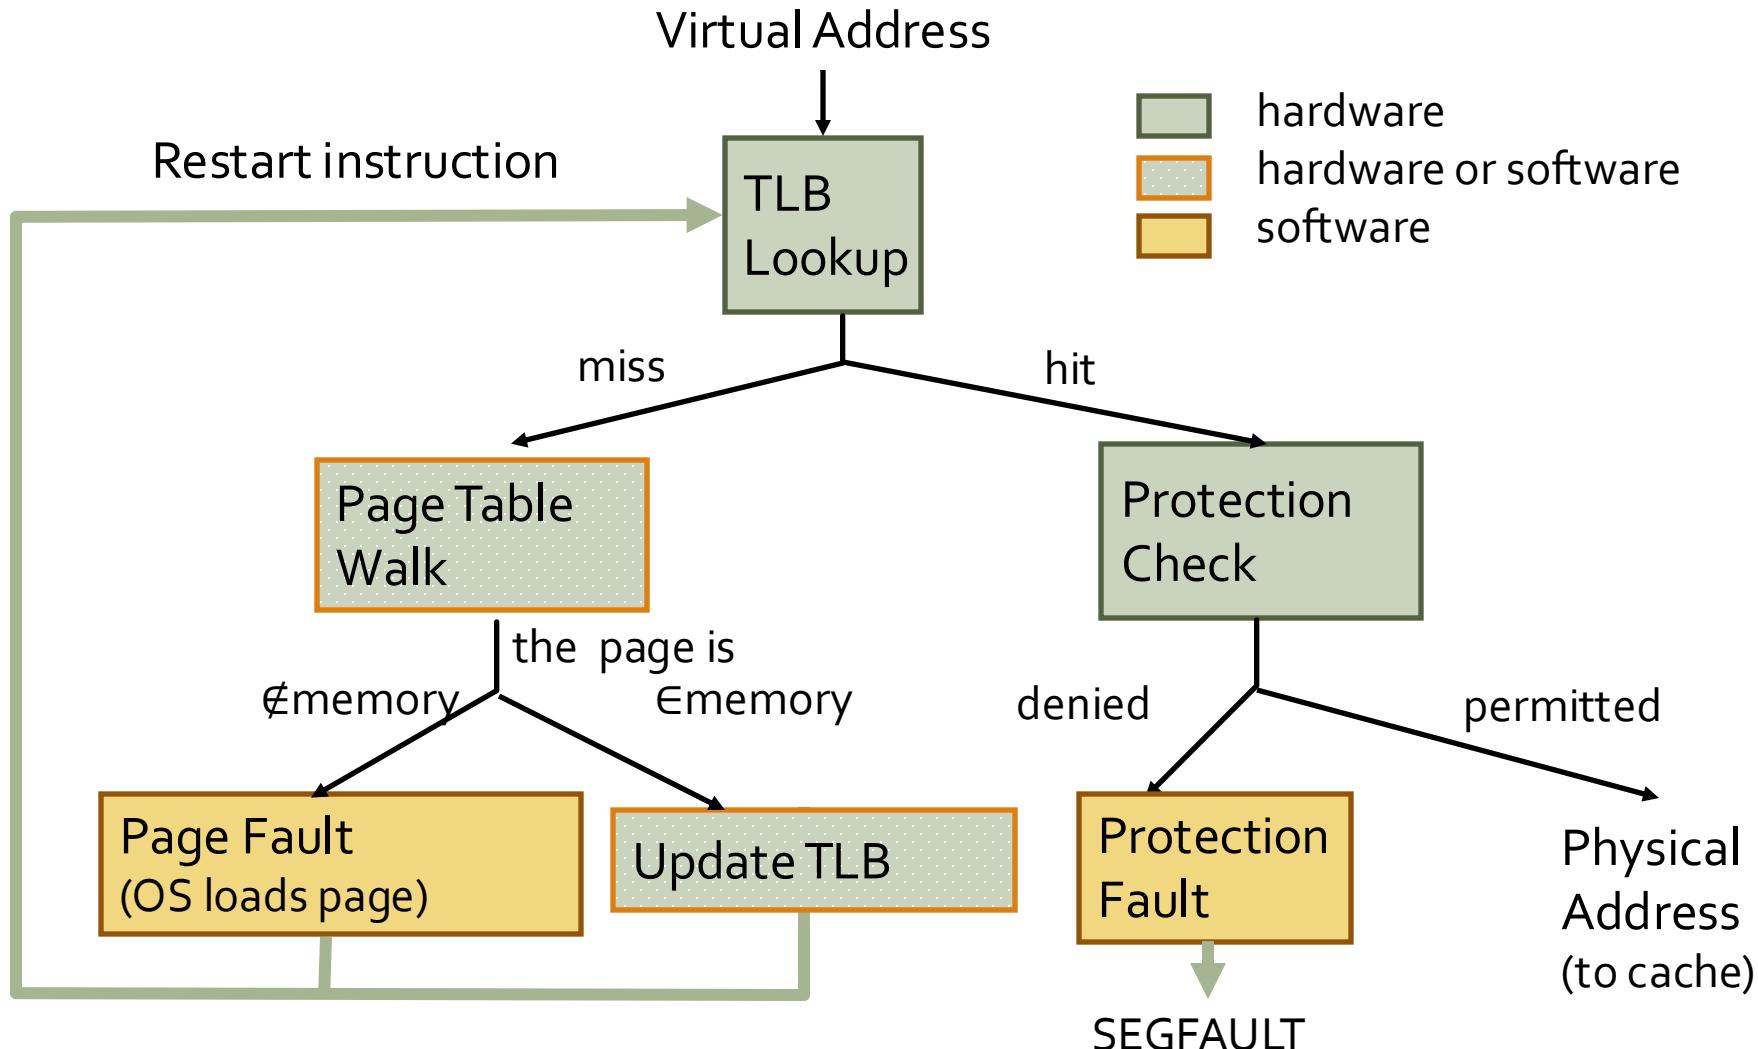


# Page-Based Virtual-Memory Machine

## (Hardware Page-Table Walk)



# Address Translation: putting it all together



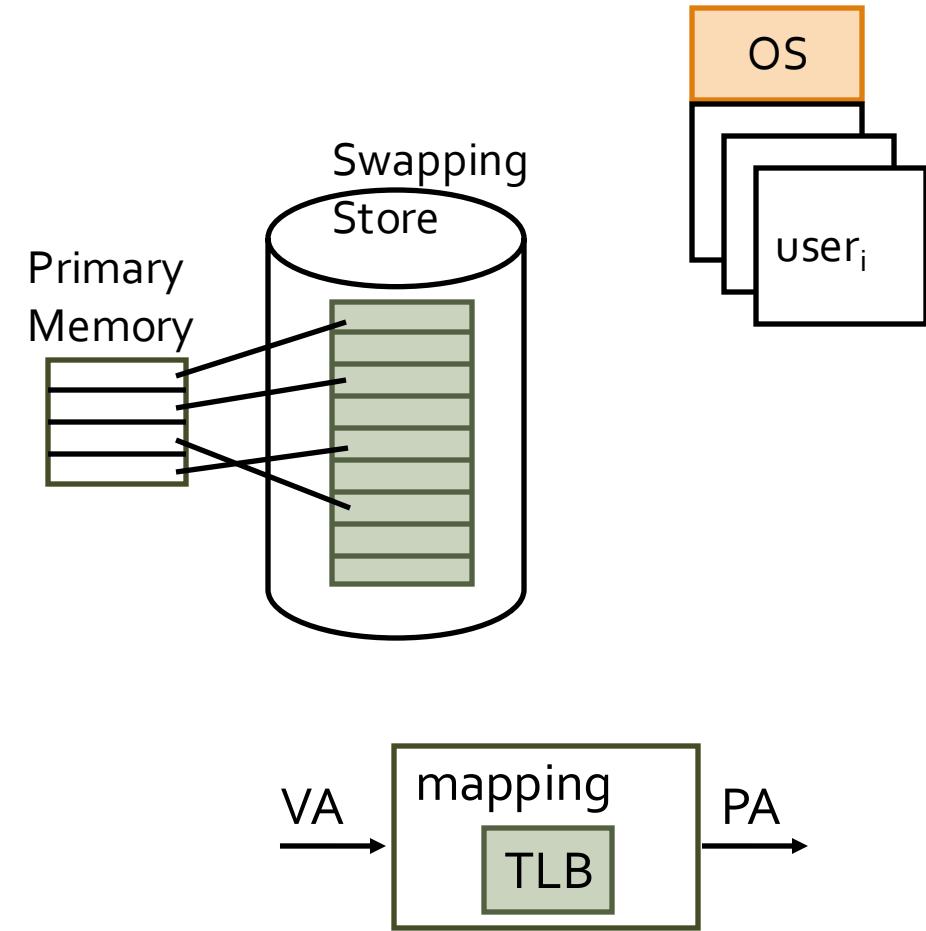
# Page Fault Handler

- When the referenced page is not in DRAM:
  - The missing page is located (or created)
  - It is brought in from disk, and page table is updated
    - Another user job may run on CPU while first job waits for the requested page to be read from disk, provided system allows architectural context to be saved and restored
  - If no free pages are left, a page is swapped out
    - Pseudo-LRU replacement policy, implemented in software
    - A set of free pages can be maintained by OS as background activity
- Since it takes a long time to transfer a page (msecs), page faults are handled completely in software by the OS

# Modern Virtual Memory Systems

Illusion of a large, private, uniform store

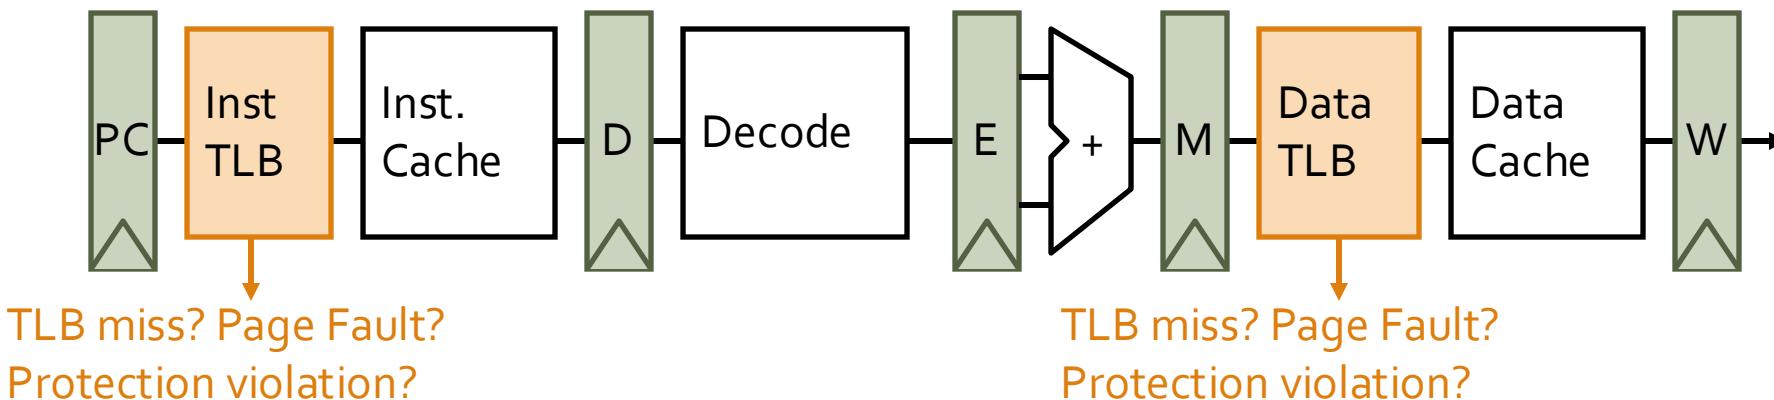
- Protection & Privacy
  - Several users, each with their private address space and one or more shared address spaces  
page table  $\equiv$  name space
- Demand Paging
  - Provides the ability to run programs larger than the primary memory
  - Hides differences in machine configurations
- The price is address translation on each memory reference



# Agenda

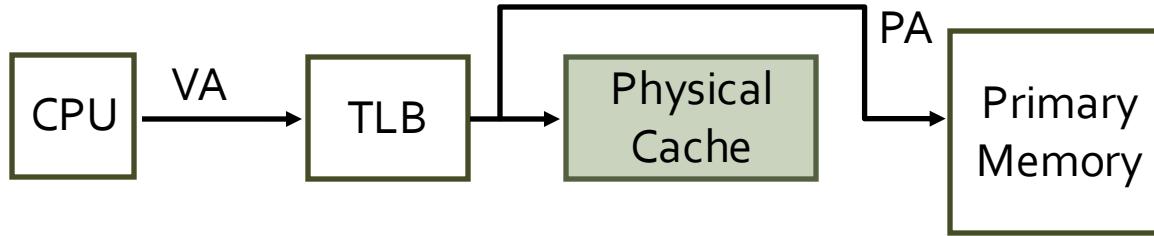
- Evolution of Virtual Memory
- Modern VM Implementation
- **TLB & Cache Organization**
- Modern Usage

# Address Translation in CPU Pipeline

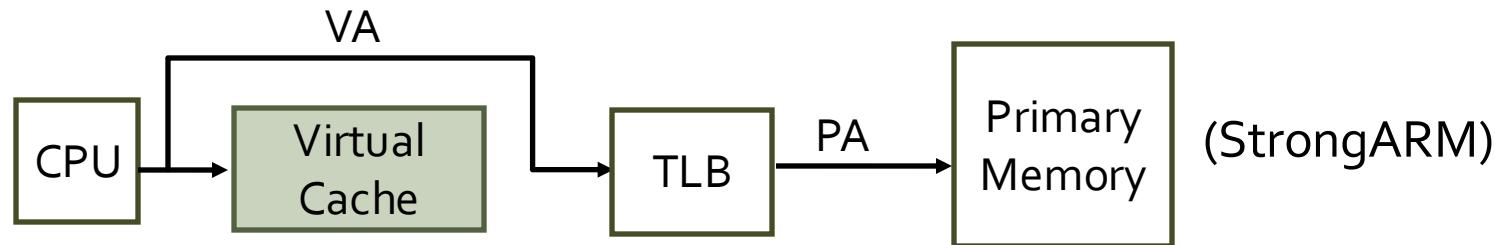


- Software handlers need **restartable** exception on TLB fault
- Handling a TLB miss needs a **hardware or software** mechanism to refill TLB
- Need to cope with additional latency of TLB:
  - slow down the clock?
  - pipeline the TLB and cache access?
  - virtual address caches
  - parallel TLB/cache access

# Virtual-Address Caches

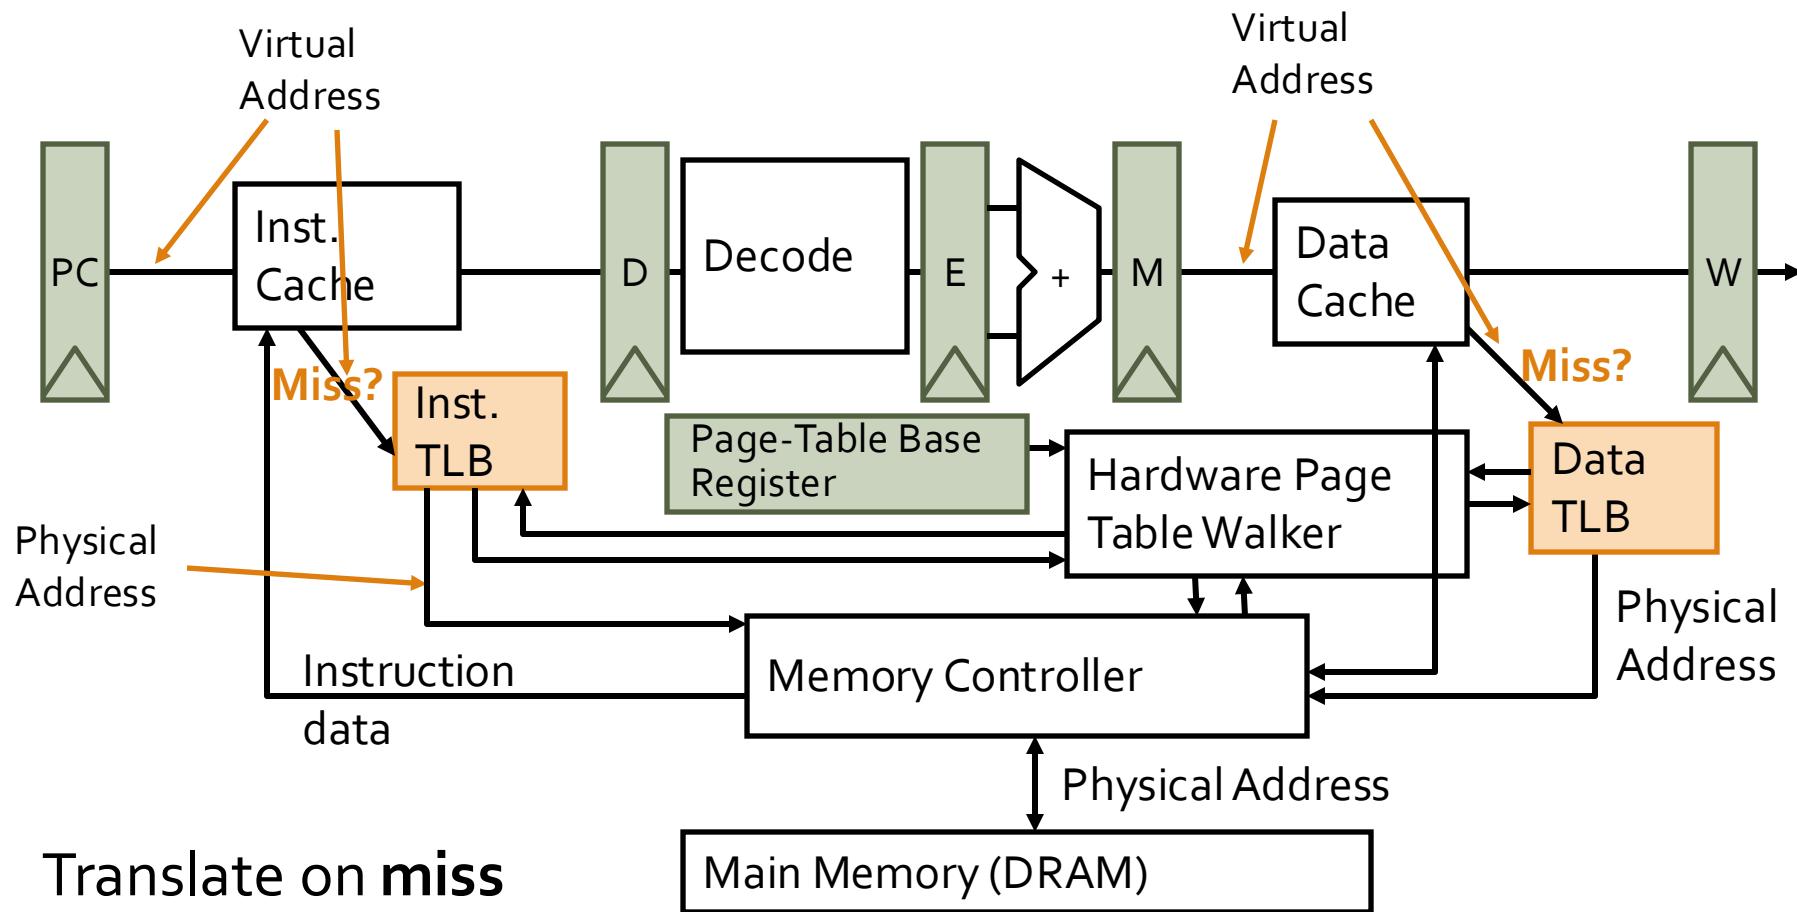


Alternative: place the cache before the TLB

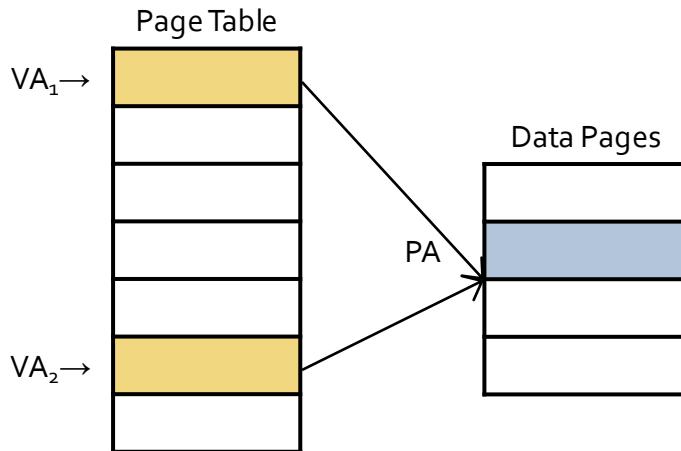


- one-step process in case of a hit (+)
- cache needs to be flushed on a context switch unless address space identifiers (ASIDs) included in tags (-)
- **aliasing problems** due to the sharing of pages (-)
- maintaining cache coherence (-) (see later in course)

# Virtually Addressed Cache (Virtual Index/Virtual Tag)



# Aliasing in Virtual-Address Caches



Two virtual pages share one physical page

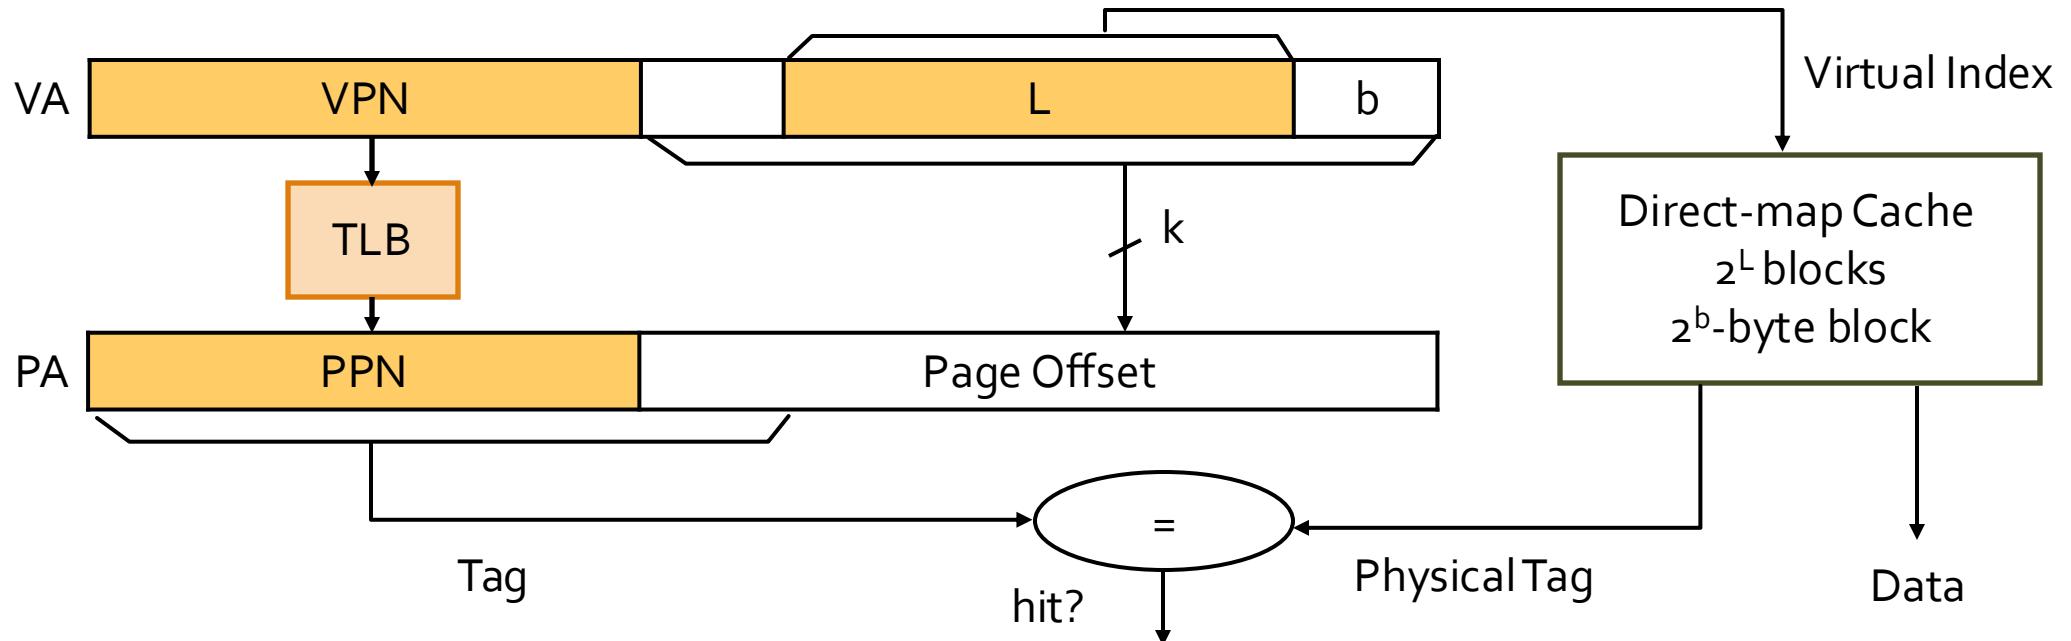
Tag	Data
$VA_1$	1st Copy of Data at PA
$VA_2$	2nd Copy of Data at PA

Virtual cache can have two copies of same physical data. Writes to one copy not visible to reads of other!

- General Solution: Prevent aliases coexisting in cache
- Software (i.e., OS) solution for direct-mapped cache
  - VAs of shared pages must agree in cache index bits; this ensures all VAs accessing same PA will conflict in direct-mapped cache (early SPARCcs)

# Concurrent Access to TLB & Cache

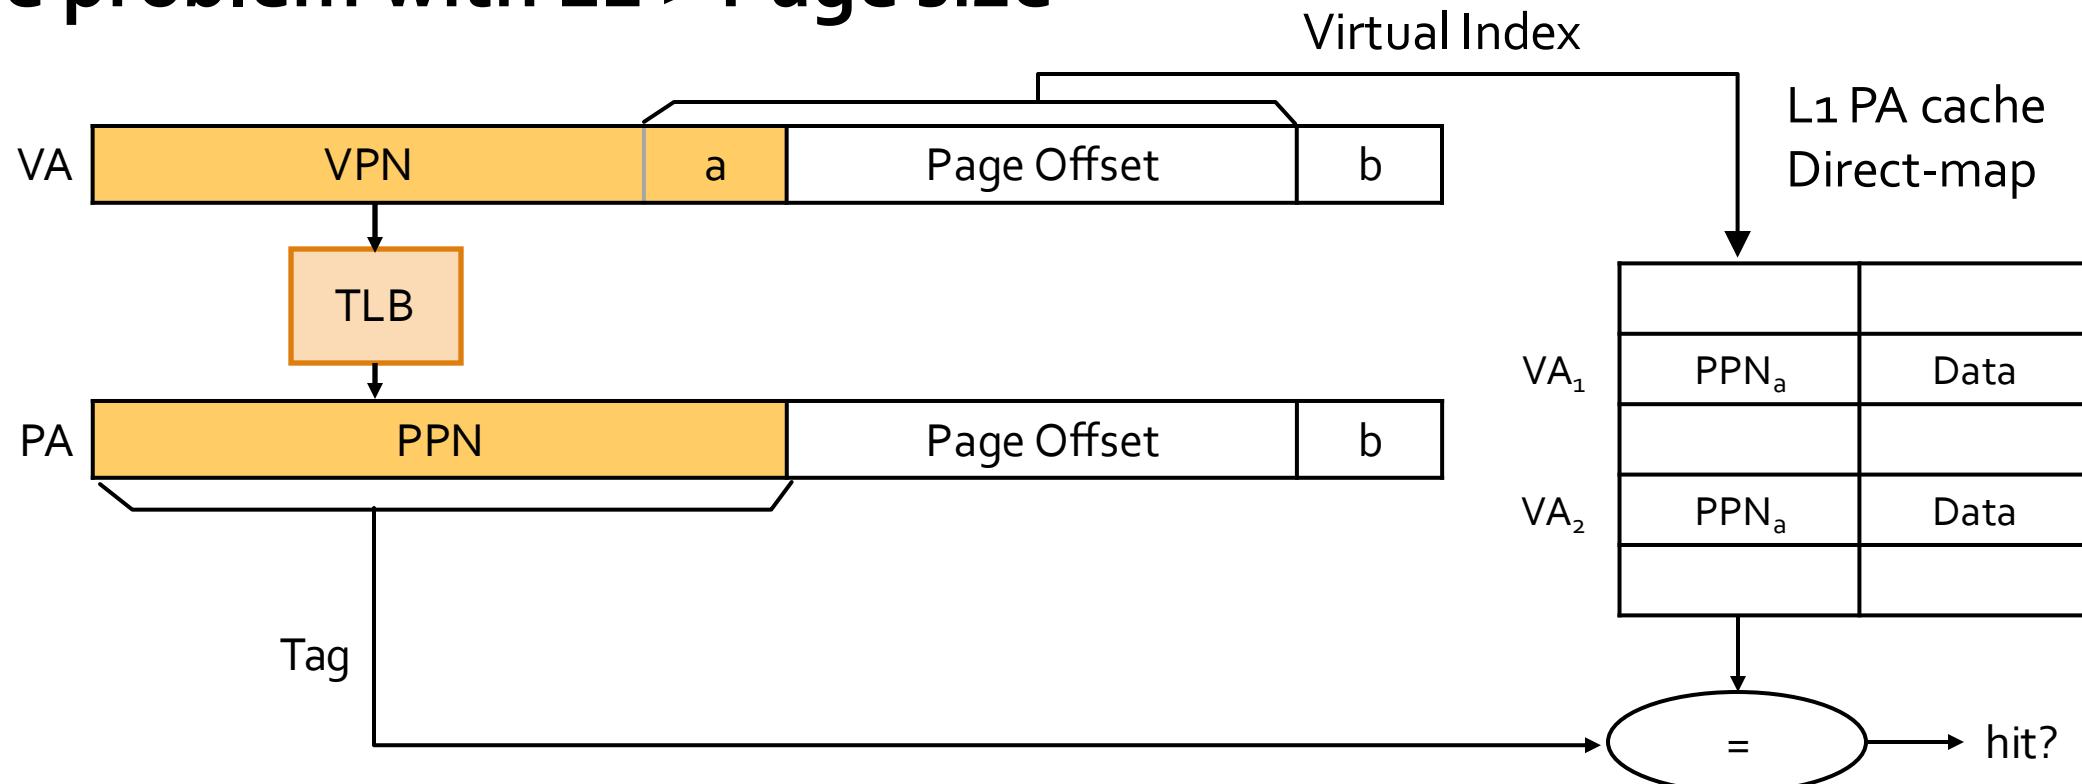
## (Virtual Index/Physical Tag)



- Index L is available without consulting the TLB
  - **cache and TLB accesses can begin simultaneously!**
- Tag comparison is made after both accesses are completed
- Cases:  $L + b = k$ ,  $L + b < k$ ,  $L + b > k$  (cache can only be as big as page size!)

# Concurrent Access to TLB & Large L1

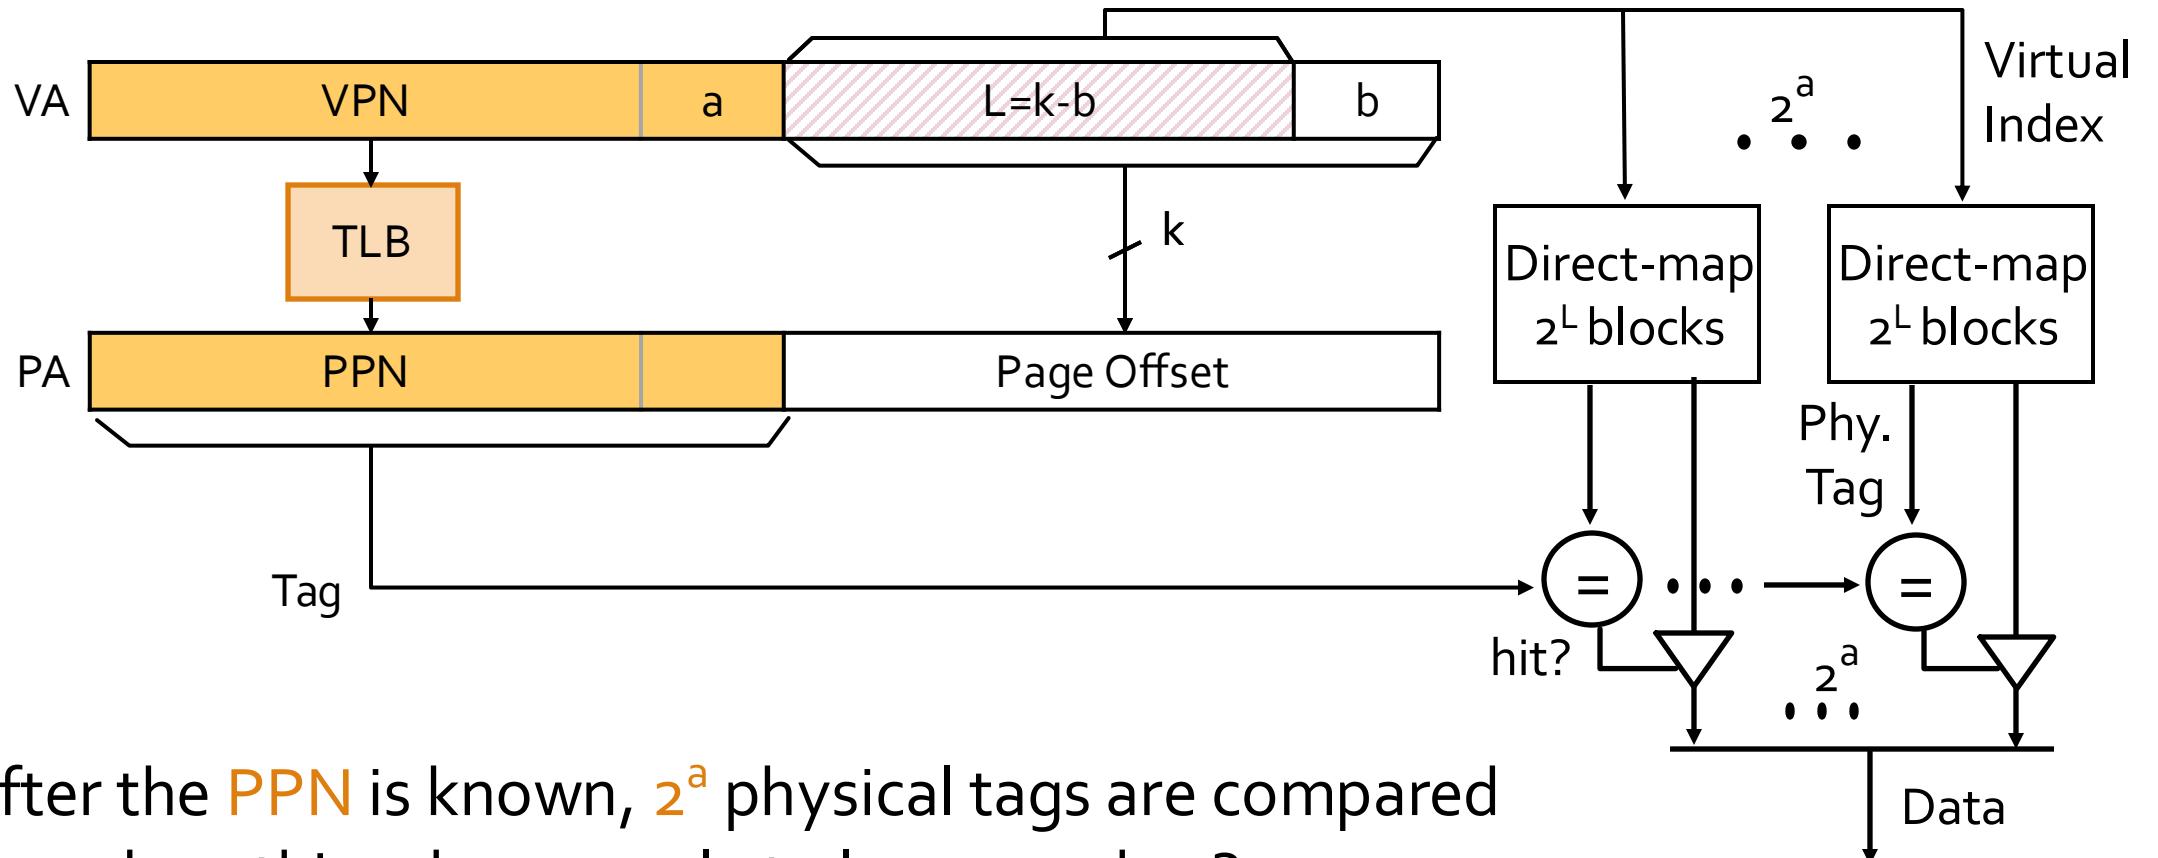
The problem with  $L_1 > \text{Page size}$



Can  $VA_1$  and  $VA_2$  both map to PA ?

# Virtual-Index Physical-Tag Caches:

## Use associativity to further increase cache capacity



- After the **PPN** is known,  $2^a$  physical tags are compared
- How does this scheme scale to larger caches?
  - 4KB page size \* 8-way ( $2^3$ ) associative = 32KB cache

# Cache-TLB Interactions

- Physically Indexed/Physically Tagged
- Virtually Indexed/Virtually Tagged
- Virtually Indexed/Physically Tagged
  - Concurrent cache access with TLB Translation
- Both Indexed/Physically Tagged
  - Small enough cache or highly associative cache will have fewer indexes than page size
  - Concurrent cache access with TLB Translation
- ~~Physically Indexed/Virtually Tagged~~

# Agenda

- Evolution of Virtual Memory
- Modern VM Implementation
- TLB & Cache Organization
- Modern Usage

# Virtual Memory Use Today - 1

- Servers/desktops/laptops/smartphones have full demand-paged virtual memory
  - Portability between machines with different memory sizes
  - Protection between multiple users or multiple tasks
  - Share small physical memory among active tasks
  - Simplifies implementation of some OS features
- Vector supercomputers have translation and protection but rarely complete demand-paging
  - (Older Crays: base&bound, Japanese & Cray X1/X2: pages)
  - Don't waste expensive CPU time thrashing to disk (make jobs fit in memory)
  - Mostly run in batch mode (run set of jobs that fits in memory)
  - Difficult to implement restartable vector instructions

# Virtual Memory Use Today - 2

- Most embedded processors and DSPs provide physical addressing only
  - Can't afford area/speed/power budget for virtual memory support
  - Often there is no secondary storage to swap to!
  - Programs custom written for particular memory configuration in product
  - Difficult to implement restartable instructions for exposed architectures

# Acknowledgements

- These slides contain material developed and copyright by:
  - Arvind (MIT)
  - Krste Asanovic (MIT/UCB)
  - Joel Emer (Intel/MIT)
  - James Hoe (CMU)
  - John Kubiatowicz (UCB)
  - David Patterson (UCB)
  - Christopher Batten (Cornell)
  - David Wentzlaff (Princeton)
- MIT material derived from course 6.823
- UCB material derived from course CS252
- Cornell material derived from course ECE 4750
- Princeton material derived from course ECE 475